
Luigi Documentation

Release 1.0

Erik Bernhardsson and Elias Freider

June 22, 2015

1	Getting Started	3
2	Documentation	5
3	More background	7
4	Dependency graph example	9
5	Background	11
6	Who uses Luigi?	13
7	Getting Help	15
8	External links	17
9	Authors	19
10	Table of Contents	21
10.1	Example – Top Artists	21
10.2	Building workflows	24
10.3	Tasks	28
10.4	Parameters	33
10.5	Running from the Command Line	35
10.6	Programmatic Execution	36
10.7	Using the Central Scheduler	36
10.8	Execution Model	40
10.9	Luigi Patterns	41
10.10	Configuration	43
10.11	More Info	49
11	API Reference	51
11.1	luigi package	51
11.2	luigi.contrib package	135
11.3	luigi.tools package	177
11.4	Indices and tables	182
	Python Module Index	183

Luigi is a Python package that helps you build complex pipelines of batch jobs. It handles dependency resolution, workflow management, visualization, handling failures, command line integration, and much more.

Run `pip install luigi` to install the latest stable version from [PyPI](#).

For bleeding edge code, `git clone https://github.com/spotify/luigi` and `python setup.py install`. Bleeding edge documentation can be found [here](#).

If you want to run the central scheduler (highly recommended), you need to install [Tornado](#) which you can do from PyPI as well: `pip install tornado`.

Getting Started

Take a look at the [Example workflow](#) and [Building workflows](#) which explains some of the most important concepts.

Documentation

Full documentation is available at [Read the Docs](#), including the [Luigi package documentation](#).

See [Configuration](#) for how to configure Luigi.

More background

The purpose of Luigi is to address all the plumbing typically associated with long-running batch processes. You want to chain many tasks, automate them, and failures *will* happen. These tasks can be anything, but are typically long running things like [Hadoop](#) jobs, dumping data to/from databases, running machine learning algorithms, or anything else.

There are other software packages that focus on lower level aspects of data processing, like [Hive](#), [Pig](#), or [Cascading](#). Luigi is not a framework to replace these. Instead it helps you stitch many tasks together, where each task can be a [Hive query](#), a [Hadoop job in Java](#), a [Spark job in Scala or Python](#) a Python snippet, [dumping a table](#) from a database, or anything else. It's easy to build up long-running pipelines that comprise thousands of tasks and take days or weeks to complete. Luigi takes care of a lot of the workflow management so that you can focus on the tasks themselves and their dependencies.

You can build pretty much any task you want, but Luigi also comes with a *toolbox* of several common task templates that you use. It includes support for running [Python mapreduce jobs](#) in Hadoop, as well as [Hive](#), and [Pig](#), jobs. It also comes with [file system abstractions for HDFS](#), and local files that ensures all file system operations are atomic. This is important because it means your data pipeline will not crash in a state containing partial data.

Dependency graph example

Just to give you an idea of what Luigi does, this is a screen shot from something we are running in production. Using Luigi's visualizer, we get a nice visual overview of the dependency graph of the workflow. Each node represents a task which has to be run. Green tasks are already completed whereas yellow tasks are yet to be run. Most of these tasks are Hadoop jobs, but there are also some things that run locally and build up data files.

Background

We use Luigi internally at [Spotify](#) to run thousands of tasks every day, organized in complex dependency graphs. Most of these tasks are Hadoop jobs. Luigi provides an infrastructure that powers all kinds of stuff including recommendations, toplists, A/B test analysis, external reports, internal dashboards, etc. Luigi grew out of the realization that powerful abstractions for batch processing can help programmers focus on the most important bits and leave the rest (the boilerplate) to the framework.

Conceptually, Luigi is similar to [GNU Make](#) where you have certain tasks and these tasks in turn may have dependencies on other tasks. There are also some similarities to [Oozie](#) and [Azkaban](#). One major difference is that Luigi is not just built specifically for Hadoop, and it's easy to extend it with other kinds of tasks.

Everything in Luigi is in Python. Instead of XML configuration or similar external data files, the dependency graph is specified *within Python*. This makes it easy to build up complex dependency graphs of tasks, where the dependencies can involve date algebra or recursive references to other versions of the same task. However, the workflow can trigger things not in Python, such as running [Pig scripts](#) or [scp'ing files](#).

Who uses Luigi?

Several companies have written blog posts or presentation about Luigi:

- Spotify (NYC Data Science)
- Foursquare
- Mortar Data
- Stripe
- Asana
- Buffer
- SeatGeek
- Treasure Data

Please let us know if your company wants to be featured on this list!

Getting Help

- Find us on *#luigi* on freenode.
- Subscribe to the [luigi-user](#) group and ask a question.

External links

- [Mailing List](#) (Google Groups)
- [Releases](#) (PyPi)
- [Source code](#) (Github)

Authors

Luigi was built at [Spotify](#), mainly by [Erik Bernhardsson](#) and [Elias Freider](#). Many other people have contributed since open sourcing in late 2012. [Arash Rouhani](#) is currently the chief maintainer of Luigi.

Table of Contents

10.1 Example – Top Artists

This is a very simplified case of something we do at Spotify a lot. All user actions are logged to HDFS where we run a bunch of Hadoop jobs to transform the data. At some point we might end up with a smaller data set that we can bulk ingest into Cassandra, Postgres, or some other format.

For the purpose of this exercise, we want to aggregate all streams, find the top 10 artists and then put the results into Postgres.

This example is also available in [examples/top_artists.py](#).

10.1.1 Step 1 - Aggregate Artist Streams

```
class AggregateArtists(luigi.Task):
    date_interval = luigi.DateIntervalParameter()

    def output(self):
        return luigi.LocalTarget("data/artist_streams_%s.tsv" % self.date_interval)

    def requires(self):
        return [Streams(date) for date in self.date_interval]

    def run(self):
        artist_count = defaultdict(int)

        for input in self.input():
            with input.open('r') as in_file:
                for line in in_file:
                    timestamp, artist, track = line.strip().split()
                    artist_count[artist] += 1

        with self.output().open('w') as out_file:
            for artist, count in artist_count.iteritems():
                print >> out_file, artist, count
```

Note that this is just a portion of the file *examples/top_artists.py*. In particular, *Streams* is defined as a *Task*, acting as a dependency for *AggregateArtists*. In addition, `luigi.run()` is called if the script is executed directly, allowing it to be run from the command line.

There are several pieces of this snippet that deserve more explanation.

- Any *Task* may be customized by instantiating one or more *Parameter* objects on the class level.
- The *output()* method tells Luigi where the result of running the task will end up. The path can be some function of the parameters.
- The *requires()* tasks specifies other tasks that we need to perform this task. In this case it's an external dump named *Streams* which takes the date as the argument.
- For plain Tasks, the *run()* method implements the task. This could be anything, including calling subprocesses, performing long running number crunching, etc. For some subclasses of *Task* you don't have to implement the run method. For instance, for the *JobTask* subclass you implement a *mapper* and *reducer* instead.
- *HdfsTarget* is a built in class that makes it easy to read/write from/to HDFS. It also makes all file operations atomic, which is nice in case your script crashes for any reason.

10.1.2 Running this Locally

Try running this using eg.

```
$ python examples/top_artists.py AggregateArtists --local-scheduler --date-interval 2012-06
```

You can also try to view the manual using *-help* which will give you an overview of the options:

```
usage: wordcount.py [-h] [--local-scheduler] [--scheduler-host SCHEDULER_HOST]
                  [--lock] [--lock-pid-dir LOCK_PID_DIR] [--workers WORKERS]
                  [--date-interval DATE_INTERVAL]

optional arguments:
  -h, --help                show this help message and exit
  --local-scheduler          Use local scheduling
  --scheduler-host SCHEDULER_HOST
                           Hostname of machine running remote scheduler [default:
                           localhost]
  --lock                    Do not run if the task is already running
  --lock-pid-dir LOCK_PID_DIR
                           Directory to store the pid file [default:
                           /var/tmp/luigi]
  --workers WORKERS         Maximum number of parallel tasks to run [default: 1]
  --date-interval DATE_INTERVAL
                           AggregateArtists.date_interval
```

Running the command again will do nothing because the output file is already created. In that sense, any task in Luigi is *idempotent* because running it many times gives the same outcome as running it once. Note that unlike Makefile, the output will not be recreated when any of the input files is modified. You need to delete the output file manually.

The *--local-scheduler* flag tells Luigi not to connect to a scheduler server. This is not recommended for other purpose than just testing things.

10.1.3 Step 1b - Running this in Hadoop

Luigi comes with native Python Hadoop mapreduce support built in, and here is how this could look like, instead of the class above.

```
class AggregateArtistsHadoop(luigi.contrib.hadoop.JobTask):
    date_interval = luigi.DateIntervalParameter()

    def output(self):
        return luigi.contrib.hdfs.HdfsTarget("data/artist_streams_%s.tsv" % self.date_interval)
```

```

def requires(self):
    return [StreamsHdfs(date) for date in self.date_interval]

def mapper(self, line):
    timestamp, artist, track = line.strip().split()
    yield artist, 1

def reducer(self, key, values):
    yield key, sum(values)

```

Note that `luigi.contrib.hadoop.JobTask` doesn't require you to implement a `run()` method. Instead, you typically implement a `mapper()` and `reducer()` method.

10.1.4 Step 2 – Find the Top Artists

At this point, we've counted the number of streams for each artists, for the full time period. We are left with a large file that contains mappings of artist -> count data, and we want to find the top 10 artists. Since we only have a few hundred thousand artists, and calculating artists is nontrivial to parallelize, we choose to do this not as a Hadoop job, but just as a plain old for-loop in Python.

```

class Top10Artists(luigi.Task):
    date_interval = luigi.DateIntervalParameter()
    use_hadoop = luigi.BoolParameter()

    def requires(self):
        if self.use_hadoop:
            return AggregateArtistsHadoop(self.date_interval)
        else:
            return AggregateArtists(self.date_interval)

    def output(self):
        return luigi.LocalTarget("data/top_artists_%s.tsv" % self.date_interval)

    def run(self):
        top_10 = nlargest(10, self._input_iterator())
        with self.output().open('w') as out_file:
            for streams, artist in top_10:
                print >> out_file, self.date_interval.date_a, self.date_interval.date_b, artist, streams

    def _input_iterator(self):
        with self.input().open('r') as in_file:
            for line in in_file:
                artist, streams = line.strip().split()
                yield int(streams), int(artist)

```

The most interesting thing here is that this task (*Top10Artists*) defines a dependency on the previous task (*AggregateArtists*). This means that if the output of *AggregateArtists* does not exist, the task will run before *Top10Artists*.

```
$ python examples/top_artists.py Top10Artists --local-scheduler --date-interval 2012-07
```

This will run both tasks.

10.1.5 Step 3 - Insert into Postgres

This mainly serves as an example of a specific subclass *Task* that doesn't require any code to be written. It's also an example of how you can define task templates that you can reuse for a lot of different tasks.

```
class ArtistToplistToDatabase(luigi.postgres.CopyToTable):
    date_interval = luigi.DateIntervalParameter()
    use_hadoop = luigi.BoolParameter()

    host = "localhost"
    database = "toplists"
    user = "luigi"
    password = "abc123" # ;)
    table = "top10"

    columns = [("date_from", "DATE"),
               ("date_to", "DATE"),
               ("artist", "TEXT"),
               ("streams", "INT")]

    def requires(self):
        return Top10Artists(self.date_interval, self.use_hadoop)
```

Just like previously, this defines a recursive dependency on the previous task. If you try to build the task, that will also trigger building all its upstream dependencies.

10.1.6 Using the Central Planner

The `-local-scheduler` flag tells Luigi not to connect to a central scheduler. This is recommended in order to get started and or for development purposes. At the point where you start putting things in production we strongly recommend running the central scheduler server. In addition to providing locking so that the same task is not run by multiple processes at the same time, this server also provides a pretty nice visualization of your current work flow.

If you drop the `-local-scheduler` flag, your script will try to connect to the central planner, by default at localhost port 8082. If you run

```
luigid
```

in the background and then run

```
$ python wordcount.py --date 2012-W03
```

then in fact your script will now do the scheduling through a centralized server. You need [Tornado](#) for this to work.

Launching `http://localhost:8082` should show something like this:

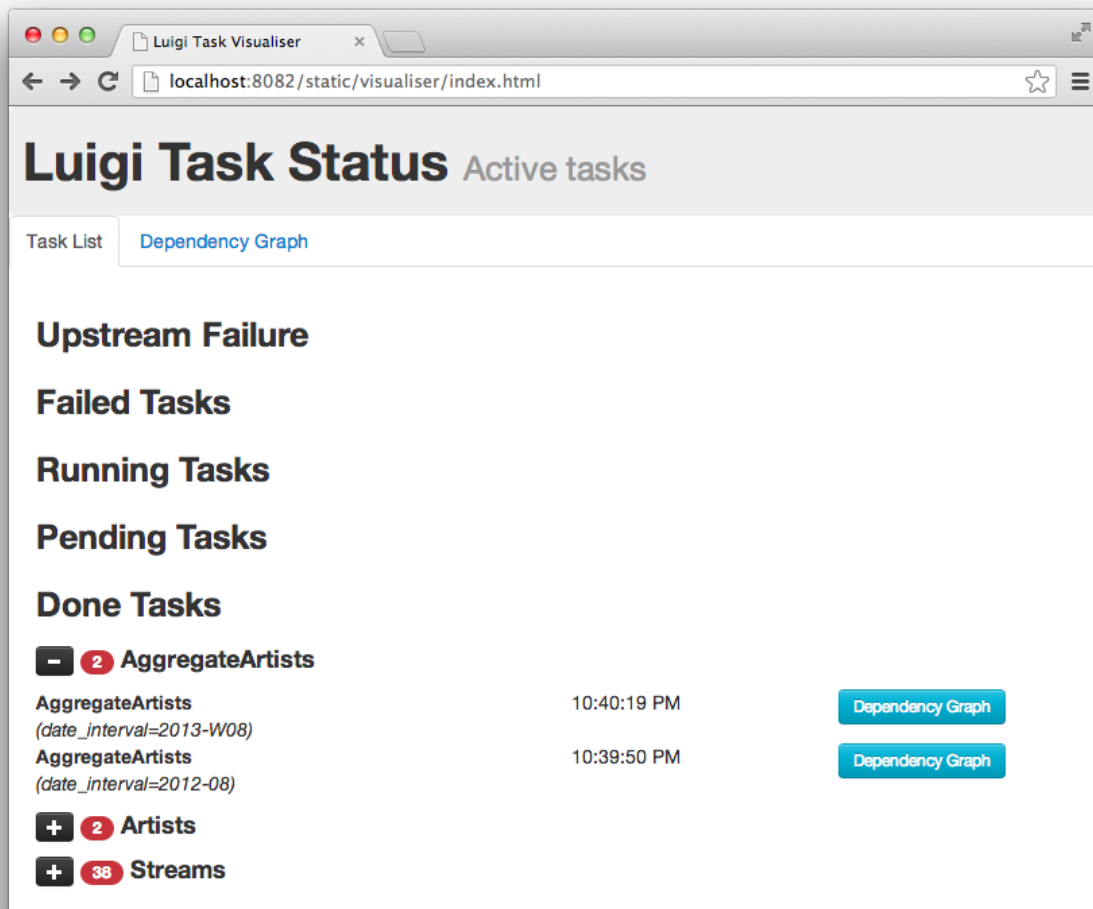
Web server screenshot Looking at the dependency graph for any of the tasks yields something like this:

Aggregate artists screenshot

In production, you'll want to run the centralized scheduler. See: [Using the Central Scheduler](#) for more information.

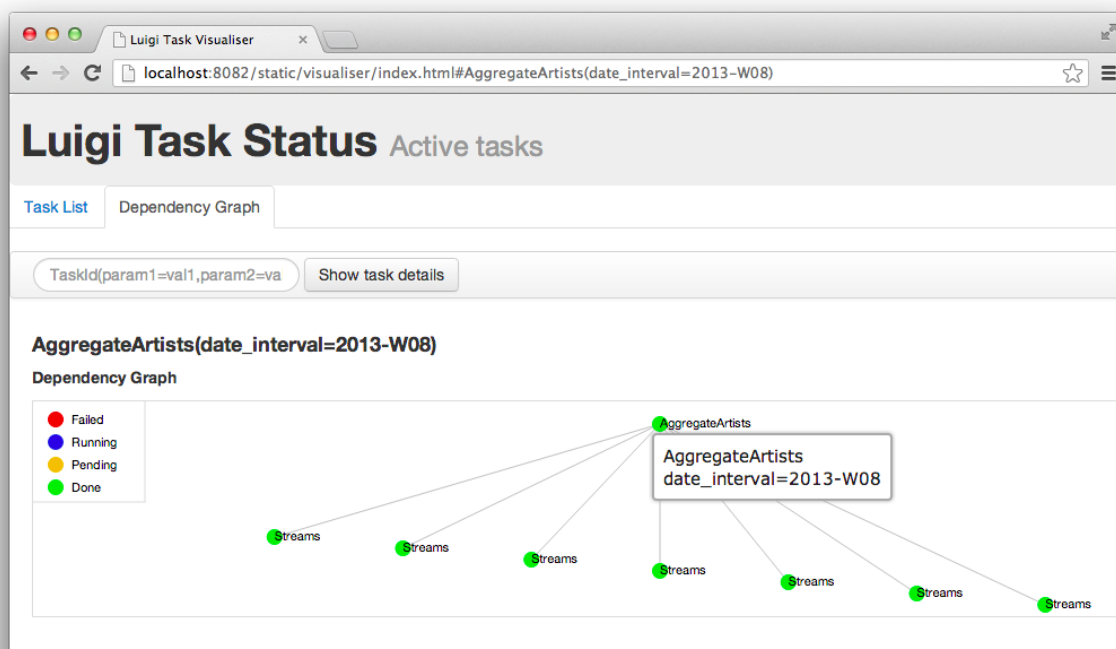
10.2 Building workflows

There are two fundamental building blocks of Luigi - the `Task` class and the `Target` class. Both are abstract classes and expect a few methods to be implemented. In addition to those two concepts, the `Parameter` class is an important concept that governs how a Task is run.



The screenshot shows a web browser window titled "Luigi Task Visualiser" with the URL "localhost:8082/static/visualiser/index.html". The page has a header "Luigi Task Status" and a sub-header "Active tasks". Below the header, there are two tabs: "Task List" (selected) and "Dependency Graph". The main content area lists several task categories: "Upstream Failure", "Failed Tasks", "Running Tasks", "Pending Tasks", and "Done Tasks". Under "Done Tasks", there are three expandable sections: "AggregateArtists" (2 tasks), "Artists" (2 tasks), and "Streams" (38 tasks). Each section has a plus icon and a red circle with the task count. The "AggregateArtists" section is expanded, showing two tasks with their timestamps and a "Dependency Graph" button.

Task Name	Timestamp	Action
AggregateArtists (date_interval=2013-W08)	10:40:19 PM	Dependency Graph
AggregateArtists (date_interval=2012-08)	10:39:50 PM	Dependency Graph



10.2.1 Target

The *Target* class corresponds to a file on a disk, a file on HDFS or some kind of a checkpoint, like an entry in a database. Actually, the only method that Targets have to implement is the *exists* method which returns True if and only if the Target exists.

In practice, implementing Target subclasses is rarely needed. Luigi comes a toolbox of several useful Targets. In particular, *LocalTarget* and *HdfsTarget*, but there is also support for other file systems: *luigi.s3.S3Target*, *luigi.contrib.ssh.RemoteTarget*, *luigi.ftp.RemoteTarget*, *luigi.contrib.mysql.MySqlTarget*, *luigi.redshift.RedshiftTarget*, and several more.

Most of these targets, are file system-like. For instance, *LocalTarget* and *HdfsTarget* map to a file on the local drive or a file in HDFS. In addition these also wrap the underlying operations to make them atomic. They both implement the *open()* method which returns a stream object that could be read (*mode='r'*) from or written to (*mode='w'*).

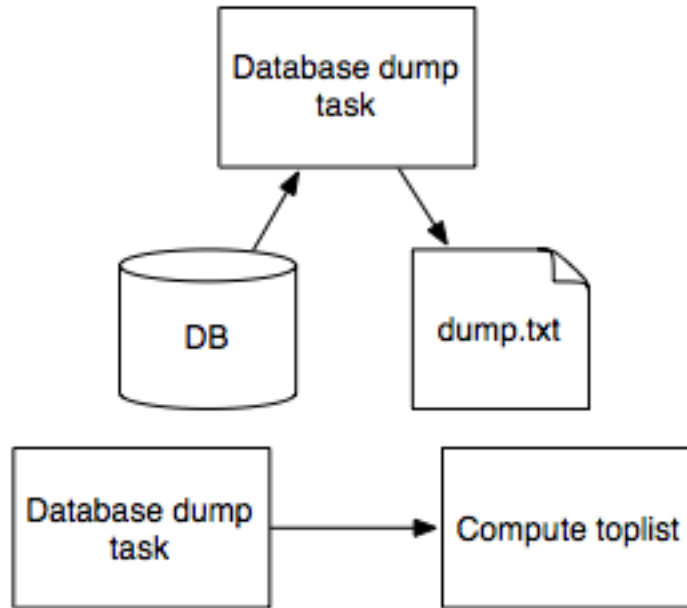
Luigi comes with Gzip support by providing *format=format.Gzip*. Adding support for other formats is pretty simple.

10.2.2 Task

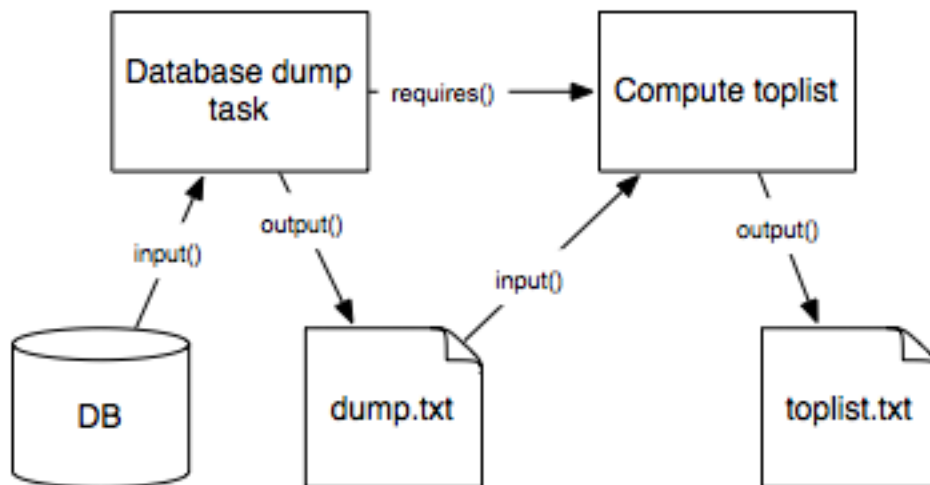
The *Task* class is a bit more conceptually interesting because this is where computation is done. There are a few methods that can be implemented to alter its behavior, most notably *run()*, *output()* and *requires()*.

Tasks consume Targets that were created by some other task. They usually also output targets:

You can define dependencies between *Tasks* using the *requires()* method. See [Tasks](#) for more info.



Each task defines its outputs using the `output()` method. Additionally, there is a helper method `input()` that returns the corresponding Target classes for each Task dependency.



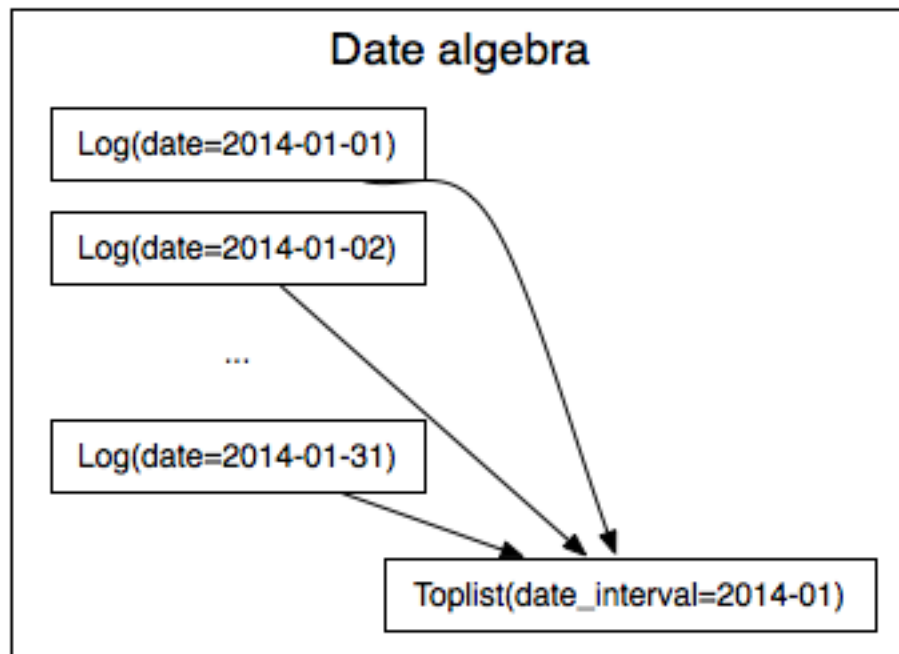
10.2.3 Parameter

The Task class corresponds to some type of job that is run, but in general you want to allow some form of parametrization of it. For instance, if your Task class runs a Hadoop job to create a report every night, you probably want to make the date a parameter of the class. See [Parameters](#) for more info.



10.2.4 Dependencies

Using tasks, targets, and parameters, Luigi lets you express arbitrary dependencies in *code*, rather than using some kind of awkward config DSL. This is really useful because in the real world, dependencies are often very messy. For instance, some examples of the dependencies you might encounter:



(These diagrams are from a Luigi presentation in late 2014 at NYC Data Science meetup)

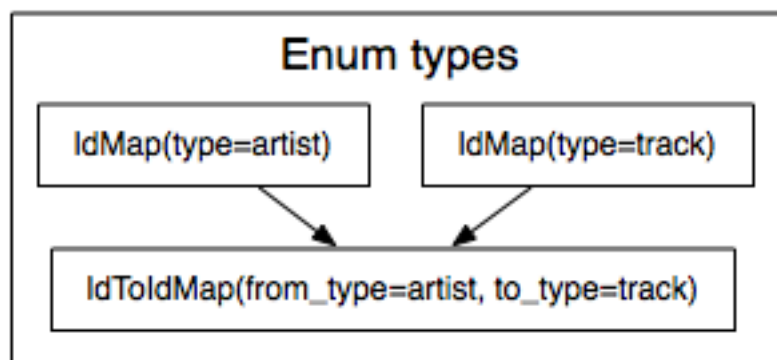
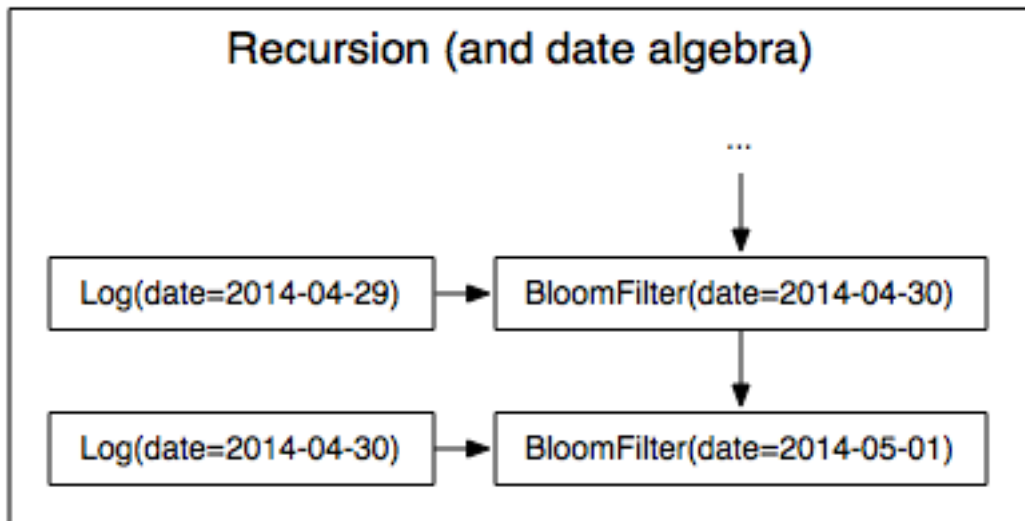
10.3 Tasks

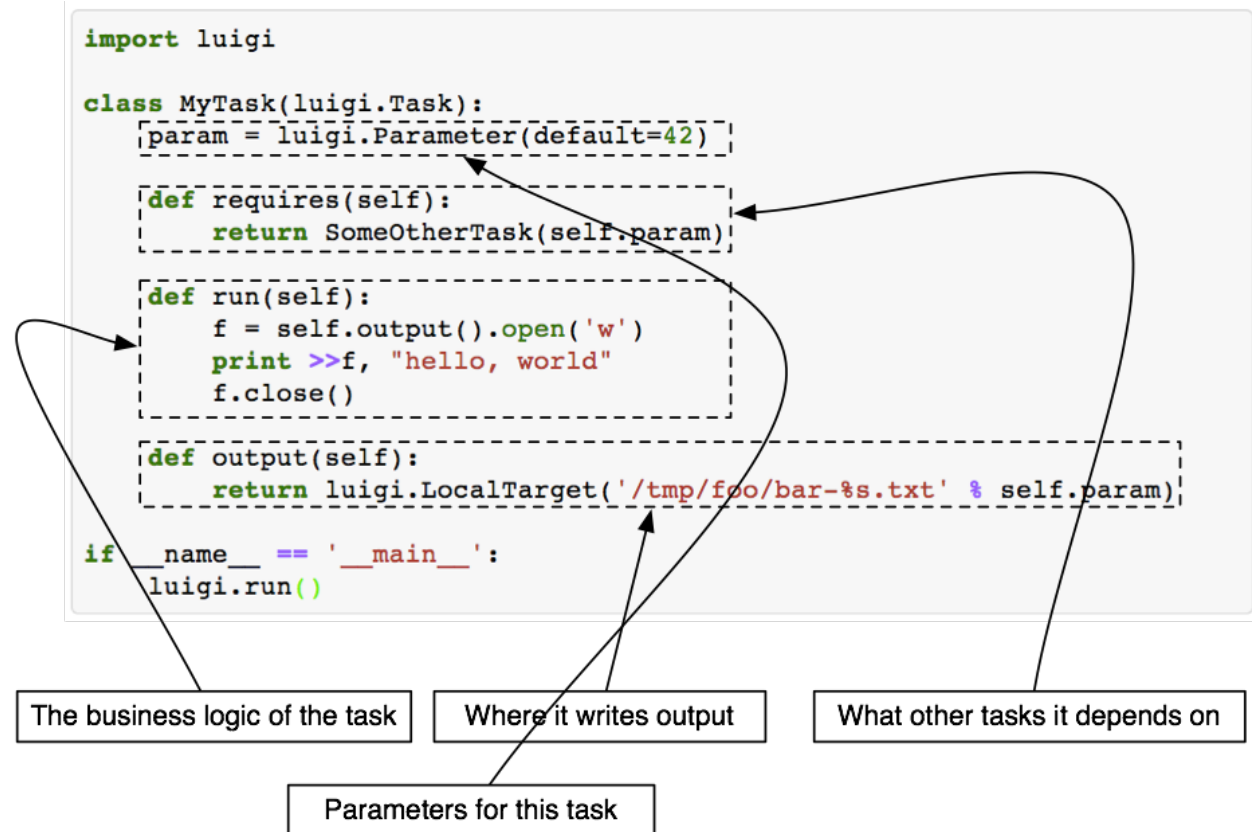
Tasks are where the execution takes place. Tasks depend on each other and output targets.

An outline of how a task can look like:

10.3.1 Task.requires

The `requires()` method is used to specify dependencies on other Task object, which might even be of the same class. For instance, an example implementation could be





```

def requires(self):
    return OtherTask(self.date), DailyReport(self.date - datetime.timedelta(1))

```

In this case, the `DailyReport` task depends on two inputs created earlier, one of which is the same class. `requires` can return other Tasks in any way wrapped up within dicts/lists/tuples/etc.

10.3.2 Requiring another Task

Note that `requires()` can *not* return a `Target` object. If you have a simple `Target` object that is created externally you can wrap it in a Task class like this:

```

class LogFiles(luigi.Task):
    def output(self):
        return luigi.contrib.hdfs.HdfsTarget('/log')

```

This also makes it easier to add parameters:

```

class LogFiles(luigi.Task):
    date = luigi.DateParameter()
    def output(self):
        return luigi.contrib.hdfs.HdfsTarget(self.date.strftime('/log/%Y-%m-%d'))

```

10.3.3 Task.output

The `output()` method returns one or more `Target` objects. Similarly to `requires`, can return wrap them up in any way that's convenient for you. However we recommend that any `Task` only return one single `Target` in output. If

multiple outputs are returned, atomicity will be lost unless the *Task* itself can ensure that each *Target* is atomically created. (If atomicity is not of concern, then it is safe to return multiple *Target* objects.)

```
class DailyReport(luigi.Task):
    date = luigi.DateParameter()
    def output(self):
        return luigi.contrib.hdfs.HdfsTarget(self.date.strftime('/reports/%Y-%m-%d'))
    # ...
```

10.3.4 Task.run

The *run()* method now contains the actual code that is run. When you are using *Task.requires* and *Task.run* Luigi breaks down everything into two stages. First it figures out all dependencies between tasks, then it runs everything. The *input()* method is an internal helper method that just replaces all Task objects in requires with their corresponding output. An example:

```
class TaskA(luigi.Task):
    def output(self):
        return luigi.LocalTarget('xyz')

class FlipLinesBackwards(luigi.Task):
    def requires(self):
        return TaskA()

    def output(self):
        return luigi.LocalTarget('abc')

    def run(self):
        f = self.input().open('r') # this will return a file stream that reads from "xyz"
        g = self.output().open('w')
        for line in f:
            g.write('%s\n', ''.join(reversed(line.strip().split())))
        g.close() # needed because files are atomic
```

10.3.5 Task.input

As seen in the example above, *Task* is a wrapper around *Task.requires* that returns the corresponding Target objects instead of Task objects. Anything returned by *Task.requires* will be transformed, including lists, nested dicts, etc. This can be useful if you have many dependencies:

```
class TaskWithManyInputs(luigi.Task):
    def requires(self):
        return {'a': TaskA(), 'b': [TaskB(i) for i in xrange(100)]}

    def run(self):
        f = self.input()['a'].open('r')
        g = [y.open('r') for y in self.input()['b']]
```

10.3.6 Dynamic dependencies

Sometimes you might not know exactly what other tasks to depend on until runtime. In that case, Luigi provides a mechanism to specify dynamic dependencies. If you yield another *Task* in the *Task.run* method, the current task will be suspended and the other task will be run. You can also return a list of tasks.

```
class MyTask(luigi.Task):
    def run(self):
        other_target = yield OtherTask()

        # dynamic dependencies resolve into targets
        f = other_target.open('r')
```

This mechanism is an alternative to *Task.requires* in case you are not able to build up the full dependency graph before running the task. It does come with some constraints: the *Task.run* method will resume from scratch each time a new task is yielded. In other words, you should make sure your *Task.run* method is idempotent. (This is good practice for all Tasks in Luigi, but especially so for tasks with dynamic dependencies).

For an example of a workflow using dynamic dependencies, see [examples/dynamic_requirements.py](#).

10.3.7 Events and callbacks

Luigi has a built-in event system that allows you to register callbacks to events and trigger them from your own tasks. You can both hook into some pre-defined events and create your own. Each event handle is tied to a Task class and will be triggered only from that class or a subclass of it. This allows you to effortlessly subscribe to events only from a specific class (e.g. for hadoop jobs).

```
@luigi.Task.event_handler(luigi.Event.SUCCESS)
def celebrate_success(task):
    """Will be called directly after a successful execution
    of `run` on any Task subclass (i.e. all luigi Tasks)
    """
    ...

@luigi.contrib.hadoop.JobTask.event_handler(luigi.Event.FAILURE)
def mourn_failure(task, exception):
    """Will be called directly after a failed execution
    of `run` on any JobTask subclass
    """
    ...

luigi.run()
```

10.3.8 But I just want to run a Hadoop job?

The Hadoop code is integrated in the rest of the Luigi code because we really believe almost all Hadoop jobs benefit from being part of some sort of workflow. However, in theory, nothing stops you from using the *JobTask* class (and also *HdfsTarget*) without using the rest of Luigi. You can simply run it manually using

```
MyJobTask('abc', 123).run()
```

You can use the *hdfs.HdfsTarget* class anywhere by just instantiating it:

```
t = luigi.contrib.hdfs.HdfsTarget('/tmp/test.gz', format=format.Gzip)
f = t.open('w')
# ...
f.close() # needed
```

10.3.9 Task priority

The scheduler decides which task to run next from the set of all task that have all their dependencies met. By default, this choice is pretty arbitrary, which is fine for most workflows and situations.

If you want to have some control on the order of execution of available tasks, you can set the `priority` property of a task, for example as follows:

```
# A static priority value as a class constant:
class MyTask(luigi.Task):
    priority = 100
    # ...

# A dynamic priority value with a "@property" decorated method:
class OtherTask(luigi.Task):
    @property
    def priority(self):
        if self.date > some_threshold:
            return 80
        else:
            return 40
    # ...
```

Tasks with a higher priority value will be picked before tasks with a lower priority value. There is no predefined range of priorities, you can choose whatever (int or float) values you want to use. The default value is 0.

Warning: task execution order in Luigi is influenced by both dependencies and priorities, but in Luigi dependencies come first. For example: if there is a task A with priority 1000 but still with unmet dependencies and a task B with priority 1 without any pending dependencies, task B will be picked first.

10.3.10 Instance caching

In addition to the stuff mentioned above, Luigi also does some metaclass logic so that if e.g. `DailyReport(datetime.date(2012, 5, 10))` is instantiated twice in the code, it will in fact result in the same object. See [Instance caching](#) for more info

10.4 Parameters

Parameters is the Luigi equivalent of creating a constructor for each Task. Luigi requires you to declare these parameters instantiating `Parameter` objects on the class scope:

```
class DailyReport(luigi.contrib.hadoop.JobTask):
    date = luigi.DateParameter(default=datetime.date.today())
    # ...
```

By doing this, Luigi can do take care of all the boilerplate code that would normally be needed in the constructor. Internally, the `DailyReport` object can now be constructed by running `DailyReport(datetime.date(2012, 5, 10))` or just `DailyReport()`. Luigi also creates a command line parser that automatically handles the conversion from strings to Python types. This way you can invoke the job on the command line eg. by passing `--date 2012-15-10`.

The parameters are all set to their values on the Task object instance, i.e.

```
d = DailyReport(datetime.date(2012, 5, 10))
print d.date
```

will return the same date that the object was constructed with. Same goes if you invoke Luigi on the command line.

10.4.1 Instance caching

Tasks are uniquely identified by their class name and values of their parameters. In fact, within the same worker, two tasks of the same class with parameters of the same values are not just equal, but the same instance:

```
>>> import luigi
>>> import datetime
>>> class DateTask(luigi.Task):
...     date = luigi.DateParameter()
...
>>> a = datetime.date(2014, 1, 21)
>>> b = datetime.date(2014, 1, 21)
>>> a is b
False
>>> c = DateTask(date=a)
>>> d = DateTask(date=b)
>>> c
DateTask(date=2014-01-21)
>>> d
DateTask(date=2014-01-21)
>>> c is d
True
```

10.4.2 Insignificant parameters

If a parameter is created with `significant=False`, it is ignored as far as the Task signature is concerned. Tasks created with only insignificant parameters differing have the same signature but are not the same instance:

```
>>> class DateTask2(DateTask):
...     other = luigi.Parameter(significant=False)
...
>>> c = DateTask2(date=a, other="foo")
>>> d = DateTask2(date=b, other="bar")
>>> c
DateTask2(date=2014-01-21)
>>> d
DateTask2(date=2014-01-21)
>>> c.other
'foo'
>>> d.other
'bar'
>>> c is d
False
>>> hash(c) == hash(d)
True
```

10.4.3 Parameter types

In the examples above, the *type* of the parameter is determined by using different subclasses of *Parameter*. There are a few of them, like *DateParameter*, *DateIntervalParameter*, *IntParameter*, *FloatParameter*, etc.

Python is not a strongly typed language and you don't have to specify the types of any of your parameters. You can simply use the base class `Parameter` if you don't care.

The reason you would use a subclass like `DateParameter` is that Luigi needs to know its type for the command line interaction. That's how it knows how to convert a string provided on the command line to the corresponding type (i.e. `datetime.date` instead of a string).

10.4.4 Setting parameter value for other classes

All parameters are also exposed on a class level on the command line interface. For instance, say you have classes `TaskA` and `TaskB`:

```
class TaskA(luigi.Task):
    x = luigi.Parameter()

class TaskB(luigi.Task):
    y = luigi.Parameter()
```

You can run `TaskB` on the command line: `python script.py TaskB --y 42`. But you can also set the class value of `TaskA` by running `python script.py TaskB --y 42 --TaskA-x 43`. This sets the value of `TaskA.x` to 43 on a *class* level. It is still possible to override it inside Python if you instantiate `TaskA` (`x=44`).

All parameters can also be set from the configuration file. For instance, you can put this in the config:

```
[TaskA]
x: 45
```

Just as in the previous case, this will set the value of `TaskA.x` to 45 on the *class* level. And likewise, it is still possible to override it inside Python if you instantiate `TaskA` (`x=44`).

10.4.5 Parameter resolution order

Parameters are resolved in the following order of decreasing priority:

1. Any value passed to the constructor, or task level value set on the command line (applies on an instance level)
2. Any value set on the command line (applies on a class level)
3. Any configuration option (applies on a class level)
4. Any default value provided to the parameter (applies on a class level)

See the `Parameter` class for more information.

10.5 Running from the Command Line

Any task can be instantiated and run from the command line:

```
import luigi

class MyTask(luigi.Task):
    x = luigi.IntParameter()
    y = luigi.IntParameter(default=45)

    def run(self):
        print self.x + self.y
```

```
if __name__ == '__main__':  
    luigi.run()
```

You can run this task from the command line like this:

```
$ python my_task.py MyTask --local-scheduler --x 123 --y 456
```

You can also pass `main_task_cls=MyTask` and `local_scheduler=True` to `luigi.run()` and that way you can invoke it simply using

```
$ python my_task.py --x 123 --y 456
```

The other way to run a Luigi task is to use the builtin *luigi* script. This will be default on your path and can be run by providing a module name. The module will be imported dynamically:

```
$ luigi --module my_module MyTask --x 123 --y 456
```

10.6 Programmatic Execution

As seen above, command line integration is achieved by simply adding

```
if __name__ == '__main__':  
    luigi.run()
```

This will read the args from the command line (using `argparse`) and invoke everything.

In case you just want to run a Luigi chain from a Python script, you can do that internally without the command line integration. The code will look something like

```
task = MyTask(123, 'xyz')  
interface.setup_interface_logging()  
sch = scheduler.CentralPlannerScheduler()  
w = worker.Worker(scheduler=sch)  
w.add(task)  
w.run()
```

10.7 Using the Central Scheduler

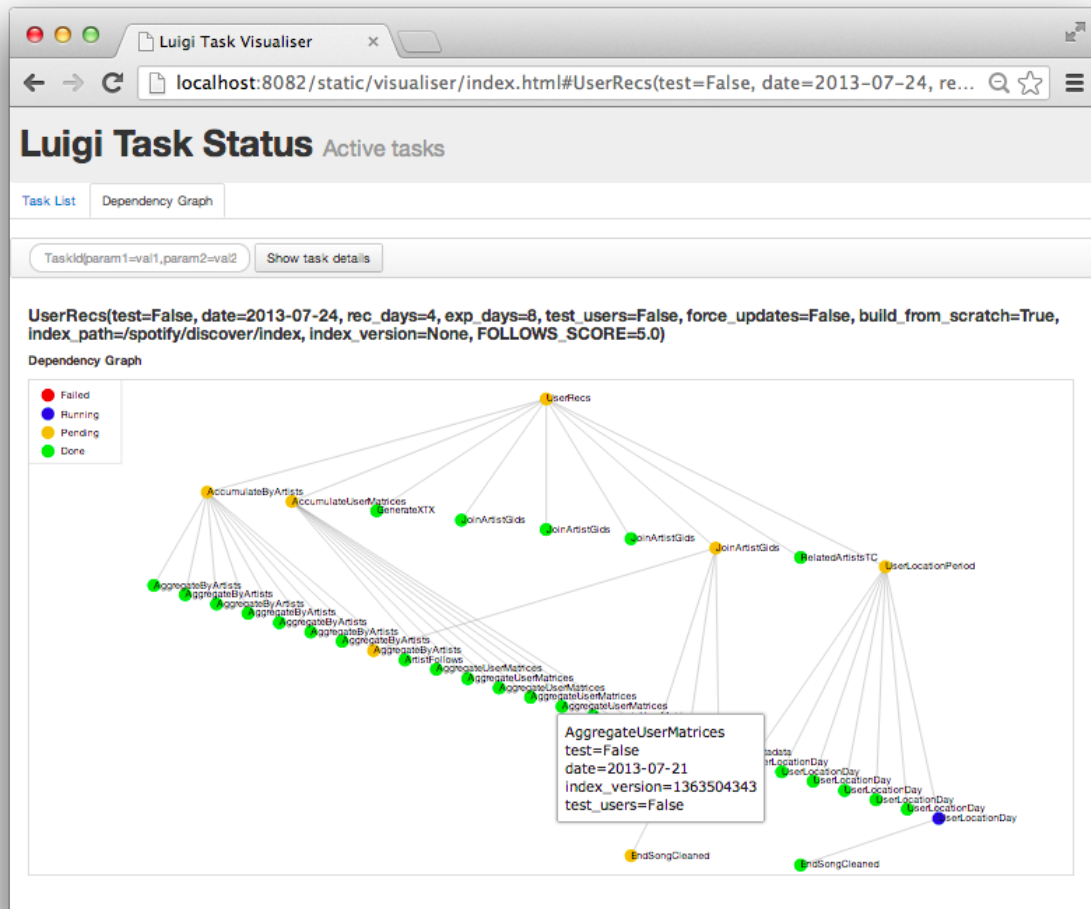
While the `--local-scheduler` flag is useful for development purposes, it's not recommended for production usage. The centralized scheduler services two purposes:

- Make sure two instances of the same task are not running simultaneously
- Provide visualization of everything that's going on.

Note that the central scheduler does not execute anything for you or help you with job parallelization. For running tasks periodically, the easiest thing to do is to trigger a Python script from cron or from a continuously running process. There is no central process that automatically triggers jobs. This model may seem limited, but we believe that it makes things far more intuitive and easy to understand.

10.7.1 The luigid server

To run the server as a daemon run:



```
luigid --background --pidfile <PATH_TO_PIDFILE> --logdir <PATH_TO_LOGDIR> --state-path <PATH_TO_STATE_PATH>
```

Note that this requires `python-daemon`. By default, the server starts on port 8082 (which can be changed with the `--port` flag) and listens on all IPs.

For a full list of configuration options and defaults, see the [scheduler configuration section](#). Note that `luigid` uses the same configuration files as the luigi client (i.e. `client.cfg` or `/etc/luigi/client.cfg` by default).

10.7.2 Enabling Task History

Task History is an experimental feature in which additional information about tasks that have been executed are recorded in a relational database for historical analysis. This information is exposed via the Central Scheduler at `/history`.

To enable the task history, specify `record_task_history = True` in the `[scheduler]` section of `client.cfg` and specify `db_connection` under `[task_history]`. The `db_connection` string is to be used to configure the [SQLAlchemy engine](#). When starting up, `luigid` will create all the necessary tables using `create_all`.

Example configuration:

```
[scheduler]
record_task_history = True
state-path = /usr/local/var/luigi-state.pickle

[task_history]
db_connection = sqlite:///usr/local/var/luigi-task-hist.db
```

The task history has the following pages:

- `/history` a reverse-cronological listing of runs from the past 24 hours. Example screenshot:

Name	Host	Last Action	Status
WordCount	None	2014-12-31 20:16:58.505362	DONE
WordCount	None	2014-12-31 20:16:56.602269	DONE
InputText	None	2014-12-31 20:16:52.233391	PENDING
WordCount	None	2014-12-31 20:16:52.210956	PENDING

- `/history/by_id/:id` detailed information about a run, including: parameter values, the host on which it ran, and timing information. Example screenshot:
- `/history/by_name/:name` a listing of all runs of a task with the given task name. Example screenshot:
- `/history/by_params/:name?data=params` a listing of all runs of a given task restricted to runs with param values matching the given data. The data is a json blob describing the parameters, e.g. `{"foo": "bar"}` looks for a task with `foo=bar`.

Info

Task Id	4
Task Name	WordCount
Host	None
More	All "WordCount" runs

Parameters

Name	Value
date_interval	2014-12-31

Actions

Status	Action Time
DONE	2014-12-31 20:16:58.505362

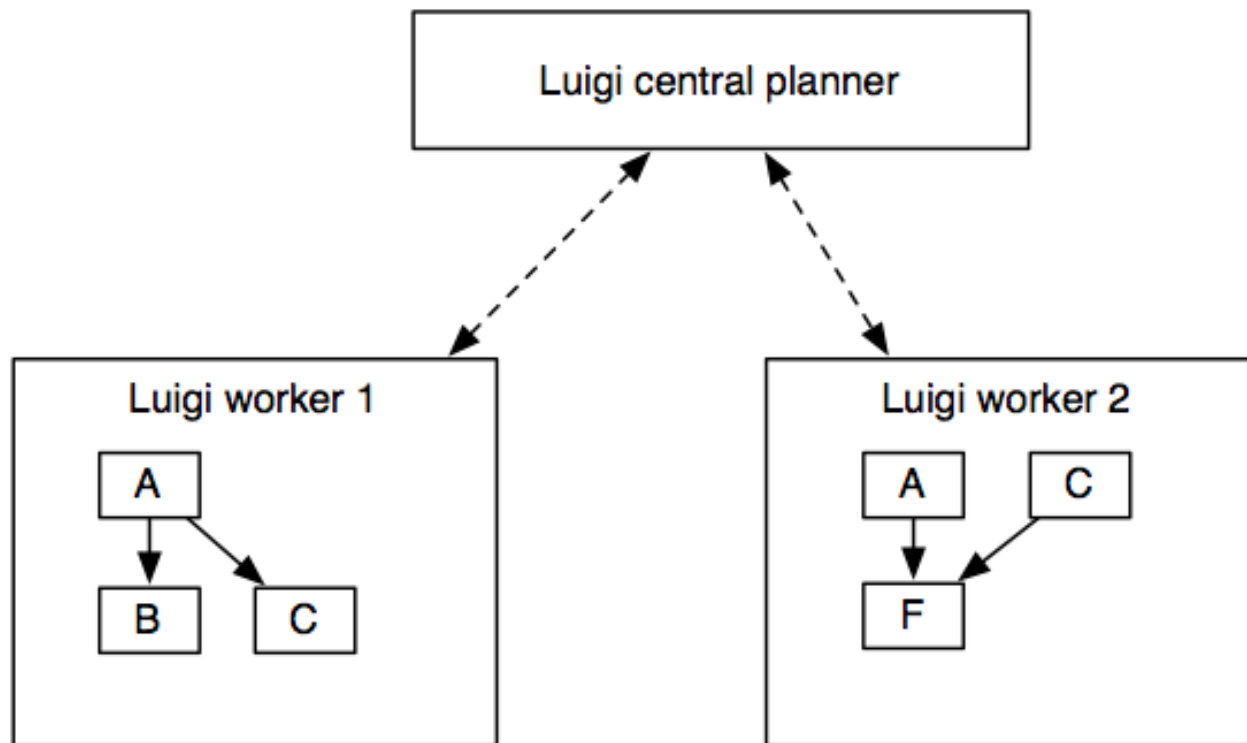
Name	Host	Last Action	Status
WordCount	None	2014-12-31 20:16:52.210956	PENDING
WordCount	None	2014-12-31 20:16:56.602269	DONE
WordCount	None	2014-12-31 20:16:58.505362	DONE

10.8 Execution Model

Luigi has a quite simple model for execution and triggering.

10.8.1 Workers and task execution

The most important aspect is that *no execution is transferred*. When you run a Luigi workflow, the worker schedules all tasks, and also executes the tasks within the process.



The benefit of this scheme is that it's super easy to debug since all execution takes place in the process. It also makes deployment a non-event. During development, you typically run the Luigi workflow from the command line, whereas when you deploy it, you can trigger it using crontab or any other scheduler.

The downside is that Luigi doesn't give you scalability for free. In practice this is not a problem until you start running thousands of tasks.

Isn't the point of Luigi to automate and schedule these workflows? To some extent. Luigi helps you *encode the dependencies* of tasks and build up chains. Furthermore, Luigi's scheduler makes sure that there's centralized view of the dependency graph and that the same job will not be executed by multiple workers simultaneously.

10.8.2 Triggering tasks

Luigi does not include its own triggering, so you have to rely on an external scheduler such as crontab to actually trigger the workflows.

In practice it's not a big hurdle because Luigi avoids all the mess typically caused by it. Scheduling a complex workflow is fairly trivial using eg. crontab.

In the future, Luigi might implement its own triggering. The dependency on crontab (or any external triggering mechanism) is a bit awkward and it would be nice to avoid.

Trigger example

For instance, if you have an external data dump that arrives every day and that your workflow depends on it, you write a workflow that depends on this data dump. Crontab can then trigger this workflow *every minute* to check if the data has arrived. If it has, it will run the full dependency graph.

```
class DataDump(luigi.ExternalTask):
    date = luigi.DateParameter()
    def output(self): return luigi.contrib.hdfs.HdfsTarget(self.date.strftime('/var/log/dump/%Y-%m-%d'))

class AggregationTask(luigi.Task):
    date = luigi.DateParameter()
    window = luigi.IntParameter()
    def requires(self): return [DataDump(self.date - datetime.timedelta(i)) for i in xrange(self.window)]
    def run(self): run_some_cool_stuff(self.input())
    def output(self): return luigi.contrib.hdfs.HdfsTarget('/aggregated-%s-%d' % (self.date, self.window))

class RunAll(luigi.Task):
    ''' Dummy task that triggers execution of a other tasks'''
    def requires(self):
        for window in [3, 7, 14]:
            for d in xrange(10): # guarantee that aggregations were run for the past 10 days
                yield AggregationTask(datetime.date.today() - datetime.timedelta(d), window)

if __name__ == '__main__':
    luigi.run(main_task_cls=RunAll)
```

You can trigger this as much as you want from crontab, and even across multiple machines, because the central scheduler will make sure at most one of each AggregationTask task is run simultaneously. Note that this might actually mean multiple tasks can be run because there are instances with different parameters, and this can gives you some form of parallelization (eg. AggregationTask(2013-01-09) might run in parallel with AggregationTask(2013-01-08)).

Of course, some Task types (eg. HadoopJobTask) can transfer execution to other places, but this is up to each Task to define.

10.9 Luigi Patterns

10.9.1 Code Reuse

One nice thing about Luigi is that it's super easy to depend on tasks defined in other repos. It's also trivial to have “forks” in the execution path, where the output of one task may become the input of many other tasks.

Currently no semantics for “intermediate” output is supported, meaning that all output will be persisted indefinitely. The upside of that is that if you try to run $X \rightarrow Y$, and Y crashes, you can resume with the previously built X . The downside is that you will have a lot of intermediate results on your file system. A useful pattern is to put these files in a special directory and have some kind of periodical garbage collection clean it up.

10.9.2 Triggering Many Tasks

A convenient pattern is to have a dummy Task at the end of several dependency chains, so you can trigger a multitude of pipelines by specifying just one task in command line, similarly to how e.g. `make` works.

```
class AllReports(luigi.Task):
    date = luigi.DateParameter(default=datetime.date.today())
    def requires(self):
        yield SomeReport(self.date)
        yield SomeOtherReport(self.date)
        yield CropReport(self.date)
        yield TPSReport(self.date)
        yield FooBarBazReport(self.date)
```

This simple task will not do anything itself, but will invoke a bunch of other tasks. Per each invocation Luigi will perform as many of the pending jobs as possible (those which have all their dependencies present).

10.9.3 Triggering recurring tasks

A common requirement is to have a daily report (or something else) produced every night. Sometimes for various reasons tasks will keep crashing or lacking their required dependencies for more than a day though, which would lead to a missing deliverable for some date. Oops.

To ensure that the above AllReports task is eventually completed for every day (value of date parameter), one could e.g. add a loop in requires method to yield dependencies on the past few days preceding self.date. Then, so long as Luigi keeps being invoked, the backlog of jobs would catch up nicely after fixing intermittent problems.

Luigi actually comes with a reusable tool for achieving this, called RangeDailyBase (resp. RangeHourlyBase). Simply putting

```
luigi --module all_reports RangeDailyBase --of AllReports --start 2015-01-01
```

in your crontab will easily keep gaps from occurring from 2015-01-01 onwards. NB - it will not always loop over everything from 2015-01-01 till current time though, but rather a maximum of 3 months ago by default - see RangeDailyBase documentation for this and more knobs for tweaking behavior. See also Monitoring below.

10.9.4 Efficiently triggering recurring tasks

RangeDailyBase, described above, is named like that because a more efficient subclass exists, RangeDaily (resp. RangeHourly), tailored for hundreds of task classes scheduled concurrently with contiguousness requirements spanning years (which would incur redundant completeness checks and scheduler overload using the naive looping approach.) Usage:

```
luigi --module all_reports RangeDaily --of AllReports --start 2015-01-01
```

It has the same knobs as RangeDailyBase, with some added requirements. Namely the task must implement an efficient bulk_complete method, or must be writing output to file system Target with date parameter value consistently represented in the file path.

10.9.5 Backfilling tasks

Also a common use case, sometimes you have tweaked existing recurring task code and you want to schedule recomputation of it over an interval of dates for that or another reason. Most conveniently it is achieved with the above described range tools, just with both start (inclusive) and stop (exclusive) parameters specified:

```
luigi --module all_reports RangeDaily --of AllReportsV2 --start 2014-10-31 --stop 2014-12-25
```

10.9.6 Monitoring task pipelines

Set error-email in [Configuration](#) to receive notifications whenever tasks crash. (This can get noisy with growing numbers of tasks and intermittent failures.)

The above mentioned range tools for recurring tasks not only implement reliable scheduling for you, but also emit events which you can use to set up delay monitoring. That way you can implement alerts for when jobs are stuck for prolonged periods lacking input data or otherwise requiring attention.

10.10 Configuration

All configuration can be done by adding a configuration files. They are looked for in:

- `/etc/luigi/client.cfg`
- `client.cfg` in your current working directory
- `LUIGI_CONFIG_PATH` environment variable

in increasing order of preference. The order only matters in case of key conflicts (see docs for [ConfigParser](#))

The config file is broken into sections, each controlling a different part of the config. Example configuration file:

```
[hadoop]
version: cdh4
streaming-jar: /usr/lib/hadoop-xyz/hadoop-streaming-xyz-123.jar

[core]
default-scheduler-host: luigi-host.mycompany.foo
error-email: foo@bar.baz
```

By default, all parameters will be overridden by matching values in the configuration file. For instance if you have a Task definition:

```
class DailyReport (luigi.contrib.hadoop.JobTask):
    date = luigi.DateParameter(default=datetime.date.today())
    # ...
```

Then you can override the default value for date by providing it in the configuration:

```
[DailyReport]
date: 2012-01-01
```

You can also use `config_path` as an argument to the `Parameter` if you want to use a specific section in the config.

10.10.1 Configurable options

Luigi comes with a lot of configurable options. Below, we describe each section and the parameters available within it.

10.10.2 [core]

These parameters control core luigi behavior, such as error e-mails and interactions between the worker and scheduler.

default-scheduler-host Hostname of the machine running the scheduler. Defaults to localhost.

default-scheduler-port Port of the remote scheduler api process. Defaults to 8082.

email-prefix Optional prefix to add to the subject line of all e-mails. For example, setting this to “[LUIGI]” would change the subject line of an e-mail from “Luigi: Framework error” to “[LUIGI] Luigi: Framework error”

email-sender User name in from field of error e-mails. Default value: luigi-client@<server_name>

email-type Type of e-mail to send. Valid values are “plain” and “html”. When set to html, tracebacks are wrapped in <pre> tags to get fixed-width font. Default value is plain.

error-email Recipient of all error e-mails. If this is not set, no error e-mails are sent when luigi crashes. If luigi is run from the command line, no e-mails will be sent unless output is redirected to a file.

hdfs-tmp-dir Base directory in which to store temporary files on hdfs. Defaults to tempfile.gettempdir()

history-filename If set, specifies a filename for Luigi to write stuff (currently just job id) to in mapreduce job’s output directory. Useful in a configuration where no history is stored in the output directory by Hadoop.

logging_conf_file Location of the logging configuration file.

max-reschedules The maximum number of times that a job can be automatically rescheduled by a worker before it will stop trying. Workers will reschedule a job if it is found to not be done when attempting to run a dependent job. This defaults to 1.

max-shown-tasks New in version 1.0.20.

The maximum number of tasks returned in a task_list api call. This will restrict the number of tasks shown in any section in the visualiser. Small values can alleviate frozen browsers when there are too many done tasks. This defaults to 100000 (one hundred thousand).

no_configure_logging If true, logging is not configured. Defaults to false.

parallel-scheduling If true, the scheduler will compute complete functions of tasks in parallel using multiprocessing. This can significantly speed up scheduling, but requires that all tasks can be pickled.

retry-external-tasks If true, incomplete external tasks (i.e. tasks where the *run()* method is NotImplemented) will be retested for completion while Luigi is running. This means that if external dependencies are satisfied after a workflow has started, any tasks dependent on that resource will be eligible for running. Note: Every time the task remains incomplete, it will count as FAILED, so normal retry logic applies (see: *disable-num-failures* and *retry-delay*). This setting works best with *worker-keep-alive*: *true*. If false, external tasks will only be evaluated when Luigi is first invoked. In this case, Luigi will not check whether external dependencies are satisfied while a workflow is in progress, so dependent tasks will remain PENDING until the workflow is reinvoked. Defaults to false for backwards compatibility.

rpc-connect-timeout Number of seconds to wait before timing out when making an API call. Defaults to 10.0

smtp_host Hostname for sending mail through smtp. Defaults to localhost.

smtp_local_hostname If specified, overrides the FQDN of localhost in the HELO/EHLO command.

smtp_login Username to log in to your smtp server, if necessary.

smtp_password Password to log in to your smtp server. Must be specified for smtp_login to have an effect.

smtp_port Port number for smtp on smtp_host. Defaults to 0.

smtp_ssl If true, connects to smtp through SSL. Defaults to false.

smtp_timeout Optionally sets the number of seconds after which smtp attempts should time out.

tmp-dir DEPRECATED - use hdfs-tmp-dir instead

worker-count-uniques If true, workers will only count unique pending jobs when deciding whether to stay alive. So if a worker can't get a job to run and other workers are waiting on all of its pending jobs, the worker will die. worker-keep-alive must be true for this to have any effect. Defaults to false.

worker-keep-alive If true, workers will stay alive when they run out of jobs to run, as long as they have some pending job waiting to be run. Defaults to false.

worker-ping-interval Number of seconds to wait between pinging scheduler to let it know that the worker is still alive. Defaults to 1.0.

worker-task-limit New in version 1.0.25.

Maximum number of tasks to schedule per invocation. Upon exceeding it, the worker will issue a warning and proceed with the workflow obtained thus far. Prevents incidents due to spamming of the scheduler, usually accidental. Default: no limit.

worker-timeout New in version 1.0.20.

Number of seconds after which to kill a task which has been running for too long. This provides a default value for all tasks, which can be overridden by setting the worker-timeout property in any task. This only works when using multiple workers, as the timeout is implemented by killing worker subprocesses. Default value is 0, meaning no timeout.

worker-wait-interval Number of seconds for the worker to wait before asking the scheduler for another job after the scheduler has said that it does not have any available jobs.

10.10.3 [elasticsearch]

These parameters control use of elasticsearch

marker-index Defaults to "update_log".

marker-doc-type Defaults to "entry".

10.10.4 [email]

These parameters control sending error e-mails through Amazon SES.

AWS_ACCESS_KEY Your AWS access key

AWS_SECRET_KEY Your AWS secret key

region Your AWS region. Defaults to us-east-1.

type If set to "ses", error e-mails will be send through Amazon SES. Otherwise, e-mails are sent via smtp.

10.10.5 [hadoop]

Parameters controlling basic hadoop tasks

command Name of command for running hadoop from the command line. Defaults to "hadoop"

python-executable Name of command for running python from the command line. Defaults to "python"

scheduler Type of scheduler to use when scheduling hadoop jobs. Can be "fair" or "capacity". Defaults to "fair".

streaming-jar Path to your streaming jar. Must be specified to run streaming jobs.

version Version of hadoop used in your cluster. Can be "cdh3", "cdh4", or "apache1". Defaults to "cdh4".

10.10.6 [hdfs]

Parameters controlling the use of snakebite to speed up hdfs queries.

client Client to use for most hadoop commands. Options are “snakebite”, “snakebite_with_hadoopcli_fallback”, and “hadoopcli”. Snakebite is much faster, so use of it is encouraged. Using snakebite requires it to be installed separately on the machine. Defaults to “hadoopcli”.

client_version Optionally specifies hadoop client version for snakebite.

effective_user Optionally specifies the effective user for snakebite.

namenode_host The hostname of the namenode. Needed for snakebite if snakebite_autoconfig is not set.

namenode_port The port used by snakebite on the namenode. Needed for snakebite if snakebite_autoconfig is not set.

snakebite_autoconfig If true, attempts to automatically detect the host and port of the namenode for snakebite queries. Defaults to false.

10.10.7 [hive]

Parameters controlling hive tasks

command Name of the command used to run hive on the command line. Defaults to “hive”.

hiverc-location Optional path to hive rc file.

metastore_host Hostname for metastore.

metastore_port Port for hive to connect to metastore host.

release If set to “apache”, uses a hive client that better handles apache hive output. All other values use the standard client Defaults to “cdh4”.

10.10.8 [mysql]

Parameters controlling use of MySQL targets

marker-table Table in which to store status of table updates. This table will be created if it doesn’t already exist. Defaults to “table_updates”.

10.10.9 [postgres]

Parameters controlling the use of Postgres targets

local-tmp-dir Directory in which to temporarily store data before writing to postgres. Uses system default if not specified.

marker-table Table in which to store status of table updates. This table will be created if it doesn’t already exist. Defaults to “table_updates”.

10.10.10 [redshift]

Parameters controlling the use of Redshift targets

marker-table Table in which to store status of table updates. This table will be created if it doesn’t already exist. Defaults to “table_updates”.

10.10.11 [resources]

This section can contain arbitrary keys. Each of these specifies the amount of a global resource that the scheduler can allow workers to use. The scheduler will prevent running jobs with resources specified from exceeding the counts in this section. Unspecified resources are assumed to have limit 1. Example resources section for a configuration with 2 hive resources and 1 mysql resource:

```
[resources]
hive: 2
mysql: 1
```

Note that it was not necessary to specify the 1 for mysql here, but it is good practice to do so when you have a fixed set of resources.

10.10.12 [scalding]

Parameters controlling running of scalding jobs

scala-home Home directory for scala on your machine. Defaults to either SCALA_HOME or /usr/share/scala if SCALA_HOME is unset.

scalding-home Home directory for scalding on your machine. Defaults to either SCALDING_HOME or /usr/share/scalding if SCALDING_HOME is unset.

scalding-provided Provided directory for scalding on your machine. Defaults to either SCALDING_HOME/provided or /usr/share/scalding/provided

scalding-libjars Libjars directory for scalding on your machine. Defaults to either SCALDING_HOME/libjars or /usr/share/scalding/libjars

10.10.13 [scheduler]

Parameters controlling scheduler behavior

disable-hard-timeout Hard time limit after which tasks will be disabled by the server if they fail again, in seconds. It will disable the task if it fails **again** after this amount of time. E.g. if this was set to 600 (i.e. 10 minutes), and the task first failed at 10:00am, the task would be disabled if it failed again any time after 10:10am. Note: This setting does not consider the values of the *disable-num-failures* or *disable-window-seconds* settings.

disable-num-failures Number of times a task can fail within disable-window-seconds before the scheduler will automatically disable it. If not set, the scheduler will not automatically disable jobs.

disable-persist-seconds Number of seconds for which an automatic scheduler disable lasts. Defaults to 86400 (1 day).

disable-window-seconds Number of seconds during which disable-num-failures failures must occur in order for an automatic disable by the scheduler. The scheduler forgets about disables that have occurred longer ago than this amount of time. Defaults to 3600 (1 hour).

record_task_history If true, stores task history in a database. Defaults to false.

remove-delay Number of seconds to wait before removing a task that has no stakeholders. Defaults to 600 (10 minutes).

retry-delay Number of seconds to wait after a task failure to mark it pending again. Defaults to 900 (15 minutes).

state-path Path in which to store the luigi scheduler's state. When the scheduler is shut down, its state is stored in this path. The scheduler must be shut down cleanly for this to work, usually with a kill command. If the kill command includes the -9 flag, the scheduler will not be able to save its state. When the scheduler is started, it

will load the state from this path if it exists. This will restore all scheduled jobs and other state from when the scheduler last shut down.

Sometimes this path must be deleted when restarting the scheduler after upgrading luigi, as old state files can become incompatible with the new scheduler. When this happens, all workers should be restarted after the scheduler both to become compatible with the updated code and to reschedule the jobs that the scheduler has now forgotten about.

This defaults to `/var/lib/luigi-server/state.pickle`

worker-disconnect-delay Number of seconds to wait after a worker has stopped pinging the scheduler before removing it and marking all of its running tasks as failed. Defaults to 60.

10.10.14 [spark]

Parameters controlling the default execution of *SparkSubmitTask* and *PySparkTask*:

Deprecated since version 1.1.1: *SparkJob*, *Spark1xJob* and *PySpark1xJob* are deprecated. Please use *SparkSubmitTask* or *PySparkTask*.

spark-submit Command to run in order to submit spark jobs. Default: `spark-submit`

master Master url to use for spark-submit. Example: `local[*]`, `spark://masterhost:7077`. Default: Spark default (Prior to 1.1.1: `yarn-client`)

deploy-mode Whether to launch the driver programs locally (“client”) or on one of the worker machines inside the cluster (“cluster”). Default: Spark default

jars Comma-separated list of local jars to include on the driver and executor classpaths. Default: Spark default

py-files Comma-separated list of `.zip`, `.egg`, or `.py` files to place on the PYTHONPATH for Python apps. Default: Spark default

files Comma-separated list of files to be placed in the working directory of each executor. Default: Spark default

conf: Arbitrary Spark configuration property in the form `Prop=Value1Prop2=Value2`. Default: Spark default

properties-file Path to a file from which to load extra properties. Default: Spark default

driver-memory Memory for driver (e.g. `1000M`, `2G`). Default: Spark default

driver-java-options Extra Java options to pass to the driver. Default: Spark default

driver-library-path Extra library path entries to pass to the driver. Default: Spark default

driver-class-path Extra class path entries to pass to the driver. Default: Spark default

executor-memory Memory per executor (e.g. `1000M`, `2G`). Default: Spark default

Configuration for Spark submit jobs on Spark standalone with cluster deploy mode only:

driver-cores Cores for driver. Default: Spark default

supervise If given, restarts the driver on failure. Default: Spark default

Configuration for Spark submit jobs on Spark standalone and Mesos only:

total-executor-cores Total cores for all executors. Default: Spark default

Configuration for Spark submit jobs on YARN only:

executor-cores Number of cores per executor. Default: Spark default

queue The YARN queue to submit to. Default: Spark default

num-executors Number of executors to launch. Default: Spark default

archives Comma separated list of archives to be extracted into the working directory of each executor. Default: Spark default

hadoop-conf-dir Location of the hadoop conf dir. Sets HADOOP_CONF_DIR environment variable when running spark. Example: /etc/hadoop/conf

Extra configuration for PySparkTask jobs:

py-packages Comma-separated list of local packages (in your python path) to be distributed to the cluster.

Parameters controlling the execution of SparkJob jobs (deprecated):

spark-jar Location of the spark jar. Sets SPARK_JAR environment variable when running spark. Example: /usr/share/spark/jars/spark-assembly-0.8.1-incubating-hadoop2.2.0.jar

spark-class Location of script to invoke. Example: /usr/share/spark/spark-class

10.10.15 [task_history]

Parameters controlling storage of task history in a database

db_connection Connection string for connecting to the task history db using sqlalchemy.

10.11 More Info

Luigi is the successor to a couple of attempts that we weren't fully happy with. We learned a lot from our mistakes and some design decisions include:

- Straightforward command line integration.
- As little boilerplate as possible.
- Focus on job scheduling and dependency resolution, not a particular platform. In particular this means no limitation to Hadoop. Though Hadoop/HDFS support is built-in and is easy to use, this is just one of many types of things you can run.
- A file system abstraction where code doesn't have to care about where files are located.
- Atomic file system operations through this abstraction. If a task crashes it won't lead to a broken state.
- The dependencies are decentralized. No big config file in XML. Each task just specifies which inputs it needs and cross-module dependencies are trivial.
- A web server that renders the dependency graph and does locking etc for free.
- Trivial to extend with new file systems, file formats and job types. You can easily write jobs that inserts a Tokyo Cabinet into Cassandra. Adding support for new systems is generally not very hard. (Feel free to send us a patch when you're done!)
- Date algebra included.
- Lots of unit tests of the most basic stuff

It wouldn't be fair not to mention some limitations with the current design:

- Its focus is on batch processing so it's probably less useful for near real-time pipelines or continuously running processes.
- The assumption is that a each task is a sizable chunk of work. While you can probably schedule a few thousand jobs, it's not meant to scale beyond tens of thousands.

- Luigi does not support distribution of execution. When you have workers running thousands of jobs daily, this starts to matter, because the worker nodes get overloaded. There are some ways to mitigate this (trigger from many nodes, use resources), but none of them is ideal
- Luigi does not come with built-in triggering, and you still need to rely on something like crontab to trigger workflows periodically.

Also it should be mentioned that Luigi is named after the world's second most famous plumber.

10.11.1 Want to Contribute?

Awesome! Let us know if you have any ideas. Feel free to contact x@y.com where x = luigi and y = spotify.

10.11.2 Running Unit Tests

You can see in `.travis.yml` how Travis CI runs the tests. Essentially, what you do is first `pip install tox`, then you can run any of these examples and change them to your needs.

```
# Run all nonhdfs tests
export TOX_ENV=nonhdfs; export PYTHONPATH=''; tox -e $TOX_ENV test

# Run specific nonhdfs tests
export TOX_ENV=nonhdfs; export PYTHONPATH=''; tox -e $TOX_ENV test/test_ssh.py

# Run specific hdp tests with hdp hadoop distrubtion
export TOX_ENV=hdp; export PYTHONPATH=''; JAVA_HOME=/usr/lib/jvm/java-1.7.0-openjdk-amd64 tox -e $TOX_ENV test/test_ssh.py
```

10.11.3 Future Ideas

- S3/EC2 - We have some old ugly code based on Boto that could be integrated in a day or two.
- Built in support for Pig/Hive.
- Better visualization tool - the layout gets pretty messy as the number of tasks grows.
- Integration with existing Hadoop frameworks like mrjob would be cool and probably pretty easy.
- Better support (without much boilerplate) for unittesting specific Tasks

API Reference

<i>luigi</i>	Package containing core luigi functionality.
<i>luigi.contrib</i>	Package containing optional and-on functionality.
<i>luigi.tools</i>	Sort of a standard library for doing stuff with Tasks at a somewhat abstract level.

11.1 luigi package

11.1.1 Subpackages

luigi.contrib package

Subpackages

luigi.contrib.hdfs package

Submodules

luigi.contrib.hdfs.abstract_client module Module containing abstract class about hdfs clients.

class `luigi.contrib.hdfs.abstract_client.HdfsFileSystem`

Bases: `luigi.target.FileSystem`

This client uses Apache 2.x syntax for file system commands, which also matched CDH4.

rename (*path*, *dest*)

Rename or move a file

rename_dont_move (*path*, *dest*)

Override this method with an implementation that uses `rename2`, which is a rename operation that never moves.

For instance, `rename2 a b` never moves `a` into `b` folder.

Currently, the hadoop cli does not support this operation.

We keep the interface simple by just aliasing this to normal `rename` and let individual implementations redefine the method.

`rename2` - <https://github.com/apache/hadoop/blob/ae91b13/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/protocol/ClientProtocol.java> (lines 483-523)

remove (*path*, *recursive=True*, *skip_trash=False*)
chmod (*path*, *permissions*, *recursive=False*)
chown (*path*, *owner*, *group*, *recursive=False*)
count (*path*)
 Count contents in a directory
copy (*path*, *destination*)
put (*local_path*, *destination*)
get (*path*, *local_destination*)
mkdir (*path*, *parents=True*, *raise_if_exists=False*)
listdir (*path*, *ignore_directories=False*, *ignore_files=False*, *include_size=False*, *include_type=False*,
 include_time=False, *recursive=False*)
touchz (*path*)

luigi.contrib.hdfs.clients module The implementations of the hdfs clients. The hadoop cli client and the snakebite client.

`luigi.contrib.hdfs.clients.get_autoconfig_client` (*show_warnings=True*)
 Creates the client as specified in the *client.cfg* configuration.

luigi.contrib.hdfs.config module You can configure what client by setting the “client” config under the “hdfs” section in the configuration, or using the `--hdfs-client` command line option. “hadoopcli” is the slowest, but should work out of the box. “snakebite” is the fastest, but requires Snakebite to be installed.

class `luigi.contrib.hdfs.config.hdfs` (**args*, ***kwargs*)
 Bases: `luigi.task.Config`

 Constructor to resolve values for all Parameters.

 For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

 can be instantiated as `MyTask(count=10)`.

client_version = <luigi.parameter.IntParameter object>

effective_user = <luigi.parameter.Parameter object>

snakebite_autoconfig = <luigi.parameter.BoolParameter object>

namenode_host = <luigi.parameter.Parameter object>

namenode_port = <luigi.parameter.IntParameter object>

client = <luigi.parameter.Parameter object>

tmp_dir = <luigi.parameter.Parameter object>

task_namespace = None

class `luigi.contrib.hdfs.config.hadoopcli` (**args*, ***kwargs*)
 Bases: `luigi.task.Config`

 Constructor to resolve values for all Parameters.

 For example, the Task:


```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

command = <luigi.parameter.Parameter object>

version = <luigi.parameter.Parameter object>

task_namespace = None

```
luigi.contrib.hdfs.config.load_hadoop_cmd()
```

```
luigi.contrib.hdfs.config.get_configured_hadoop_version()
```

CDH4 (hadoop 2+) has a slightly different syntax for interacting with hdfs via the command line.

The default version is CDH4, but one can override this setting with “cdh3” or “apache1” in the hadoop section of the config in order to use the old syntax.

```
luigi.contrib.hdfs.config.get_configured_hdfs_client(show_warnings=True)
```

This is a helper that fetches the configuration value for ‘client’ in the [hdfs] section. It will return the client that retains backwards compatibility when ‘client’ isn’t configured.

```
luigi.contrib.hdfs.config.tmppath(path=None, include_unix_username=True)
```

@param path: target path for which it is needed to generate temporary location @type path: str @type include_unix_username: bool @rtype: str

Note that include_unix_username might work on windows too.

luigi.contrib.hdfs.error module The implementations of the hdfs clients. The hadoop cli client and the snakebite client.

```
exception luigi.contrib.hdfs.error.HDFSClientError(command, returncode, stdout, stderr)
```

Bases: `exceptions.Exception`

luigi.contrib.hdfs.format module

```
class luigi.contrib.hdfs.format.HdfsReadPipe(path)
```

Bases: `luigi.format.InputPipeProcessWrapper`

```
class luigi.contrib.hdfs.format.HdfsAtomicWritePipe(path)
```

Bases: `luigi.format.OutputPipeProcessWrapper`

File like object for writing to HDFS

The referenced file is first written to a temporary location and then renamed to final location on `close()`. If `close()` isn’t called the temporary file will be cleaned up when this object is garbage collected

TODO: if this is buggy, change it so it first writes to a local temporary file and then uploads it on completion

abort()

close()

```
class luigi.contrib.hdfs.format.HdfsAtomicWriteDirPipe(path, data_extension='')
```

Bases: `luigi.format.OutputPipeProcessWrapper`

Writes a data<data_extension> file to a directory at <path>.

abort()

close()

```
class luigi.contrib.hdfs.format.PlainFormat
```

Bases: `luigi.format.Format`

```
    input = 'bytes'
    output = 'hdfs'
    hdfs_writer(path)
    hdfs_reader(path)
    pipe_reader(path)
    pipe_writer(output_pipe)

class luigi.contrib.hdfs.format.PlainDirFormat
    Bases: luigi.format.Format
    input = 'bytes'
    output = 'hdfs'
    hdfs_writer(path)
    hdfs_reader(path)
    pipe_reader(path)
    pipe_writer(path)

class luigi.contrib.hdfs.format.CompatibleHdfsFormat(writer, reader, input=None)
    Bases: luigi.format.Format
    output = 'hdfs'
    pipe_writer(output)
    pipe_reader(input)
    hdfs_writer(output)
    hdfs_reader(input)
```

luigi.contrib.hdfs.hadoopcli_clients module The implementations of the hdfs clients. The hadoop cli client and the snakebite client.

```
luigi.contrib.hdfs.hadoopcli_clients.create_hadoopcli_client()
    Given that we want one of the hadoop cli clients (unlike snakebite), this one will return the right one.
```

```
class luigi.contrib.hdfs.hadoopcli_clients.HdfsClient
    Bases: luigi.contrib.hdfs.abstract_client.HdfsFileSystem

    This client uses Apache 2.x syntax for file system commands, which also matched CDH4.

    recursive_listdir_cmd = ['-ls', '-R']
    static call_check(command)
    exists(path)
        Use hadoop fs -stat to check file existence.
    rename(path, dest)
    remove(path, recursive=True, skip_trash=False)
    chmod(path, permissions, recursive=False)
    chown(path, owner, group, recursive=False)
    count(path)
```

```

copy (path, destination)
put (local_path, destination)
get (path, local_destination)
getmerge (path, local_destination, new_line=False)
mkdir (path, parents=True, raise_if_exists=False)
listdir (path, ignore_directories=False, ignore_files=False, include_size=False, include_type=False,
          include_time=False, recursive=False)
touchz (path)

```

```

class luigi.contrib.hdfs.hadoopcli_clients.HdfsClientCdh3
    Bases: luigi.contrib.hdfs.hadoopcli_clients.HdfsClient

```

This client uses CDH3 syntax for file system commands.

```

mkdir (path)
    No -p switch, so this will fail creating ancestors.

remove (path, recursive=True, skip_trash=False)

```

```

class luigi.contrib.hdfs.hadoopcli_clients.HdfsClientApache1
    Bases: luigi.contrib.hdfs.hadoopcli_clients.HdfsClientCdh3

```

This client uses Apache 1.x syntax for file system commands, which are similar to CDH3 except for the file existence check.

```

recursive_listdir_cmd = ['lsr']

exists (path)

```

luigi.contrib.hdfs.snakebite_client module A luigi file system client that wraps around snakebite

Originally written by Alan Brenner <alan@magnetic.com> github.com/alanbbr

```

class luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient
    Bases: luigi.contrib.hdfs.abstract_client.HdfsFileSystem

```

A hdfs client using snakebite. Since Snakebite has a python API, it'll be about 100 times faster than the hadoop cli client, which does shell out to a java program on each file system operation.

```

static list_path (path)

get_bite ()
    If Luigi has forked, we have a different PID, and need to reconnect.

exists (path)
    Use snakebite.test to check file existence.

```

Parameters *path* (*string*) – path to test

Returns boolean, True if path exists in HDFS

```

rename (path, dest)
    Use snakebite.rename, if available.

```

Parameters

- **path** (*either a string or sequence of strings*) – source file(s)
- **dest** (*string*) – destination file (single input) or directory (multiple)

Returns list of renamed items

rename_dont_move (*path, dest*)

Use snakebite.rename_dont_move, if available.

Parameters

- **path** (*string*) – source path (single input)
- **dest** (*string*) – destination path

Returns True if succeeded

Raises snakebite.errors.FileAlreadyExistsException

remove (*path, recursive=True, skip_trash=False*)

Use snakebite.delete, if available.

Parameters

- **path** (*either a string or a sequence of strings*) – delete-able file(s) or directory(ies)
- **recursive** (*boolean, default is True*) – delete directories trees like *nix: rm -r
- **skip_trash** (*boolean, default is False (use trash)*) – do or don't move deleted items into the trash first

Returns list of deleted items

chmod (*path, permissions, recursive=False*)

Use snakebite.chmod, if available.

Parameters

- **path** (*either a string or sequence of strings*) – update-able file(s)
- **permissions** (*octal*) – *nix style permission number
- **recursive** (*boolean, default is False*) – change just listed entry(ies) or all in directories

Returns list of all changed items

chown (*path, owner, group, recursive=False*)

Use snakebite.chown/chgrp, if available.

One of owner or group must be set. Just setting group calls chgrp.

Parameters

- **path** (*either a string or sequence of strings*) – update-able file(s)
- **owner** (*string*) – new owner, can be blank
- **group** (*string*) – new group, can be blank
- **recursive** (*boolean, default is False*) – change just listed entry(ies) or all in directories

Returns list of all changed items

count (*path*)

Use snakebite.count, if available.

Parameters **path** (*string*) – directory to count the contents of

Returns dictionary with content_size, dir_count and file_count keys

copy (*path, destination*)

Raise a NotImplementedError exception.

put (*local_path, destination*)

Raise a NotImplementedError exception.

get (*path*, *local_destination*)

Use `snakebite.copyToLocal`, if available.

Parameters

- **path** (*string*) – HDFS file
- **local_destination** (*string*) – path on the system running Luigi

mkdir (*path*, *parents=True*, *mode=493*, *raise_if_exists=False*)

Use `snakebite.mkdir`, if available.

Snakebite's `mkdir` method allows control over full path creation, so by default, tell it to build a full path to work like `hadoop fs -mkdir`.

Parameters

- **path** (*string*) – HDFS path to create
- **parents** (*boolean*, *default is True*) – create any missing parent directories
- **mode** (*octal*, *default 0755*) – *nix style owner/group/other permissions

listdir (*path*, *ignore_directories=False*, *ignore_files=False*, *include_size=False*, *include_type=False*, *include_time=False*, *recursive=False*)

Use `snakebite.ls` to get the list of items in a directory.

Parameters

- **path** (*string*) – the directory to list
- **ignore_directories** (*boolean*, *default is False*) – if True, do not yield directory entries
- **ignore_files** (*boolean*, *default is False*) – if True, do not yield file entries
- **include_size** (*boolean*, *default is False (do not include)*) – include the size in bytes of the current item
- **include_type** (*boolean*, *default is False (do not include)*) – include the type (d or f) of the current item
- **include_time** (*boolean*, *default is False (do not include)*) – include the last modification time of the current item
- **recursive** (*boolean*, *default is False (do not recurse)*) – list subdirectory contents

Returns yield with a string, or if any of the `include_*` settings are true, a tuple starting with the path, and `include_*` items in order

touchz (*path*)

Raise a `NotImplementedError` exception.

luigi.contrib.hdfs.target module Provides access to HDFS using the `HdfsTarget`, a subclass of `Target`.

class luigi.contrib.hdfs.target.HdfsTarget (*path=None*, *format=None*, *is_tmp=False*, *fs=None*)

Bases: `luigi.target.FileSystemTarget`

fs

glob_exists (*expected_files*)

open (*mode='r'*)

remove (*skip_trash=False*)

rename (*path*, *raise_if_exists=False*)

Rename does not change self.path, so be careful with assumptions.

Not recommended for directories. Use move_dir. [spotify/luigi#522](#)

move (*path*, *raise_if_exists=False*)

Move does not change self.path, so be careful with assumptions.

Not recommended for directories. Use move_dir. [spotify/luigi#522](#)

move_dir (*path*)

Rename a directory.

The implementation uses *rename_dont_move*, which on some clients is just a normal *mv* operation, which can cause nested directories.

One could argue that the implementation should use the mkdir+raise_if_exists approach, but we at Spotify have had more trouble with that over just using plain mv. See [spotify/luigi#557](#)

is_writable ()

Currently only works with hadoopcli

Module contents Provides access to HDFS using the `HdfsTarget`, a subclass of `Target`. You can configure what client by setting the “client” config under the “hdfs” section in the configuration, or using the `--hdfs-client` command line option. “hadoopcli” is the slowest, but should work out of the box. “snakebite” is the fastest, but requires Snakebite to be installed.

Currently (4th May) the `luigi.contrib.hdfs` module is under reorganization. We recommend importing the reexports from `luigi.contrib.hdfs` instead of the sub-modules, as we’re not yet sure how the final structure of the sub-modules will be. Eventually this module will be empty and you’ll have to import directly from the sub-modules like `luigi.contrib.hdfs.config`.

Submodules

luigi.contrib.bigquery module

class `luigi.contrib.bigquery.CreateDisposition`

Bases: `object`

CREATE_IF_NEEDED = ‘CREATE_IF_NEEDED’

CREATE_NEVER = ‘CREATE_NEVER’

class `luigi.contrib.bigquery.WriteDisposition`

Bases: `object`

WRITE_TRUNCATE = ‘WRITE_TRUNCATE’

WRITE_APPEND = ‘WRITE_APPEND’

WRITE_EMPTY = ‘WRITE_EMPTY’

class `luigi.contrib.bigquery.QueryMode`

Bases: `object`

INTERACTIVE = ‘INTERACTIVE’

BATCH = ‘BATCH’

class `luigi.contrib.bigquery.SourceFormat`

Bases: `object`

CSV = ‘CSV’

```

DATASTORE_BACKUP = 'DATASTORE_BACKUP'
NEWLINE_DELIMITED_JSON = 'NEWLINE_DELIMITED_JSON'

class luigi.contrib.bigquery.BQDataset (project_id, dataset_id)
    Bases: tuple

    dataset_id
        Alias for field number 1

    project_id
        Alias for field number 0

class luigi.contrib.bigquery.BQTable
    Bases: luigi.contrib.bigquery.BQTable

    dataset

class luigi.contrib.bigquery.BigqueryClient (oauth_credentials=None, descriptor='',
                                             http_=None)
    Bases: object

    A client for Google BigQuery.

    For details of how authentication and the descriptor work, see the documentation for the GCS client. The
    descriptor URL for BigQuery is https://www.googleapis.com/discovery/v1/apis/bigquery/v2/rest

    dataset_exists (dataset)
        Returns whether the given dataset exists.

        Parameters dataset (BQDataset) –

    table_exists (table)
        Returns whether the given table exists.

        Parameters table (BQTable) –

    make_dataset (dataset, raise_if_exists=False, body={})
        Creates a new dataset with the default permissions.

        Parameters

        • dataset (BQDataset) –

        • raise_if_exists – whether to raise an exception if the dataset already exists.

        Raises luigi.target.FileAlreadyExists if raise_if_exists=True and the dataset exists

    delete_dataset (dataset, delete_nonempty=True)
        Deletes a dataset (and optionally any tables in it), if it exists.

        Parameters

        • dataset (BQDataset) –

        • delete_nonempty – if true, will delete any tables before deleting the dataset

    delete_table (table)
        Deletes a table, if it exists.

        Parameters table (BQTable) –

    list_datasets (project_id)
        Returns the list of datasets in a given project.

        Parameters project_id (str) –

```

list_tables (*dataset*)

Returns the list of tables in a given dataset.

Parameters *dataset* ([BQDataset](#)) –

run_job (*project_id*, *body*, *dataset=None*)

Runs a bigquery “job”. See the documentation for the format of body.

Note: You probably don’t need to use this directly. Use the tasks defined below.

Parameters *dataset* ([BQDataset](#)) –

copy (*source_table*, *dest_table*, *create_disposition='CREATE_IF_NEEDED'*,
write_disposition='WRITE_TRUNCATE')
Copies (or appends) a table to another table.

Parameters

- **source_table** ([BQTable](#)) –
- **dest_table** ([BQTable](#)) –
- **create_disposition** ([CreateDisposition](#)) – whether to create the table if needed
- **write_disposition** ([WriteDisposition](#)) – whether to append/truncate/fail if the table exists

class `luigi.contrib.bigquery.BigqueryTarget` (*project_id*, *dataset_id*, *table_id*, *client=None*)

Bases: [luigi.target.Target](#)

classmethod `from_bqtable` (*table*, *client=None*)

A constructor that takes a [BQTable](#).

Parameters *table* ([BQTable](#)) –

exists ()

class `luigi.contrib.bigquery.BigqueryLoadTask` (**args*, ***kwargs*)

Bases: [luigi.task.Task](#)

Load data into bigquery from GCS.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

source_format

The source format to use (see [SourceFormat](#)).

write_disposition

What to do if the table already exists. By default this will fail the job.

See [WriteDisposition](#)

schema

Schema in the format defined at <https://cloud.google.com/bigquery/docs/reference/v2/jobs#configuration.load.schema>.

If the value is falsy, it is omitted and inferred by bigquery, which only works for CSV inputs.

max_bad_records

source_uris
Source data which should be in GCS.

run()

task_namespace = None

class `luigi.contrib.bigquery.BigqueryRunQueryTask(*args, **kwargs)`

Bases: `luigi.task.Task`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

write_disposition

What to do if the table already exists. By default this will fail the job.

See `WriteDisposition`

create_disposition

Whether to create the table or not. See `CreateDisposition`

query

The query, in text form.

query_mode

The query mode. See `QueryMode`.

run()

task_namespace = None

luigi.contrib.esindex module Support for Elasticsearch (1.0.0 or newer).

Provides an `ElasticsearchTarget` and a `CopyToIndex` template task.

Modeled after `luigi.contrib.rdbms.CopyToTable`.

A minimal example (assuming elasticsearch is running on localhost:9200):

```
class ExampleIndex(CopyToIndex):
    index = 'example'

    def docs(self):
        return [{'_id': 1, 'title': 'An example document.'}]

if __name__ == '__main__':
    task = ExampleIndex()
    luigi.build([task], local_scheduler=True)
```

All options:

```
class ExampleIndex(CopyToIndex):
    host = 'localhost'
    port = 9200
    index = 'example'
    doc_type = 'default'
    purge_existing_index = True
```

```
marker_index_hist_size = 1

def docs(self):
    return [{'_id': 1, 'title': 'An example document.'}]

if __name__ == '__main__':
    task = ExampleIndex()
    luigi.build([task], local_scheduler=True)
```

Host, port, index, doc_type parameters are standard elasticsearch.

purge_existing_index will delete the index, whenever an update is required. This is useful, when one deals with “dumps” that represent the whole data, not just updates.

marker_index_hist_size sets the maximum number of entries in the ‘marker’ index:

- 0 (default) keeps all updates,
- 1 to only remember the most recent update to the index.

This can be useful, if an index needs to be recreated, even though the corresponding indexing task has been run sometime in the past - but a later indexing task might have altered the index in the meantime.

There are two luigi *client.cfg* configuration options:

```
[elasticsearch]

marker-index = update_log
marker-doc-type = entry
```

```
class luigi.contrib.esindex.ElasticsearchTarget(host, port, index, doc_type, up-
                                             date_id, marker_index_hist_size=0,
                                             http_auth=None, timeout=10, ex-
                                             tra_elasticsearch_args={})
```

Bases: *luigi.target.Target*

Target for a resource in Elasticsearch.

Parameters

- **host** (*str*) – Elasticsearch server host
- **port** (*int*) – Elasticsearch server port
- **index** (*str*) – index name
- **doc_type** (*str*) – doctype name
- **update_id** (*str*) – an identifier for this data set
- **marker_index_hist_size** (*int*) – list of changes to the index to remember
- **timeout** (*int*) – Elasticsearch connection timeout
- **extra_elasticsearch_args** – extra args for Elasticsearch

marker_index = ‘update_log’

marker_doc_type = ‘entry’

marker_index_document_id()

Generate an id for the indicator document.

touch()

Mark this update as complete.

The document id would be sufficient but, for documentation, we index the parameters *update_id*, *target_index*, *target_doc_type* and *date* as well.

exists()

Test, if this task has been run.

create_marker_index()

Create the index that will keep track of the tasks if necessary.

ensure_hist_size()

Shrink the history of updates for a *index/doc_type* combination down to *self.marker_index_hist_size*.

class `luigi.contrib.esindex.CopyToIndex(*args, **kwargs)`

Bases: `luigi.task.Task`

Template task for inserting a data set into Elasticsearch.

Usage:

1. Subclass and override the required *index* attribute.
2. Implement a custom *docs* method, that returns an iterable over the documents. A document can be a JSON string, e.g. from a newline-delimited JSON (ldj) file (default implementation) or some dictionary.

Optional attributes:

- *doc_type* (default),
- *host* (localhost),
- *port* (9200),
- *settings* ({'settings': {}})
- *mapping* (None),
- *chunk_size* (2000),
- *raise_on_error* (True),
- *purge_existing_index* (False),
- *marker_index_hist_size* (0)

If settings are defined, they are only applied at index creation time.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

host

ES hostname.

port

ES port.

http_auth

ES optional http auth information as either ':' separated string or a tuple, e.g. ('user', 'pass') or "user:pass".

index

The target index.

May exist or not.

doc_type

The target doc_type.

mapping

Dictionary with custom mapping or *None*.

settings

Settings to be used at index creation time.

chunk_size

Single API call for this number of docs.

raise_on_error

Whether to fail fast.

purge_existing_index

Whether to delete the *index* completely before any indexing.

marker_index_hist_size

Number of event log entries in the marker index. 0: unlimited.

timeout

Timeout.

extra_elasticsearch_args

Extra arguments to pass to the Elasticsearch constructor

docs ()

Return the documents to be indexed.

Beside the user defined fields, the document may contain an *_index*, *_type* and *_id*.

create_index ()

Override to provide code for creating the target index.

By default it will be created without any special settings or mappings.

delete_index ()

Delete the index, if it exists.

update_id ()

This id will be a unique identifier for this indexing task.

output ()

Returns a ElasticsearchTarget representing the inserted dataset.

Normally you don't override this.

run ()

Run task, namely:

- purge existing index, if requested (*purge_existing_index*),
- create the index, if missing,
- apply mappings, if given,
- set refresh interval to -1 (disable) for performance reasons,
- bulk index in batches of size *chunk_size* (2000),

- set refresh interval to 1s,
- refresh Elasticsearch,
- create entry in marker index.

task_namespace = None

luigi.contrib.ftp module This library is a wrapper of ftplib. It is convenient to move data from/to FTP.

There is an example on how to use it (example/ftp_experiment_outputs.py)

You can also find unittest for each class.

Be aware that normal ftp do not provide secure communication.

class `luigi.contrib.ftp.RemoteFileSystem` (*host, username=None, password=None, port=21, tls=False*)

Bases: `luigi.target.FileSystem`

exists (*path, mtime=None*)

Return *True* if file or directory at *path* exist, *False* otherwise.

Additional check on modified time when *mtime* is passed in.

Return *False* if the file's modified time is older *mtime*.

remove (*path, recursive=True*)

Remove file or directory at location *path*.

Parameters

- **path** (*str*) – a path within the `FileSystem` to remove.
- **recursive** (*bool*) – if the path is a directory, recursively remove the directory and all of its descendants. Defaults to *True*.

put (*local_path, path*)

get (*path, local_path*)

class `luigi.contrib.ftp.AtomicFtpFile` (*fs, path*)

Bases: `luigi.target.AtomicLocalFile`

Simple class that writes to a temp file and upload to ftp on `close()`.

Also cleans up the temp file if `close` is not invoked.

Initializes an `AtomicFtpfile` instance. :param fs: :param path: :type path: str

move_to_final_destination ()

fs

class `luigi.contrib.ftp.RemoteTarget` (*path, host, format=None, username=None, password=None, port=21, mtime=None, tls=False*)

Bases: `luigi.target.FileSystemTarget`

Target used for reading from remote files.

The target is implemented using ssh commands streaming data over the network.

fs

open (*mode*)

Open the `FileSystem` target.

This method returns a file-like object which can either be read from or written to depending on the specified mode.

Parameters *mode* (*str*) – the mode *r* opens the `FileSystemTarget` in read-only mode, whereas *w* will open the `FileSystemTarget` in write mode. Subclasses can implement additional options.

exists ()

put (*local_path*)

get (*local_path*)

luigi.contrib.gcs module luigi bindings for Google Cloud Storage

exception `luigi.contrib.gcs.InvalidDeleteException`

Bases: `luigi.target.FileSystemException`

class `luigi.contrib.gcs.GCSClient` (*oauth_credentials=None*, *descriptor=''*, *http_=None*)

Bases: `luigi.target.FileSystem`

An implementation of a `FileSystem` over Google Cloud Storage.

There are several ways to use this class. By default it will use the app default credentials, as described at <https://developers.google.com/identity/protocols/application-default-credentials>. Alternatively, you may pass an `oauth2client` credentials object. e.g. to use a service account:

```
credentials = oauth2client.client.SignedJwtAssertionCredentials(
    '012345678912-ThisIsARandomServiceAccountEmail@developer.gserviceaccount.com',
    'These are the contents of the p12 file that came with the service account',
    scope='https://www.googleapis.com/auth/devstorage.read_write')
client = GCSClient(oauth_credentials=credentials)
```

Warning: By default this class will use “automated service discovery” which will require a connection to the web. The google api client downloads a JSON file to “create” the library interface on the fly. If you want a more hermetic build, you can pass the contents of this file (currently found at <https://www.googleapis.com/discovery/v1/apis/storage/v1/rest>) as the `descriptor` argument.

exists (*path*)

isdir (*path*)

remove (*path*, *recursive=True*)

put (*filename*, *dest_path*, *mimetype=None*)

put_string (*contents*, *dest_path*, *mimetype=None*)

makedirs (*path*, *parents=True*, *raise_if_exists=False*)

copy (*source_path*, *destination_path*)

rename (*source_path*, *destination_path*)

Rename/move an object from one S3 location to another.

listdir (*path*)

Get an iterable with S3 folder contents. Iterable contains paths relative to queried path.

download (*path*)

class `luigi.contrib.gcs.AtomicGCSFile` (*path*, *gcs_client*)

Bases: `luigi.target.AtomicLocalFile`

A GCS file that writes to a temp file and put to GCS on close.

```
move_to_final_destination()
```

```
class luigi.contrib.gcs.GCSTarget(path, format=None, client=None)
    Bases: luigi.target.FileSystemTarget
```

```
fs = None
```

```
open(mode='r')
```

```
class luigi.contrib.gcs.GCSFlagTarget(path, format=None, client=None, flag='_SUCCESS')
    Bases: luigi.contrib.gcs.GCSTarget
```

Defines a target directory with a flag-file (defaults to `_SUCCESS`) used to signify job success.

This checks for two things:

- the path exists (just like the `GCSTarget`)
- the `_SUCCESS` file exists within the directory.

Because Hadoop outputs into a directory and not a single file, the path is assumed to be a directory.

This is meant to be a handy alternative to `AtomicGCSFile`.

The `AtomicFile` approach can be burdensome for GCS since there are no directories, per se.

If we have 1,000,000 output files, then we have to rename 1,000,000 objects.

Initializes a `S3FlagTarget`.

Parameters

- **path** (*str*) – the directory where the files are stored.
- **client** –
- **flag** (*str*) –

```
fs = None
```

```
exists()
```

luigi.contrib.hadoop module Run Hadoop Mapreduce jobs using Hadoop Streaming. To run a job, you need to subclass `luigi.contrib.hadoop.JobTask` and implement a `mapper` and `reducer` methods. See [Example – Top Artists](#) for an example of how to run a Hadoop job.

```
class luigi.contrib.hadoop.BaseHadoopJobTask(*args, **kwargs)
    Bases: luigi.task.Task
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
batch_counter_default = 1
```

```
deps()
```

```
final_combiner = NotImplemented
```

```
final_mapper = NotImplemented
```

```
final_reducer = NotImplemented
```

init_hadoop()

init_local()

Implement any work to setup any internal datastructure etc here.

You can add extra input using the `requires_local/input_local` methods.

Anything you set on the object will be pickled and available on the Hadoop nodes.

input_hadoop()

input_local()

job_runner()

jobconfs()

mr_priority = NotImplemented

on_failure (*exception*)

pool = <luigi.parameter.Parameter object>

requires_hadoop()

requires_local()

Default impl - override this method if you need any local input to be accessible in `init()`.

run()

task_id = None

task_namespace = None

class `luigi.contrib.hadoop.DefaultHadoopJobRunner`

Bases: `luigi.contrib.hadoop.HadoopJobRunner`

The default job runner just reads from config and sets stuff.

exception `luigi.contrib.hadoop.HadoopJobError` (*message, out=None, err=None*)

Bases: `exceptions.RuntimeError`

class `luigi.contrib.hadoop.HadoopJobRunner` (*streaming_jar, modules=None, streaming_args=None, libjars=None, libjars_in_hdfs=None, jobconfs=None, input_format=None, output_format=None, end_job_with_atomic_move_dir=True*)

Bases: `luigi.contrib.hadoop.JobRunner`

Takes care of uploading & executing a Hadoop job using Hadoop streaming.

TODO: add code to support Elastic Mapreduce (using boto) and local execution.

finish()

run_job (*job*)

class `luigi.contrib.hadoop.HadoopRunContext`

Bases: `object`

kill_job (*captured_signal=None, stack_frame=None*)

class `luigi.contrib.hadoop.JobRunner`

Bases: `object`

run_job = NotImplemented

class `luigi.contrib.hadoop.JobTask(*args, **kwargs)`
 Bases: `luigi.contrib.hadoop.BaseHadoopJobTask`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

add_link (*src, dst*)

combiner = `NotImplemented`

data_interchange_format = `'python'`

deserialize

dump (*directory=''*)

Dump instance to file.

extra_files ()

Can be overridden in subclass.

Each element is either a string, or a pair of two strings (*src, dst*).

- *src* can be a directory (in which case everything will be copied recursively).

- *dst* can include subdirectories (foo/bar/baz.txt etc)

Uses Hadoop's `-files` option so that the same file is reused across tasks.

extra_modules ()

incr_counter (**args, **kwargs*)

Increments a Hadoop counter.

Since counters can be a bit slow to update, this batches the updates.

init_combiner ()

init_mapper ()

init_reducer ()

internal_reader (*input_stream*)

Reader which uses python eval on each part of a tab separated string. Yields a tuple of python objects.

internal_serialize

internal_writer (*outputs, stdout*)

Writer which outputs the python repr for each item.

job_runner ()

Get the MapReduce runner for this job.

If all outputs are `HdfsTargets`, the `DefaultHadoopJobRunner` will be used. Otherwise, the `LocalJobRunner` which streams all data through the local machine will be used (great for testing).

jobconfs ()

mapper (*item*)

Re-define to process an input item (usually a line of input data).

Defaults to identity mapper that sends all lines to the same reducer.

n_reduce_tasks = 25

reader (*input_stream*)

Reader is a method which iterates over input lines and outputs records.

The default implementation yields one argument containing the line for each line in the input.

reducer = NotImplemented

run_combiner (*stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>*)

run_mapper (*stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>*)

Run the mapper on the hadoop node.

run_reducer (*stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>*)

Run the reducer on the hadoop node.

serialize

task_namespace = None

writer (*outputs, stdout, stderr=<open file '<stderr>', mode 'w'>*)

Writer format is a method which iterates over the output records from the reducer and formats them for output.

The default implementation outputs tab separated items.

class `luigi.contrib.hadoop.LocalJobRunner` (*samplelines=None*)

Bases: `luigi.contrib.hadoop.JobRunner`

Will run the job locally.

This is useful for debugging and also unit testing. Tries to mimic Hadoop Streaming.

TODO: integrate with JobTask

group (*input_stream*)

run_job (*job*)

sample (*input_stream, n, output*)

`luigi.contrib.hadoop.attach` (**packages*)

Attach a python package to hadoop map reduce tarballs to make those packages available on the hadoop cluster.

`luigi.contrib.hadoop.create_packages_archive` (*packages, filename*)

Create a tar archive which will contain the files for the packages listed in packages.

`luigi.contrib.hadoop.dereference` (*f*)

`luigi.contrib.hadoop.fetch_task_failures` (*tracking_url*)

Uses mechanize to fetch the actual task logs from the task tracker.

This is highly opportunistic, and we might not succeed. So we set a low timeout and hope it works. If it does not, it's not the end of the world.

TODO: Yarn has a REST API that we should probably use instead:
<http://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/WebServicesIntro.html>

`luigi.contrib.hadoop.flatten` (*sequence*)

A simple generator which flattens a sequence.

Only one level is flattened.

`(1, (2, 3), 4) -> (1, 2, 3, 4)`

`luigi.contrib.hadoop.get_extra_files` (*extra_files*)

class `luigi.contrib.hadoop.hadoop(*args, **kwargs)`
 Bases: `luigi.task.Config`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

pool = `<luigi.parameter.Parameter object>`

task_namespace = `None`

`luigi.contrib.hadoop.run_and_track_hadoop_job(arglist, tracking_url_callback=None, env=None)`

Runs the job by invoking the command from the given arglist. Finds tracking urls from the output and attempts to fetch errors using those urls if the job fails. Throws `HadoopJobError` with information about the error (including stdout and stderr from the process) on failure and returns normally otherwise.

Parameters

- **arglist** –
- **tracking_url_callback** –
- **env** –

Returns

luigi.contrib.hadoop_jar module Provides functionality to run a Hadoop job using a Jar

`luigi.contrib.hadoop_jar.fix_paths(job)`

Coerce input arguments to use temporary files when used for output.

Return a list of temporary file pairs (tmpfile, destination path) and a list of arguments.

Converts each `HdfsTarget` to a string for the path.

exception `luigi.contrib.hadoop_jar.HadoopJarJobError`

Bases: `exceptions.Exception`

class `luigi.contrib.hadoop_jar.HadoopJarJobRunner`

Bases: `luigi.contrib.hadoop.JobRunner`

JobRunner for *hadoop jar* commands. Used to run a `HadoopJarJobTask`.

run_job (*job*)

class `luigi.contrib.hadoop_jar.HadoopJarJobTask(*args, **kwargs)`

Bases: `luigi.contrib.hadoop.BaseHadoopJobTask`

A job task for *hadoop jar* commands that define a jar and (optional) main method.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

jar()
Path to the jar for this Hadoop Job.

main()
optional main method for this Hadoop Job.

job_runner()

atomic_output()
If True, then rewrite output arguments to be temp locations and atomically move them into place after the job finishes.

ssh()
Set this to run hadoop command remotely via ssh. It needs to be a dict that looks like {"host": "myhost", "key_file": None, "username": None, ["no_host_key_check": False]}

args()
Returns an array of args to pass to the job (after hadoop jar <jar> <main>).

task_namespace = None

luigi.contrib.hive module

exception `luigi.contrib.hive.HiveCommandError` (*message*, *out=None*, *err=None*)

Bases: `exceptions.RuntimeError`

`luigi.contrib.hive.load_hive_cmd()`

`luigi.contrib.hive.get_hive_syntax()`

`luigi.contrib.hive.run_hive` (*args*, *check_return_code=True*)

Runs the *hive* from the command line, passing in the given args, and returning stdout.

With the apache release of Hive, so of the table existence checks (which are done using DESCRIBE do not exit with a return code of 0 so we need an option to ignore the return code and just return stdout for parsing

`luigi.contrib.hive.run_hive_cmd` (*hivecmd*, *check_return_code=True*)

Runs the given hive query and returns stdout.

`luigi.contrib.hive.run_hive_script` (*script*)

Runs the contents of the given script in hive and returns stdout.

class `luigi.contrib.hive.HiveClient`

Bases: `object`

table_location (*table*, *database='default'*, *partition=None*)

Returns location of db.table (or db.table.partition). partition is a dict of partition key to value.

table_schema (*table*, *database='default'*)

Returns list of [(name, type)] for each column in database.table.

table_exists (*table*, *database='default'*, *partition=None*)

Returns true if db.table (or db.table.partition) exists. partition is a dict of partition key to value.

partition_spec (*partition*)

Turn a dict into a string partition specification

class `luigi.contrib.hive.HiveCommandClient`

Bases: `luigi.contrib.hive.HiveClient`

Uses *hive* invocations to find information.

table_location (*table*, *database='default'*, *partition=None*)

table_exists (*table*, *database='default'*, *partition=None*)

table_schema (*table*, *database*='default')

partition_spec (*partition*)

Turns a dict into the a Hive partition specification string.

class `luigi.contrib.hive.ApacheHiveCommandClient`

Bases: `luigi.contrib.hive.HiveCommandClient`

A subclass for the HiveCommandClient to (in some cases) ignore the return code from the hive command so that we can just parse the output.

table_schema (*table*, *database*='default')

class `luigi.contrib.hive.MetastoreClient`

Bases: `luigi.contrib.hive.HiveClient`

table_location (*table*, *database*='default', *partition*=None)

table_exists (*table*, *database*='default', *partition*=None)

table_schema (*table*, *database*='default')

partition_spec (*partition*)

class `luigi.contrib.hive.HiveThriftContext`

Bases: `object`

Context manager for hive metastore client.

`luigi.contrib.hive.get_default_client()`

class `luigi.contrib.hive.HiveQueryTask` (**args*, ***kwargs*)

Bases: `luigi.contrib.hadoop.BaseHadoopJobTask`

Task to run a hive query.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

n_reduce_tasks = None

bytes_per_reducer = None

reducers_max = None

query ()

Text of query to run in hive

hiverc ()

Location of an rc file to run before the query if hiverc-location key is specified in client.cfg, will default to the value there otherwise returns None.

Returning a list of rc files will load all of them in order.

hiveconfs ()

Returns an dict of key=value settings to be passed along to the hive command line via `-hiveconf`. By default, sets `mapred.job.name` to `task_id` and if not None, sets:

- `mapred.reduce.tasks` (`n_reduce_tasks`)
- `mapred.fairscheduler.pool` (`pool`) or `mapred.job.queue.name` (`pool`)

- hive.exec.reducers.bytes.per.reducer (bytes_per_reducer)

- hive.exec.reducers.max (reducers_max)

job_runner()

task_namespace = None

class `luigi.contrib.hive.HiveQueryRunner`

Bases: `luigi.contrib.hadoop.JobRunner`

Runs a HiveQueryTask by shelling out to hive.

prepare_outputs (*job*)

Called before job is started.

If output is a *FileSystemTarget*, create parent directories so the hive command won't fail

run_job (*job*)

class `luigi.contrib.hive.HiveTableTarget` (*table*, *database='default'*, *client=None*)

Bases: `luigi.target.Target`

exists returns true if the table exists.

exists ()

path

Returns the path to this table in HDFS.

open (*mode*)

class `luigi.contrib.hive.HivePartitionTarget` (*table*, *partition*, *database='default'*,
fail_missing_table=True, *client=None*)

Bases: `luigi.target.Target`

exists returns true if the table's partition exists.

exists ()

path

Returns the path for this HiveTablePartitionTarget's data.

open (*mode*)

class `luigi.contrib.hive.ExternalHiveTask` (**args*, ***kwargs*)

Bases: `luigi.task.ExternalTask`

External task that depends on a Hive table/partition.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

database = <luigi.parameter.Parameter object>

table = <luigi.parameter.Parameter object>

partition = <luigi.parameter.Parameter object>

task_namespace = None

output ()

luigi.contrib.mysqlldb module

class `luigi.contrib.mysqlldb.MySqlTarget` (*host, database, user, password, table, update_id*)

Bases: `luigi.target.Target`

Target for a resource in MySQL.

Initializes a MySqlTarget instance.

Parameters

- **host** (*str*) – MySQL server address. Possibly a host:port string.
- **database** (*str*) – database name.
- **user** (*str*) – database user
- **password** (*str*) – password for specified user.
- **update_id** (*str*) – an identifier for this data set.

marker_table = 'table_updates'

touch (*connection=None*)

Mark this update as complete.

IMPORTANT, If the marker table doesn't exist, the connection transaction will be aborted and the connection reset. Then the marker table will be created.

exists (*connection=None*)

connect (*autocommit=False*)

create_marker_table ()

Create marker table if it doesn't exist.

Using a separate connection since the transaction might have to be reset.

luigi.contrib.pig module Apache Pig support. Example configuration section in client.cfg:

```
[pig]
# pig home directory
home: /usr/share/pig
```

class `luigi.contrib.pig.PigJobTask` (**args, **kwargs*)

Bases: `luigi.task.Task`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

pig_home ()

pig_command_path ()

pig_env_vars ()

Dictionary of environment variables that should be set when running Pig.

Ex:: return { 'PIG_CLASSPATH': '/your/path' }

pig_properties()

Dictionary of properties that should be set when running Pig.

Example:

```
return { 'pig.additional.jars': '/path/to/your/jar' }
```

pig_parameters()

Dictionary of parameters that should be set for the Pig job.

Example:

```
return { 'YOUR_PARAM_NAME': 'Your param value' }
```

pig_options()

List of options that will be appended to the Pig command.

Example:

```
return ['-x', 'local']
```

output()

pig_script_path()

Return the path to the Pig script to be run.

run()

track_and_progress(cmd)

task_namespace = None

class `luigi.contrib.pig.PigRunContext`

Bases: `object`

kill_job(*captured_signal=None, stack_frame=None*)

exception `luigi.contrib.pig.PigJobError`(*message, out=None, err=None*)

Bases: `exceptions.RuntimeError`

luigi.contrib.pyspark_runner module The pyspark program.

This module will be run by spark-submit for PySparkTask jobs.

The first argument is a path to the pickled instance of the PySparkTask, other arguments are the ones returned by `PySparkTask.app_options()`

class `luigi.contrib.pyspark_runner.PySparkRunner`(*job, *args*)

Bases: `object`

run()

luigi.contrib.rdbms module A common module for posgres like databases, such as postgres or redshift

class `luigi.contrib.rdbms.CopyToTable`(**args, **kwargs*)

Bases: `luigi.task.MixinNaiveBulkComplete, luigi.task.Task`

An abstract task for inserting a data set into RDBMS.

Usage:

Subclass and override the following attributes:

- host*,

- database*,
- user*,
- password*,
- table*
- columns*

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

host

database

user

password

table

columns = []

null_values = (None,)

column_separator = '\t'

create_table (*connection*)

Override to provide code for creating the target table.

By default it will be created using types (optionally) specified in columns.

If overridden, use the provided connection object for setting up the table in order to create the table and insert data using the same transaction.

update_id ()

This update id will be a unique identifier for this insert on this table.

output ()

init_copy (*connection*)

Override to perform custom queries.

Any code here will be formed in the same transaction as the main copy, just prior to copying data. Example use cases include truncating the table or removing all data older than X in the database to keep a rolling window of data available in the table.

copy (*cursor*, *file*)

task_namespace = None

luigi.contrib.redis_store module

class `luigi.contrib.redis_store.RedisTarget` (*host*, *port*, *db*, *update_id*, *password*=None, *socket_timeout*=None, *expire*=None)

Bases: `luigi.target.Target`

Target for a resource in Redis.

Parameters

- **host** (*str*) – Redis server host
- **port** (*int*) – Redis server port
- **db** (*int*) – database index
- **update_id** (*str*) – an identifier for this data hash
- **password** (*str*) – a password to connect to the redis server
- **socket_timeout** (*int*) – client socket timeout
- **expire** (*int*) – timeout before the target is deleted

marker_prefix = <luigi.parameter.Parameter object>

marker_key ()

Generate a key for the indicator hash.

touch ()

Mark this update as complete.

We index the parameters *update_id* and *date*.

exists ()

Test, if this task has been run.

luigi.contrib.redshift module

class luigi.contrib.redshift.**RedshiftTarget** (*host, database, user, password, table, update_id*)

Bases: *luigi.postgres.PostgresTarget*

Target for a resource in Redshift.

Redshift is similar to postgres with a few adjustments required by redshift.

Args: *host* (*str*): Postgres server address. Possibly a host:port string. *database* (*str*): Database name *user* (*str*): Database user *password* (*str*): Password for specified user *update_id* (*str*): An identifier for this data set

marker_table = 'table_updates'

use_db_timestamps = False

class luigi.contrib.redshift.**S3CopyToTable** (**args, **kwargs*)

Bases: *luigi.contrib.rdbms.CopyToTable*

Template task for inserting a data set into Redshift from s3.

Usage:

•Subclass and override the required attributes: ** host, * database, * user, * password, * table, * columns, * aws_access_key_id, * aws_secret_access_key, * s3_load_path*.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

s3_load_path

Override to return the load path.

aws_access_key_id

Override to return the key id.

aws_secret_access_key

Override to return the secret access key.

copy_options

Add extra copy options, for example:

- TIMEFORMAT 'auto'
- IGNOREHEADER 1
- TRUNCATECOLUMNS
- IGNOREBLANKLINES

table_attributes()

Add extra table attributes, for example: DISTSTYLE KEY DISTKEY (MY_FIELD) SORTKEY (MY_FIELD_2, MY_FIELD_3)

do_truncate_table()

Return True if table should be truncated before copying new data in.

truncate_table(connection)**create_table(connection)**

Override to provide code for creating the target table.

By default it will be created using types (optionally) specified in columns.

If overridden, use the provided connection object for setting up the table in order to create the table and insert data using the same transaction.

run()

If the target table doesn't exist, self.create_table will be called to attempt to create the table.

copy(cursor, f)

Defines copying from s3 into redshift.

output()

Returns a RedshiftTarget representing the inserted dataset.

Normally you don't override this.

does_table_exist(connection)

Determine whether the table already exists.

task_namespace = None

class luigi.contrib.redshift.S3CopyJSONToTable(*args, **kwargs)

Bases: `luigi.contrib.redshift.S3CopyToTable`

Template task for inserting a JSON data set into Redshift from s3.

Usage:

- Subclass and override the required attributes:

- host,
- database,
- user,
- password,
- table,
- columns,

```
-aws_access_key_id,  
-aws_secret_access_key,  
-s3_load_path,  
-jsonpath,  
-copy_json_options.
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

jsonpath

Override the jsonpath schema location for the table.

copy_json_options

Add extra copy options, for example:

- GZIP
- LZOP

copy (cursor, f)

Defines copying JSON from s3 into redshift.

task_namespace = None

class `luigi.contrib.redshift.RedshiftManifestTask(*args, **kwargs)`

Bases: `luigi.s3.S3PathTask`

Generic task to generate a manifest file that can be used in `S3CopyToTable` in order to copy multiple files from your s3 folder into a redshift table at once.

For full description on how to use the manifest file see <http://docs.aws.amazon.com/redshift/latest/dg/loading-data-files-using-manifest.html>

Usage:

•requires parameters

- **path - s3 path to the generated manifest file, including the** name of the generated file to be copied into a redshift table
- **folder_paths - s3 paths to the folders containing files you wish to be copied**

Output:

- generated manifest file

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

folder_paths = <luigi.parameter.Parameter object>

text_target = True

```
run()
```

```
task_namespace = None
```

```
class luigi.contrib.redshift.KillOpenRedshiftSessions(*args, **kwargs)
```

Bases: `luigi.task.Task`

An task for killing any open Redshift sessions in a given database. This is necessary to prevent open user sessions with transactions against the table from blocking drop or truncate table commands.

Usage:

Subclass and override the required *host*, *database*, *user*, and *password* attributes.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
connection_reset_wait_seconds = <luigi.parameter.IntParameter object>
```

```
host
```

```
database
```

```
user
```

```
password
```

```
update_id()
```

This update id will be a unique identifier for this insert on this table.

```
output()
```

Returns a `RedshiftTarget` representing the inserted dataset.

Normally you don't override this.

```
run()
```

Kill any open Redshift sessions for the given database.

```
task_namespace = None
```

luigi.contrib.scalding module

```
luigi.contrib.scalding.logger = <logging.Logger object>
```

Scalding support for Luigi.

Example configuration section in `client.cfg`:

```
[scalding]
# scala home directory, which should include a lib subdir with scala jars.
scala-home: /usr/share/scala

# scalding home directory, which should include a lib subdir with
# scalding-*-assembly-* jars as built from the official Twitter build script.
scalding-home: /usr/share/scalding

# provided dependencies, e.g. jars required for compiling but not executing
# scalding jobs. Currently required jars:
# org.apache.hadoop/hadoop-core/0.20.2
# org.slf4j/slf4j-log4j12/1.6.6
```

```
# log4j/log4j/1.2.15
# commons-httpclient/commons-httpclient/3.1
# commons-cli/commons-cli/1.2
# org.apache.zookeeper/zookeeper/3.3.4
scalding-provided: /usr/share/scalding/provided

# additional jars required.
scalding-libjars: /usr/share/scalding/libjars
```

class `luigi.contrib.scalding.ScaldingJobRunner`

Bases: `luigi.contrib.hadoop.JobRunner`

JobRunner for *pyscald* commands. Used to run a ScaldingJobTask.

get_scala_jars (*include_compiler=False*)

get_scalding_jars ()

get_scalding_core ()

get_provided_jars ()

get_libjars ()

get_tmp_job_jar (*source*)

get_build_dir (*source*)

get_job_class (*source*)

build_job_jar (*job*)

run_job (*job*)

class `luigi.contrib.scalding.ScaldingJobTask` (**args, **kwargs*)

Bases: `luigi.contrib.hadoop.BaseHadoopJobTask`

A job task for Scalding that define a scala source and (optional) main method.

`requires()` should return a dictionary where the keys are Scalding argument names and values are sub tasks or lists of subtasks.

For example:

```
{'input1': A, 'input2': C} => --input1 <Aoutput> --input2 <Coutput>
{'input1': [A, B], 'input2': [C]} => --input1 <Aoutput> <Boutput> --input2 <Coutput>
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

relpath (*current_file, rel_path*)

Compute path given current file and relative path.

source ()

Path to the scala source for this Scalding Job

Either one of `source()` or `jar()` must be specified.

jar ()

Path to the jar file for this Scalding Job

Either one of `source()` or `jar()` must be specified.

extra_jars()

Extra jars for building and running this Scalding Job.

job_class()

optional main job class for this Scalding Job.

job_runner()

atomic_output()

If True, then rewrite output arguments to be temp locations and atomically move them into place after the job finishes.

requires()

job_args()

Extra arguments to pass to the Scalding job.

task_namespace = None

args()

Returns an array of args to pass to the job.

luigi.contrib.spark module

class `luigi.contrib.spark.SparkRunContext` (*proc*)

Bases: `object`

kill_job (*captured_signal=None, stack_frame=None*)

exception `luigi.contrib.spark.SparkJobError` (*message, out=None, err=None*)

Bases: `exceptions.RuntimeError`

class `luigi.contrib.spark.SparkSubmitTask` (**args, **kwargs*)

Bases: `luigi.task.Task`

Template task for running a Spark job

Supports running jobs on Spark local, standalone, Mesos or Yarn

See <http://spark.apache.org/docs/latest/submitting-applications.html> for more information

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

name = None

entry_class = None

app = None

app_options()

Subclass this method to map your task parameters to the app's arguments

spark_submit

master

deploy_mode

```
jars
py_files
files
conf
properties_file
driver_memory
driver_java_options
driver_library_path
driver_class_path
executor_memory
driver_cores
supervise
total_executor_cores
executor_cores
queue
num_executors
archives
hadoop_conf_dir
get_environment()
spark_command()
app_command()
run()
task_namespace = None
```

```
class luigi.contrib.spark.PySparkTask(*args, **kwargs)
    Bases: luigi.contrib.spark.SparkSubmitTask
```

Template task for running an inline PySpark job

Simply implement the `main` method in your subclass

You can optionally define package names to be distributed to the cluster with `py_packages` (uses luigi's global py-packages configuration by default)

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
app = '/home/docs/checkouts/readthedocs.org/user_builds/luigi/checkouts/latest/luigi/contrib/pyspark_runner.py'
```

```
deploy_mode = 'client'
```

```
name
```


py_packages**setup** (*conf*)

Called by the pyspark_runner with a SparkConf instance that will be used to instantiate the SparkContext

Parameters *conf* – SparkConf

setup_remote (*sc*)**main** (*sc*, **args*)

Called by the pyspark_runner with a SparkContext and any arguments returned by app_options()

Parameters

- **sc** – SparkContext
- **args** – arguments list

app_command ()**run** ()**task_namespace** = None

class luigi.contrib.spark.**SparkJob** (**args*, ***kwargs*)

Bases: [luigi.task.Task](#)

Deprecated since version 1.1.1: Use SparkSubmitTask or PySparkTask instead.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as MyTask(count=10).

spark_workers = None**spark_master_memory** = None**spark_worker_memory** = None**queue** = <luigi.parameter.Parameter object>**temp_hadoop_output_file** = None**requires_local** ()

Default impl - override this method if you need any local input to be accessible in init().

requires_hadoop ()**input_local** ()**input** ()**deps** ()**jar** ()**job_class** ()**job_args** ()**output** ()**run** ()**track_progress** (*proc*)

task_namespace = None

class `luigi.contrib.spark.Spark1xBackwardCompat` (*args, **kwargs)

Bases: `luigi.contrib.spark.SparkSubmitTask`

Adapts SparkSubmitTask interface to (Py)Spark1xJob interface

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

master

output()

spark_options()

dependency_jars()

job_args()

jars

app_options()

spark_command()

task_namespace = None

class `luigi.contrib.spark.Spark1xJob` (*args, **kwargs)

Bases: `luigi.contrib.spark.Spark1xBackwardCompat`

Deprecated since version 1.1.1: Use SparkSubmitTask or PySparkTask instead.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

job_class()

jar()

entry_class

app

run()

task_namespace = None

class `luigi.contrib.spark.PySpark1xJob` (*args, **kwargs)

Bases: `luigi.contrib.spark.Spark1xBackwardCompat`

Deprecated since version 1.1.1: Use SparkSubmitTask or PySparkTask instead.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
program()

app

run()

task_namespace = None
```

luigi.contrib.sparkey module

```
class luigi.contrib.sparkey.SparkeyExportTask(*args, **kwargs)
    Bases: luigi.task.Task
```

A luigi task that writes to a local sparkey log file.

Subclasses should implement the `requires` and `output` methods. The output must be a `luigi.LocalTarget`.

The resulting sparkey log file will contain one entry for every line in the input, mapping from the first value to a tab-separated list of the rest of the line.

To generate a simple key-value index, yield “key”, “value” pairs from the input(s) to this task.

```
separator = '\t'

run()

task_namespace = None
```

luigi.contrib.sqla module Support for SQLAlchemy. Provides `SQLAlchemyTarget` for storing in databases supported by SQLAlchemy. The user would be responsible for installing the required database driver to connect using SQLAlchemy.

Minimal example of a job to copy data to database using SQLAlchemy is as shown below:

```
from sqlalchemy import String
import luigi
from luigi.contrib import sqla

class SQLATask(sqla.CopyToTable):
    # columns defines the table schema, with each element corresponding
    # to a column in the format (args, kwargs) which will be sent to
    # the sqlalchemy.Column(*args, **kwargs)
    columns = [
        (["item", String(64)], {"primary_key": True}),
        (["property", String(64)], {})
    ]
    connection_string = "sqlite://" # in memory SQLite database
    table = "item_property" # name of the table to store data

    def rows(self):
        for row in [("item1", "property1"), ("item2", "property2")]:
            yield row

if __name__ == '__main__':
    task = SQLATask()
    luigi.build([task], local_scheduler=True)
```

If the target table where the data needs to be copied already exists, then the column schema definition can be skipped and instead the reflect flag can be set as True. Here is a modified version of the above example:

```
from sqlalchemy import String
import luigi
from luigi.contrib import sqlla

class SQLATask(sqlla.CopyToTable):
    # If database table is already created, then the schema can be loaded
    # by setting the reflect flag to True
    reflect = True
    connection_string = "sqlite://" # in memory SQLite database
    table = "item_property" # name of the table to store data

    def rows(self):
        for row in [("item1", "property1"), ("item2", "property2")]:
            yield row

if __name__ == '__main__':
    task = SQLATask()
    luigi.build([task], local_scheduler=True)
```

In the above examples, the data that needs to be copied was directly provided by overriding the rows method. Alternatively, if the data comes from another task, the modified example would look as shown below:

```
from sqlalchemy import String
import luigi
from luigi.contrib import sqlla
from luigi.mock import MockFile

class BaseTask(luigi.Task):
    def output(self):
        return MockFile("BaseTask")

    def run(self):
        out = self.output().open("w")
        TASK_LIST = ["item%d\tproperty%d\n" % (i, i) for i in range(10)]
        for task in TASK_LIST:
            out.write(task)
        out.close()

class SQLATask(sqlla.CopyToTable):
    # columns defines the table schema, with each element corresponding
    # to a column in the format (args, kwargs) which will be sent to
    # the sqlalchemy.Column(*args, **kwargs)
    columns = [
        ("item", String(64), {"primary_key": True}),
        ("property", String(64), {})
    ]
    connection_string = "sqlite://" # in memory SQLite database
    table = "item_property" # name of the table to store data

    def requires(self):
        return BaseTask()

if __name__ == '__main__':
    task1, task2 = SQLATask(), BaseTask()
    luigi.build([task1, task2], local_scheduler=True)
```

In the above example, the output from *BaseTask* is copied into the database. Here we did not have to implement the *rows* method because by default *rows* implementation assumes every line is a row with column values separated by a tab. One can define *column_separator* option for the task if the values are say comma separated instead of tab separated.

You can pass in database specific connection arguments by setting the *connect_args* dictionary. The options will be passed directly to the DBAPI's connect method as keyword arguments.

The other option to *sqla.CopyToTable* that can be of help with performance aspect is the *chunk_size*. The default is 5000. This is the number of rows that will be inserted in a transaction at a time. Depending on the size of the inserts, this value can be tuned for performance.

See here for a [tutorial on building task pipelines using luigi](#) and using [SQLAlchemy in workflow pipelines](#).

Author: Gouthaman Balaraman Date: 01/02/2015

```
class luigi.contrib.sqla.SQLAlchemyTarget (connection_string,    target_table,    update_id,
                                           echo=False, connect_args=None)
```

Bases: *luigi.target.Target*

Database target using SQLAlchemy.

This will rarely have to be directly instantiated by the user.

Typical usage would be to override *luigi.contrib.sqla.CopyToTable* class to create a task to write to the database.

Constructor for the SQLAlchemyTarget.

Parameters

- **connection_string** (*str*) – SQLAlchemy connection string
- **target_table** (*str*) – The table name for the data
- **update_id** (*str*) – An identifier for this data set
- **echo** (*bool*) – Flag to setup SQLAlchemy logging
- **connect_args** (*dict*) – A dictionary of connection arguments

Returns

marker_table = None

```
class Connection (engine, pid)
```

Bases: tuple

engine

Alias for field number 0

pid

Alias for field number 1

```
SQLAlchemyTarget.connect_args = {}
```

```
SQLAlchemyTarget.engine
```

Return an engine instance, creating it if it doesn't exist.

Recreate the engine connection if it wasn't originally created by the current process.

```
SQLAlchemyTarget.touch()
```

Mark this update as complete.

```
SQLAlchemyTarget.exists()
```

SQLAlchemyTarget.**create_marker_table**()

Create marker table if it doesn't exist.

Using a separate connection since the transaction might have to be reset.

SQLAlchemyTarget.**open**(mode)

class luigi.contrib.sqa.**CopyToTable**(*args, **kwargs)

Bases: *luigi.task.Task*

An abstract task for inserting a data set into SQLAlchemy RDBMS

Usage:

- subclass and override the required *connection_string*, *table* and *columns* attributes.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

echo = False

connect_args = {}

connection_string()

table

columns = []

column_separator = '\t'

chunk_size = 5000

reflect = False

create_table(engine)

Override to provide code for creating the target table.

By default it will be created using types specified in columns. If the table exists, then it binds to the existing table.

If overridden, use the provided connection object for setting up the table in order to create the table and insert data using the same transaction. :param engine: The sqlalchemy engine instance :type engine: object

update_id()

This update id will be a unique identifier for this insert on this table.

output()

rows()

Return/yield tuples or lists corresponding to each row to be inserted.

This method can be overridden for custom file types or formats.

run()

copy(conn, ins_rows, table_bound)

This method does the actual insertion of the rows of data given by ins_rows into the database. A task that needs row updates instead of insertions should overload this method. :param conn: The sqlalchemy connection object :param ins_rows: The dictionary of rows with the keys in the format `<column_name>`. For example if you have a table with a column name "property", then the key in the dictionary would be

“_property”. This format is consistent with the bindparam usage in sqlalchemy. :param table_bound: The object referring to the table :return:

```
task_namespace = None
```

luigi.contrib.ssh module Light-weight remote execution library and utilities.

There are some examples in the unittest, but I added another more luigi-specific in the examples directory (examples/ssh_remote_execution.py)

RemoteContext is meant to provide functionality similar to that of the standard library subprocess module, but where the commands executed are run on a remote machine instead, without the user having to think about prefixing everything with “ssh” and credentials etc.

Using this mini library (which is just a convenience wrapper for subprocess), *RemoteTarget* is created to let you stream data from a remotely stored file using the luigi *FileSystemTarget* semantics.

As a bonus, *RemoteContext* also provides a really cool feature that let’s you set up ssh tunnels super easily using a python context manager (there is an example in the integration part of unittests).

This can be super convenient when you want secure communication using a non-secure protocol or circumvent firewalls (as long as they are open for ssh traffic).

```
class luigi.contrib.ssh.RemoteContext (host, **kwargs)
```

Bases: object

```
Popen (cmd, **kwargs)
```

Remote Popen.

```
check_output (cmd)
```

Execute a shell command remotely and return the output.

Simplified version of Popen when you only want the output as a string and detect any errors.

```
tunnel (*args, **kws)
```

Open a tunnel between localhost:local_port and remote_host:remote_port via the host specified by this context.

Remember to close() the returned “tunnel” object in order to clean up after yourself when you are done with the tunnel.

```
class luigi.contrib.ssh.RemoteFileSystem (host, **kwargs)
```

Bases: *luigi.target.FileSystem*

```
exists (path)
```

Return *True* if file or directory at *path* exist, False otherwise.

```
listdir (path)
```

```
isdir (path)
```

Return *True* if directory at *path* exist, False otherwise.

```
remove (path, recursive=True)
```

Remove file or directory at location *path*.

```
mkdir (path, parents=True, raise_if_exists=False)
```

```
put (local_path, path)
```

```
get (path, local_path)
```

```
class luigi.contrib.ssh.AtomicRemoteFileWriter (fs, path)
```

Bases: *luigi.format.OutputPipeProcessWrapper*

```
close()
```

```
tmp_path
```

```
fs
```

```
class luigi.contrib.ssh.RemoteTarget (path, host, format=None, **kwargs)
```

Bases: [luigi.target.FileSystemTarget](#)

Target used for reading from remote files.

The target is implemented using ssh commands streaming data over the network.

```
fs
```

```
open (mode='r')
```

```
put (local_path)
```

```
get (local_path)
```

luigi.contrib.target module

```
class luigi.contrib.target.CascadingClient (clients, method_names=None)
```

Bases: object

A FileSystemClient that will cascade failing function calls through a list of clients.

Which clients are used are specified at time of construction.

```
ALL_METHOD_NAMES = ['exists', 'rename', 'remove', 'chmod', 'chown', 'count', 'copy', 'get', 'put', 'mkdir', 'list', 'listdir']
```

luigi.contrib.webhdfs module Provides a [WebHdfsTarget](#) and [WebHdfsClient](#) using the [Python hdfs](#)

```
class luigi.contrib.webhdfs.WebHdfsTarget (path, client=None, format=None)
```

Bases: [luigi.target.FileSystemTarget](#)

```
fs = None
```

```
open (mode='r')
```

```
class luigi.contrib.webhdfs.ReadableWebHdfsFile (path, client)
```

Bases: object

```
read()
```

```
readlines (char='\n')
```

```
close()
```

```
class luigi.contrib.webhdfs.AtomicWebHdfsFile (path, client)
```

Bases: [luigi.target.AtomicLocalFile](#)

An Hdfs file that writes to a temp file and put to WebHdfs on close.

```
move_to_final_destination()
```

```
class luigi.contrib.webhdfs.WebHdfsClient (host=None, port=None, user=None)
```

Bases: object

```
get_config (key)
```

```
walk (path, depth=1)
```

```
exists (path)
```

Returns true if the path exists and false otherwise.


```

upload (hdfs_path, local_path, overwrite=False)
download (hdfs_path, local_path, overwrite=False, n_threads=-1)
remove (hdfs_path, recursive=False)
read (hdfs_path, offset=0, length=None, buffer_size=None, chunk_size=1024, buffer_char=None)

```

Module contents

Package containing optional and-on functionality.

luigi.tools package

Submodules

luigi.tools.deps module

```

luigi.tools.deps.get_task_requires (task)
luigi.tools.deps.dfs_paths (start_task, goal_task_family, path=None)

```

```

class luigi.tools.deps.upstream (*args, **kwargs)
    Bases: luigi.task.Config

```

Used to provide the parameter upstream-task-family

Constructor to resolve values for all Parameters.

For example, the Task:

```

class MyTask(luigi.Task):
    count = luigi.IntParameter()

```

can be instantiated as `MyTask(count=10)`.

family = <luigi.parameter.Parameter object>

task_namespace = None

```

luigi.tools.deps.find_deps (task, upstream_task_family)
    Finds all dependencies that start with the given task and have a path to upstream_task_family

    Returns all deps on all paths between task and upstream

```

```

luigi.tools.deps.find_deps_cli ()
    Finds all tasks on all paths from provided CLI task

```

```

luigi.tools.deps.main ()

```

luigi.tools.luigi_grep module

```

class luigi.tools.luigi_grep.LuigiGrep (host, port)
    Bases: object

```

graph_url

```

prefix_search (job_name_prefix)
    searches for jobs matching the given job_name_prefix.

```

```

status_search (status)
    searches for jobs matching the given status

```

```

luigi.tools.luigi_grep.main ()

```

luigi.tools.parse_task module`luigi.tools.parse_task.id_to_name_and_params (task_id)`

Turn a task_id into a (task_family, {params}) tuple.

E.g. calling with `Foo(bar=bar, baz=baz)` returns `('Foo', {'bar': 'bar', 'baz': 'baz'})`.

luigi.tools.range module Produces contiguous completed ranges of recurring tasks.

See `RangeDaily` and `RangeHourly` for basic usage.

Caveat - if gaps accumulate, their causes (e.g. missing dependencies) going unmonitored/unmitigated, then this will eventually keep retrying the same gaps over and over and make no progress to more recent times. (See 'task_limit' and 'reverse' parameters.) TODO foolproof against that kind of misuse?

class `luigi.tools.range.RangeEvent`

Bases: `luigi.event.Event`

Events communicating useful metrics.

`COMPLETE_COUNT` would normally be nondecreasing, and its derivative would describe performance (how many instances complete invocation-over-invocation).

`COMPLETE_FRACTION` reaching 1 would be a telling event in case of a backfill with defined start and stop. Would not be strikingly useful for a typical recurring task without stop defined, fluctuating close to 1.

`DELAY` is measured from the first found missing datehour till (current time + hours_forward), or till stop if it is defined. In hours for Hourly. TBD different units for other frequencies? TODO any different for reverse mode? From first missing till last missing? From last gap till stop?

`COMPLETE_COUNT` = 'event.tools.range.complete.count'

`COMPLETE_FRACTION` = 'event.tools.range.complete.fraction'

`DELAY` = 'event.tools.range.delay'

class `luigi.tools.range.RangeBase (*args, **kwargs)`

Bases: `luigi.task.WrapperTask`

Produces a contiguous completed range of a recurring task.

Made for the common use case where a task is parameterized by e.g. `DateParameter`, and assurance is needed that any gaps arising from downtime are eventually filled.

Emits events that one can use to monitor gaps and delays.

At least one of start and stop needs to be specified.

(This is quite an abstract base class for subclasses with different datetime parameter class, e.g. `DateParameter`, `DateHourParameter`, ..., and different parameter naming, e.g. `days_back/forward`, `hours_back/forward`, ..., as well as different documentation wording, for good user experience.)

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

of = <luigi.parameter.Parameter object>

start = <luigi.parameter.Parameter object>

stop = <luigi.parameter.Parameter object>

reverse = <luigi.parameter.BoolParameter object>

task_limit = <luigi.parameter.IntParameter object>

now = <luigi.parameter.IntParameter object>

datetime_to_parameter (*dt*)

parameter_to_datetime (*p*)

moving_start (*now*)

Returns a datetime from which to ensure contiguousness in the case when start is None or unfeasibly far back.

moving_stop (*now*)

Returns a datetime till which to ensure contiguousness in the case when stop is None or unfeasibly far forward.

finite_datetimes (*finite_start*, *finite_stop*)

Returns the individual datetimes in interval [*finite_start*, *finite_stop*) for which task completeness should be required, as a sorted list.

requires ()

missing_datetimes (*task_cls*, *finite_datetimes*)

Override in subclasses to do bulk checks.

Returns a sorted list.

This is a conservative base implementation that brutally checks completeness, instance by instance.

Inadvisable as it may be slow.

task_namespace = None

class luigi.tools.range.**RangeDailyBase** (*args, **kwargs)

Bases: [luigi.tools.range.RangeBase](#)

Produces a contiguous completed range of a daily recurring task.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

start = <luigi.parameter.DateParameter object>

stop = <luigi.parameter.DateParameter object>

days_back = <luigi.parameter.IntParameter object>

days_forward = <luigi.parameter.IntParameter object>

datetime_to_parameter (*dt*)

parameter_to_datetime (*p*)

moving_start (*now*)

moving_stop (*now*)

finite_datetimes (*finite_start*, *finite_stop*)

Simply returns the points in time that correspond to turn of day.

task_namespace = None

class `luigi.tools.range.RangeHourlyBase` (**args*, ***kwargs*)

Bases: `luigi.tools.range.RangeBase`

Produces a contiguous completed range of an hourly recurring task.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

start = <luigi.parameter.DateHourParameter object>

stop = <luigi.parameter.DateHourParameter object>

hours_back = <luigi.parameter.IntParameter object>

hours_forward = <luigi.parameter.IntParameter object>

datetime_to_parameter (*dt*)

parameter_to_datetime (*p*)

moving_start (*now*)

moving_stop (*now*)

finite_datetimes (*finite_start*, *finite_stop*)

Simply returns the points in time that correspond to whole hours.

task_namespace = None

`luigi.tools.range.most_common` (*items*)

Wanted functionality from Counters (new in Python 2.7).

`luigi.tools.range.infer_bulk_complete_from_fs` (*datetimes*, *datetime_to_task*, *date-*
time_to_re)

Efficiently determines missing datetimes by filesystem listing.

The current implementation works for the common case of a task writing output to a `FileSystemTarget` whose path is built using `strftime` with format like `'...%Y...%m...%d...%H...'`, without custom `complete()` or `exists()`.

(Eventually Luigi could have ranges of completion as first-class citizens. Then this listing business could be factored away/be provided for explicitly in target API or some kind of a history server.)

class `luigi.tools.range.RangeDaily` (**args*, ***kwargs*)

Bases: `luigi.tools.range.RangeDailyBase`

Efficiently produces a contiguous completed range of a daily recurring task that takes a single `DateParameter`.

Falls back to infer it from output filesystem listing to facilitate the common case usage.

Convenient to use even from command line, like:

```
luigi --module your.module RangeDaily --of YourActualTask --start 2014-01-01
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

missing_datetimes (*task_cls, finite_datetimes*)

task_namespace = None

```
class luigi.tools.range.RangeHourly(*args, **kwargs)
    Bases: luigi.tools.range.RangeHourlyBase
```

Efficiently produces a contiguous completed range of an hourly recurring task that takes a single DateHourParameter.

Benefits from bulk_complete information to efficiently cover gaps.

Falls back to infer it from output filesystem listing to facilitate the common case usage.

Convenient to use even from command line, like:

```
luigi --module your.module RangeHourly --of YourActualTask --start 2014-01-01T00
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

missing_datetimes (*task_cls, finite_datetimes*)

task_namespace = None

Module contents

Sort of a standard library for doing stuff with Tasks at a somewhat abstract level.

Submodule introduced to stop growing util.py unstructured.

11.1.2 Submodules

luigi.cmdline module

```
luigi.cmdline.luigi_run(argv=['-b', 'latex', '-D', 'language=en', '-d', '_build/doctrees', '.',
                             '_build/latex'])
```

```
luigi.cmdline.luigid(argv=['-b', 'latex', '-D', 'language=en', '-d', '_build/doctrees', '.',
                           '_build/latex'])
```

luigi.configuration module

luigi.configuration provides some convenience wrappers around Python's ConfigParser to get configuration options from config files.

The default location for configuration files is client.cfg in the current working directory, then /etc/luigi/client.cfg.

Configuration has largely been superseded by parameters since they can do essentially everything configuration can do, plus a tighter integration with the rest of Luigi.

See [Configuration](#) for more info.

```
class luigi.configuration.LuigiConfigParser (defaults=None, dict_type=<class 'collections.OrderedDict'>, allow_no_value=False)
```

```
    Bases: ConfigParser.ConfigParser
```

```
    NO_DEFAULT = <object object>
```

```
    classmethod add_config_path (path)
```

```
    classmethod instance (*args, **kwargs)
        Singleton getter
```

```
    classmethod reload ()
```

```
    get (section, option, default=<object object>, **kwargs)
```

```
    getboolean (section, option, default=<object object>)
```

```
    getint (section, option, default=<object object>)
```

```
    getfloat (section, option, default=<object object>)
```

```
    getintdict (section)
```

```
    set (section, option, value=None)
```

```
luigi.configuration.get_config ()
```

```
    Convenience method (for backwards compatibility) for accessing config singleton.
```

luigi.date_interval module

luigi.date_interval provides convenient classes for date algebra. Everything uses ISO 8601 notation, i.e. YYYY-MM-DD for dates, etc. There is a corresponding *luigi.parameter.DateIntervalParameter* that you can use to parse date intervals.

Example:

```
class MyTask (luigi.Task):
    date_interval = luigi.DateIntervalParameter()
```

Now, you can launch this from the command line using `--date-interval 2014-05-10` or `--date-interval 2014-W26` (using week notation) or `--date-interval 2014` (for a year) and some other notations.

```
class luigi.date_interval.DateInterval (date_a, date_b)
```

```
    Bases: object
```

The *DateInterval* is the base class with subclasses *Date*, *Week*, *Month*, *Year*, and *Custom*. Note that the *DateInterval* is abstract and should not be used directly: use *Custom* for arbitrary date intervals. The base class features a couple of convenience methods, such as `next ()` which returns the next consecutive date interval.

Example:

```
x = luigi.date_interval.Week (2013, 52)
print x.prev ()
```

This will print 2014-W01.

All instances of *DateInterval* have attributes `date_a` and `date_b` set. This represents the half open range of the date interval. For instance, a May 2014 is represented as `date_a = 2014-05-01`, `date_b = 2014-06-01`.

dates()

Returns a list of dates in this date interval.

hours()

Same as `dates()` but returns 24 times more info: one for each hour.

prev()

Returns the preceding corresponding date interval (eg. May -> April).

next()

Returns the subsequent corresponding date interval (eg. 2014 -> 2015).

to_string()

classmethod from_date(d)

Abstract class method.

For instance, `Month.from_date(datetime.date(2012, 6, 6))` returns a `Month(2012, 6)`.

classmethod parse(s)

Abstract class method.

For instance, `Year.parse("2014")` returns a `Year(2014)`.

class `luigi.date_interval.Date(y, m, d)`

Bases: `luigi.date_interval.DateInterval`

Most simple *DateInterval* where `date_b == date_a + datetime.timedelta(1)`.

to_string()

classmethod from_date(d)

classmethod parse(s)

class `luigi.date_interval.Week(y, w)`

Bases: `luigi.date_interval.DateInterval`

ISO 8601 week. Note that it has some counterintuitive behavior around new year. For instance Monday 29 December 2008 is week 2009-W01, and Sunday 3 January 2010 is week 2009-W53 This example was taken from http://en.wikipedia.org/wiki/ISO_8601#Week_dates

Python datetime does not have a method to convert from ISO weeks, so the constructor uses some stupid brute force

to_string()

classmethod from_date(d)

classmethod parse(s)

class `luigi.date_interval.Month(y, m)`

Bases: `luigi.date_interval.DateInterval`

to_string()

classmethod from_date(d)

classmethod parse(s)

```
class luigi.date_interval.Year(y)
    Bases: luigi.date_interval.DateInterval

    to_string()

    classmethod from_date(d)

    classmethod parse(s)

class luigi.date_interval.Custom(date_a, date_b)
    Bases: luigi.date_interval.DateInterval

    Custom date interval (does not implement prev and next methods)

    Actually the ISO 8601 specifies <start>/<end> as the time interval format Not sure if this goes for date intervals
    as well. In any case slashes will most likely cause problems with paths etc.

    to_string()

    classmethod parse(s)
```

luigi.db_task_history module

Provides a database backend to the central scheduler. This lets you see historical runs. See [Enabling Task History](#) for information about how to turn out the task history feature.

```
class luigi.db_task_history.DbTaskHistory
    Bases: luigi.task_history.TaskHistory

    Task History that writes to a database using sqlalchemy. Also has methods for useful db queries.

    task_scheduled(task_id)

    task_finished(task_id, successful)

    task_started(task_id, worker_host)

    find_all_by_parameters(task_name, session=None, **task_params)
        Find tasks with the given task_name and the same parameters as the kwargs.

    find_all_by_name(task_name, session=None)
        Find all tasks with the given task_name.

    find_latest_runs(session=None)
        Return tasks that have been updated in the past 24 hours.

    find_all_runs(session=None)
        Return all tasks that have been updated.

    find_all_events(session=None)
        Return all running/failed/done events.

    find_task_by_id(id, session=None)
        Find task with the given record ID.

class luigi.db_task_history.TaskParameter(**kwargs)
    Bases: sqlalchemy.ext.declarative.api.Base

    Table to track luigi.Parameter(s) of a Task.

    A simple constructor that allows initialization from kwargs.

    Sets attributes on the constructed instance using the names and values in kwargs.
```


Only keys that are present as attributes of the instance's class are allowed. These could be, for example, any mapped columns or relationships.

task_id

name

value

```
class luigi.db_task_history.TaskEvent (**kwargs)
    Bases: sqlalchemy.ext.declarative.api.Base
```

Table to track when a task is scheduled, starts, finishes, and fails.

A simple constructor that allows initialization from kwargs.

Sets attributes on the constructed instance using the names and values in `kwargs`.

Only keys that are present as attributes of the instance's class are allowed. These could be, for example, any mapped columns or relationships.

id

task_id

event_name

ts

```
class luigi.db_task_history.TaskRecord (**kwargs)
    Bases: sqlalchemy.ext.declarative.api.Base
```

Base table to track information about a `luigi.Task`.

References to other tables are available through `task.events`, `task.parameters`, etc.

A simple constructor that allows initialization from kwargs.

Sets attributes on the constructed instance using the names and values in `kwargs`.

Only keys that are present as attributes of the instance's class are allowed. These could be, for example, any mapped columns or relationships.

id

name

host

parameters

events

luigi.deprecate_kwarg module

```
luigi.deprecate_kwarg.deprecate_kwarg (old_name, new_name, kw_value)
```

Rename keyword arguments, but keep backwards compatibility.

Usage:

luigi.event module

Definitions needed for events. See *Events and callbacks* for info on how to use it.

```
class luigi.event.Event
    Bases: object

    DEPENDENCY_DISCOVERED = 'event.core.dependency.discovered'
    DEPENDENCY_MISSING = 'event.core.dependency.missing'
    DEPENDENCY_PRESENT = 'event.core.dependency.present'
    BROKEN_TASK = 'event.core.task.broken'
    START = 'event.core.start'
    FAILURE = 'event.core.failure'
    SUCCESS = 'event.core.success'
    PROCESSING_TIME = 'event.core.processing_time'
```

luigi.file module

LocalTarget provides a concrete implementation of a *Target* class that uses files on the local file system

```
class luigi.file.atomic_file(path)
    Bases: luigi.target.AtomicLocalFile

    Simple class that writes to a temp file and moves it on close() Also cleans up the temp file if close is not invoked

    move_to_final_destination()
    generate_tmp_path(path)

class luigi.file.LocalFileSystem
    Bases: luigi.target.FileSystem

    Wrapper for access to file system operations.

    Work in progress - add things as needed.

    exists(path)
    mkdir(path, parents=True, raise_if_exists=False)
    isdir(path)
    listdir(path)
    remove(path, recursive=True)

class luigi.file.LocalTarget(path=None, format=None, is_tmp=False)
    Bases: luigi.target.FileSystemTarget

    fs = <luigi.file.LocalFileSystem object>

    makedirs()
        Create all parent folders if they do not exist.

    open(mode='r')
    move(new_path, raise_if_exists=False)
    move_dir(new_path)
```

```

    remove()
    copy(new_path, raise_if_exists=False)
    fn
class luigi.file.File(*args, **kwargs)
    Bases: luigi.file.LocalTarget

```

luigi.format module

```

class luigi.format.FileWrapper(file_object)
    Bases: object

    Wrap file in a “real” so stuff can be added to it after creation.

class luigi.format.InputPipeProcessWrapper(command, input_pipe=None)
    Bases: object

    Initializes a InputPipeProcessWrapper instance.

    Parameters command – a subprocess.Popen instance with stdin=input_pipe and std-
        out=subprocess.PIPE. Alternatively, just its args argument as a convenience.

    create_subprocess(command)
        http://www.chiark.greenend.org.uk/ucgi/~cjwatson/blosxom/2009-07-02-python-sigpipe.html

    close()

    readable()

    writable()

    seekable()

class luigi.format.OutputPipeProcessWrapper(command, output_pipe=None)
    Bases: object

    WRITES_BEFORE_FLUSH = 10000

    write(*args, **kwargs)

    writeLine(line)

    close()

    abort()

    readable()

    writable()

    seekable()

class luigi.format.BaseWrapper(stream, *args, **kwargs)
    Bases: object

class luigi.format.NewlineWrapper(stream, newline=None)
    Bases: luigi.format.BaseWrapper

    read(n=-1)

    writelines(lines)

    write(b)

```

```
class luigi.format.MixedUnicodeBytesWrapper (stream, encoding=None)
    Bases: luigi.format.BaseWrapper

    write (b)

    writelines (lines)

class luigi.format.Format
    Bases: object

    Interface for format specifications.

    classmethod pipe_reader (input_pipe)

    classmethod pipe_writer (output_pipe)

class luigi.format.ChainFormat (*args, **kwargs)
    Bases: luigi.format.Format

    pipe_reader (input_pipe)

    pipe_writer (output_pipe)

class luigi.format.TextWrapper (stream, *args, **kwargs)
    Bases: _io.TextIOWrapper

class luigi.format.NopFormat
    Bases: luigi.format.Format

    pipe_reader (input_pipe)

    pipe_writer (output_pipe)

class luigi.format.WrappedFormat (*args, **kwargs)
    Bases: luigi.format.Format

    pipe_reader (input_pipe)

    pipe_writer (output_pipe)

class luigi.format.TextFormat (*args, **kwargs)
    Bases: luigi.format.WrappedFormat

    input = 'unicode'

    output = 'bytes'

    wrapper_cls
        alias of TextWrapper

class luigi.format.MixedUnicodeBytesFormat (*args, **kwargs)
    Bases: luigi.format.WrappedFormat

    output = 'bytes'

    wrapper_cls
        alias of MixedUnicodeBytesWrapper

class luigi.format.NewlineFormat (*args, **kwargs)
    Bases: luigi.format.WrappedFormat

    input = 'bytes'

    output = 'bytes'

    wrapper_cls
        alias of NewlineWrapper
```

```

class luigi.format.GzipFormat (compression_level=None)
    Bases: luigi.format.Format

    input = 'bytes'
    output = 'bytes'
    pipe_reader (input_pipe)
    pipe_writer (output_pipe)
class luigi.format.Bzip2Format
    Bases: luigi.format.Format

    input = 'bytes'
    output = 'bytes'
    pipe_reader (input_pipe)
    pipe_writer (output_pipe)
luigi.format.get_default_format ()

```

luigi.hadoop module

luigi.hadoop has moved to *luigi.contrib.hadoop*

luigi.hadoop_jar module

luigi.hadoop_jar has moved to *luigi.contrib.hadoop_jar*

luigi.hdfs module

luigi.hdfs has moved to *luigi.contrib.hdfs*

luigi.hive module

The hive module has been moved to *luigi.contrib.hive*

luigi.interface module

This module contains the bindings for command line integration and dynamic loading of tasks

```
luigi.interface.setup_interface_logging (conf_file=None)
```

```

class luigi.interface.core (*args, **kwargs)
    Bases: luigi.task.Config

```

Keeps track of a bunch of environment params.

Uses the internal luigi parameter mechanism. The nice thing is that we can instantiate this class and get an object with all the environment variables set. This is arguably a bit of a hack.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
use_cmdline_section = False
local_scheduler = <luigi.parameter.BoolParameter object>
scheduler_host = <luigi.parameter.Parameter object>
scheduler_port = <luigi.parameter.IntParameter object>
lock_size = <luigi.parameter.IntParameter object>
no_lock = <luigi.parameter.BoolParameter object>
lock_pid_dir = <luigi.parameter.Parameter object>
workers = <luigi.parameter.IntParameter object>
logging_conf_file = <luigi.parameter.Parameter object>
module = <luigi.parameter.Parameter object>
parallel_scheduling = <luigi.parameter.BoolParameter object>
assistant = <luigi.parameter.BoolParameter object>
task_namespace = None
```

```
class luigi.interface.WorkerSchedulerFactory
    Bases: object

    create_local_scheduler()
    create_remote_scheduler(host, port)
    create_worker(scheduler, worker_processes, assistant=False)
```

```
class luigi.interface.Interface
    Bases: object

    parse()

    static run(tasks, worker_scheduler_factory=None, override_defaults=None)
```

Parameters

- **tasks** –
- **worker_scheduler_factory** –
- **override_defaults** –

Returns True if all tasks and their dependencies were successfully run (or already completed); False if any error occurred.

```
luigi.interface.error_task_names(task_name, task_names)
luigi.interface.add_task_parameters(parser, task_cls, optparse=False)
luigi.interface.get_global_parameters()
luigi.interface.add_global_parameters(parser, optparse=False)
luigi.interface.get_task_parameters(task_cls, args)
luigi.interface.set_global_parameters(args)
```

class `luigi.interface.ArgParseInterface`

Bases: `luigi.interface.Interface`

Takes the task as the command, with parameters specific to it.

parse_task (*cmdline_args=None, main_task_cls=None*)

parse (*cmdline_args=None, main_task_cls=None*)

class `luigi.interface.DynamicArgParseInterface`

Bases: `luigi.interface.ArgParseInterface`

Uses `--module` as a way to load modules dynamically

Usage:

```
python whatever.py --module foo_module FooTask --blah xyz --x 123
```

This will dynamically import `foo_module` and then try to create `FooTask` from this.

parse (*cmdline_args=None, main_task_cls=None*)

class `luigi.interface.PassThroughOptionParser` (*usage=None, option_list=None, option_class=<class optparse.Option>, version=None, conflict_handler='error', description=None, formatter=None, add_help_option=True, prog=None, epilog=None*)

Bases: `optparse.OptionParser`

An unknown option pass-through implementation of `OptionParser`.

When unknown arguments are encountered, bundle with `largs` and try again, until `rargs` is depleted.

`sys.exit(status)` will still be called if a known argument is passed incorrectly (e.g. missing arguments or bad argument types, etc.)

class `luigi.interface.OptParseInterface` (*existing_optparse*)

Bases: `luigi.interface.Interface`

Supported for legacy reasons where it's necessary to interact with an existing parser.

Takes the task using `--task`. All parameters to all possible tasks will be defined globally in a big unordered soup.

parse (*cmdline_args=None, main_task_cls=None*)

luigi.interface.run (*cmdline_args=None, existing_optparse=None, use_optparse=False, main_task_cls=None, worker_scheduler_factory=None, use_dynamic_argparse=False, local_scheduler=False*)

Run from cmdline.

The default parser uses `argparse` however, for legacy reasons, we support `optparse` that optionally allows for overriding an existing option parser with new args.

Parameters

- **cmdline_args** –
- **existing_optparse** –
- **use_optparse** –
- **main_task_cls** –
- **worker_scheduler_factory** –

- `use_dynamic_argparse` –
- `local_scheduler` –

`luigi.interface.build(tasks, worker_scheduler_factory=None, **env_params)`

Run internally, bypassing the cmdline parsing.

Useful if you have some luigi code that you want to run internally. Example:

```
luigi.build([MyTask1(), MyTask2()], local_scheduler=True)
```

One notable difference is that *build* defaults to not using the identical process lock. Otherwise, *build* would only be callable once from each process.

Parameters

- `tasks` –
- `worker_scheduler_factory` –
- `env_params` –

Returns

luigi.lock module

Locking functionality when launching things from the command line. Uses a pidfile. This prevents multiple identical workflows to be launched simultaneously.

`luigi.lock.getpcmd(pid)`

Returns command of process.

Parameters `pid` –

`luigi.lock.get_info(pid_dir, my_pid=None)`

`luigi.lock.acquire_for(pid_dir, num_available=1)`

Makes sure the process is only run once at the same time with the same name.

Notice that we since we check the process name, different parameters to the same command can spawn multiple processes at the same time, i.e. running “/usr/bin/my_process” does not prevent anyone from launching “/usr/bin/my_process –foo bar”.

luigi.mock module

This module provides a class *MockTarget*, an implementation of *Target*. *MockTarget* contains all data in-memory. The main purpose is unit testing workflows without writing to disk.

class `luigi.mock.MockFileSystem`

Bases: `luigi.target.FileSystem`

MockFileSystem inspects/modifies `_data` to simulate file system operations.

get_all_data()

get_data(*fn*)

exists(*path*)

remove(*path*, *recursive=True*, *skip_trash=True*)

Removes the given mockfile. `skip_trash` doesn’t have any meaning.


```

listdir (path)
    listdir does a prefix match of self.get_all_data(), but doesn't yet support globs.

isdir (path)

mkdir (path, parents=True, raise_if_exists=False)
    mkdir is a noop.

clear ()

class luigi.mock.MockTarget (fn, is_tmp=None, mirror_on_stderr=False, format=None)
    Bases: luigi.target.FileSystemTarget

    fs = <luigi.mock.MockFileSystem object>

    exists ()

    rename (path, raise_if_exists=False)

    path

    open (mode)

class luigi.mock.MockFile (*args, **kwargs)
    Bases: luigi.mock.MockTarget

```

luigi.mrrunner module

The hadoop runner.

This module contains the main() method which will be used to run the mapper and reducer on the Hadoop nodes.

```

class luigi.mrrunner.Runner (job=None)
    Bases: object

    Run the mapper or reducer on hadoop nodes.

    extract_packages_archive ()

    run (kind, stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>)

luigi.mrrunner.main (args=None, stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<std-
                                out>', mode 'w'>, print_exception=<function print_exception>)
    Run either the mapper or the reducer from the class instance in the file "job-instance.pickle".

    Arguments:

    kind – is either map or reduce

luigi.mrrunner.print_exception (exc)

```

luigi.notifications module

Supports sending emails when tasks fail.

This needs some more documentation. See [Configuration](#) for configuration options. In particular using the config *error-email* should set up Luigi so that it will send emails when tasks fail.

```

[core]
error-email: foo@bar.baz

```

```

luigi.notifications.email_type ()

luigi.notifications.generate_email (sender, subject, message, recipients, image_png)

```

```
luigi.notifications.wrap_traceback (traceback)
luigi.notifications.send_email_smtp (config, sender, subject, message, recipients, image_png)
luigi.notifications.send_email_ses (config, sender, subject, message, recipients, image_png)
luigi.notifications.send_email_sendgrid (config, sender, subject, message, recipients, image_png)
luigi.notifications.send_email (subject, message, sender, recipients, image_png=None)
luigi.notifications.send_error_email (subject, message)
    Sends an email to the configured error-email.
    If no error-email is configured, then a message is logged.
```

luigi.parameter module

Parameters are one of the core concepts of Luigi. All Parameters sit on *Task* classes. See *Parameter* for more info on how to define parameters.

exception `luigi.parameter.ParameterException`

Bases: `exceptions.Exception`

Base exception.

exception `luigi.parameter.MissingParameterException`

Bases: `luigi.parameter.ParameterException`

Exception signifying that there was a missing Parameter.

exception `luigi.parameter.UnknownParameterException`

Bases: `luigi.parameter.ParameterException`

Exception signifying that an unknown Parameter was supplied.

exception `luigi.parameter.DuplicateParameterException`

Bases: `luigi.parameter.ParameterException`

Exception signifying that a Parameter was specified multiple times.

exception `luigi.parameter.UnknownConfigException`

Bases: `luigi.parameter.ParameterException`

Exception signifying that the `config_path` for the Parameter could not be found.

class `luigi.parameter.Parameter` (**args, **kwargs*)

Bases: `object`

An untyped Parameter

Parameters are objects set on the Task class level to make it possible to parameterize tasks. For instance:

```
class MyTask(luigi.Task): foo = luigi.Parameter()
```

This makes it possible to instantiate multiple tasks, eg `MyTask(foo='bar')` and `MyTask(foo='baz')`. The task will then have the `foo` attribute set appropriately.

There are subclasses of `Parameter` that define what type the parameter has. This is not enforced within Python, but are used for command line interaction.

The `config_path` argument lets you specify a place where the parameter is read from config in case no value is provided.

When a task is instantiated, it will first use any argument as the value of the parameter, eg. if you instantiate `a = TaskA(x=44)` then `a.x == 44`. If this does not exist, it will use the value of the Parameter object, which is defined on a class level. This will be resolved in this order of falling priority:

- Any value provided on the command line on the class level (eg. `--TaskA-param xyz`)
- Any value provided via config (using the `config_path` argument)
- Any default value set using the `default` flag.

counter = 67

non-atomically increasing counter used for ordering parameters.

has_value

True if a default was specified or if `config_path` references a valid entry in the conf.

Note that “value” refers to the Parameter object itself - it can be either

- 1.The default value for this parameter
- 2.A value read from the config
- 3.A global value

Any Task instance can have its own value set that overrides this.

value

The value for this Parameter.

This refers to any value defined by a default, a config option, or a global value.

Raises `MissingParameterException` if a value is not set.

Returns the parsed value.

has_task_value (*task_name*, *param_name*)

task_value (*task_name*, *param_name*)

set_global (*value*)

Set the global value of this Parameter.

Parameters **value** – the new global value.

reset_global ()

parse (*x*)

Parse an individual value from the input.

The default implementation is an identify (it returns *x*), but subclasses should override this method for specialized parsing. This method is called by `parse_from_input()` if *x* exists. If this Parameter was specified with `is_list=True`, then `parse` is called once for each item in the list.

Parameters **x** (*str*) – the value to parse.

Returns the parsed value.

serialize (*x*)

Opposite of `parse()`.

Converts the value *x* to a string.

Parameters **x** – the value to serialize.

parse_from_input (*param_name*, *x*, *task_name=None*)

Parses the parameter value from input *x*, handling defaults and `is_list`.

Parameters

- **param_name** – the name of the parameter. This is used for the message in `MissingParameterException`.
- **x** – the input value to parse.

Raises `MissingParameterException` if `x` is false-y and no default is specified.

serialize_to_input (*x*)

parser_dest (*param_name, task_name, glob=False, is_without_section=False*)

add_to_cmdline_parser (*parser, param_name, task_name, optparse=False, glob=False, is_without_section=False*)

parse_from_args (*param_name, task_name, args, params*)

set_global_from_args (*param_name, task_name, args, is_without_section=False*)

class `luigi.parameter.DateHourParameter` (**args, **kwargs*)

Bases: `luigi.parameter.Parameter`

Parameter whose value is a datetime specified to the hour.

A `DateHourParameter` is a [ISO 8601](#) formatted date and time specified to the hour. For example, `2013-07-10T19` specifies July 10, 2013 at 19:00.

date_format = `'%Y-%m-%dT%H'`

parse (*s*)

Parses a string to a datetime using the format string `%Y-%m-%dT%H`.

serialize (*dt*)

Converts the datetime to a string usnig the format string `%Y-%m-%dT%H`.

class `luigi.parameter.DateMinuteParameter` (**args, **kwargs*)

Bases: `luigi.parameter.DateHourParameter`

Parameter whose value is a datetime specified to the minute.

A `DateMinuteParameter` is a [ISO 8601](#) formatted date and time specified to the minute. For example, `2013-07-10T19H07` specifies July 10, 2013 at 19:07.

date_format = `'%Y-%m-%dT%HH%M'`

class `luigi.parameter.DateParameter` (**args, **kwargs*)

Bases: `luigi.parameter.Parameter`

Parameter whose value is a date.

A `DateParameter` is a Date string formatted `YYYY-MM-DD`. For example, `2013-07-10` specifies July 10, 2013.

parse (*s*)

Parses a date string formatted as `YYYY-MM-DD`.

class `luigi.parameter.IntParameter` (**args, **kwargs*)

Bases: `luigi.parameter.Parameter`

Parameter whose value is an int.

parse (*s*)

Parses an int from the string using `int()`.

class `luigi.parameter.FloatParameter` (**args, **kwargs*)

Bases: `luigi.parameter.Parameter`

Parameter whose value is a float.

parse (s)

Parses a float from the string using `float()`.

class `luigi.parameter.BoolParameter` (*args, **kwargs)

Bases: `luigi.parameter.Parameter`

A Parameter whose value is a `bool`.

This constructor passes along args and kwargs to ctor for `Parameter` but specifies `is_bool=True`.

parse (s)

Parses a `bool` from the string, matching ‘true’ or ‘false’ ignoring case.

class `luigi.parameter.BooleanParameter` (*args, **kwargs)

Bases: `luigi.parameter.BoolParameter`

class `luigi.parameter.DateIntervalParameter` (*args, **kwargs)

Bases: `luigi.parameter.Parameter`

A Parameter whose value is a `DateInterval`.

Date Intervals are specified using the ISO 8601 `Time Interval` notation.

parse (s)

Parses a `:py:class:~luigi.date_interval.DateInterval` from the input.

see `luigi.date_interval` for details on the parsing of DateIntervals.

class `luigi.parameter.TimeDeltaParameter` (*args, **kwargs)

Bases: `luigi.parameter.Parameter`

Class that maps to `timedelta` using strings in any of the following forms:

- `n {w[EEK[S]]|d[AY[S]]|h[OUR[S]]|m[INUTE[S]]|s[SECOND[S]]}` (e.g. “1 week 2 days” or “1 h”)

Note: multiple arguments must be supplied in longest to shortest unit order
- ISO 8601 duration `PnDTnHnMnS` (each field optional, years and months not supported)
- ISO 8601 duration `PnW`

See https://en.wikipedia.org/wiki/ISO_8601#Durations

parse (input)

Parses a time delta from the input.

See `TimeDeltaParameter` for details on supported formats.

luigi.postgres module

Implements a subclass of `Target` that writes data to Postgres. Also provides a helper task to copy data into a Postgres table.

class `luigi.postgres.MultiReplacer` (replace_pairs)

Bases: `object`

Object for one-pass replace of multiple words

Substituted parts will not be matched against other replace patterns, as opposed to when using multipass replace. The order of the items in the `replace_pairs` input will dictate replacement precedence.

Constructor arguments: `replace_pairs` – list of 2-tuples which hold strings to be replaced and replace string

Usage:

```
>>> replace_pairs = [("a", "b"), ("b", "c")]
>>> MultiReplacer(replace_pairs) ("abcd")
'bccd'
>>> replace_pairs = [("ab", "x"), ("a", "x")]
>>> MultiReplacer(replace_pairs) ("ab")
'x'
>>> replace_pairs.reverse()
>>> MultiReplacer(replace_pairs) ("ab")
'xb'
```

Initializes a MultiReplacer instance.

Parameters **replace_pairs** (*tuple*) – list of 2-tuples which hold strings to be replaced and replace string.

class `luigi.postgres.PostgresTarget` (*host, database, user, password, table, update_id*)

Bases: `luigi.target.Target`

Target for a resource in Postgres.

This will rarely have to be directly instantiated by the user.

Args: *host* (str): Postgres server address. Possibly a host:port string. *database* (str): Database name *user* (str): Database user *password* (str): Password for specified user *update_id* (str): An identifier for this data set

marker_table = 'table_updates'

use_db_timestamps = True

touch (*connection=None*)

Mark this update as complete.

Important: If the marker table doesn't exist, the connection transaction will be aborted and the connection reset. Then the marker table will be created.

exists (*connection=None*)

connect ()

Get a psycopg2 connection object to the database where the table is.

create_marker_table ()

Create marker table if it doesn't exist.

Using a separate connection since the transaction might have to be reset.

open (*mode*)

class `luigi.postgres.CopyToTable` (**args, **kwargs*)

Bases: `luigi.contrib.rdbms.CopyToTable`

Template task for inserting a data set into Postgres

Usage: Subclass and override the required *host, database, user, password, table* and *columns* attributes.

To customize how to access data from an input task, override the *rows* method with a generator that yields each row as a tuple with fields ordered according to *columns*.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

rows ()
Return/yield tuples or lists corresponding to each row to be inserted.

map_column (*value*)
Applied to each column of every row returned by *rows*.
Default behaviour is to escape special characters and identify any self.null_values.

output ()
Returns a PostgresTarget representing the inserted dataset.
Normally you don't override this.

copy (*cursor*, *file*)

run ()
Inserts data generated by rows() into target table.
If the target table doesn't exist, self.create_table will be called to attempt to create the table.
Normally you don't want to override this.

task_namespace = None

luigi.process module

Contains some helper functions to run luigid in daemon mode

```
luigi.process.check_pid(pidfile)
luigi.process.write_pid(pidfile)
luigi.process.get_log_format()
luigi.process.get_spool_handler(filename)
luigi.process.daemonize(cmd, pidfile=None, logdir=None, api_port=8082, address=None)
```

luigi.rpc module

Implementation of the REST interface between the workers and the server. rpc.py implements the client side of it, server.py implements the server side. See [Using the Central Scheduler](#) for more info.

```
exception luigi.rpc.RPCError(message, sub_exception=None)
    Bases: exceptions.Exception

class luigi.rpc.RemoteScheduler(host='localhost', port=8082, connect_timeout=None,
                                url_prefix='')
    Bases: luigi.scheduler.Scheduler

    Scheduler proxy object. Talks to a RemoteSchedulerResponder.

    ping(worker)

    add_task(worker, task_id, status='PENDING', runnable=True, deps=None, new_deps=None,
             expl=None, resources=None, priority=0, family='', module=None, params=None, assis-
             tant=False)

    get_work(worker, host=None, assistant=False)

    graph()

    dep_graph(task_id)
```

```
inverse_dep_graph (task_id)  
task_list (status, upstream_status, search=None)  
worker_list ()  
task_search (task_str)  
fetch_error (task_id)  
add_worker (worker, info)  
update_resources (**resources)  
prune ()  
re_enable_task (task_id)
```

luigi.s3 module

Implementation of Simple Storage Service support. *S3Target* is a subclass of the *Target* class to support S3 file system operations

```
exception luigi.s3.InvalidDeleteException  
    Bases: luigi.target.FileSystemException
```

```
exception luigi.s3.FileNotFoundException  
    Bases: luigi.target.FileSystemException
```

```
class luigi.s3.S3Client (aws_access_key_id=None, aws_secret_access_key=None, **kwargs)  
    Bases: luigi.target.FileSystem
```

boto-powered S3 client.

```
exists (path)  
    Does provided path exist on S3?
```

```
remove (path, recursive=True)  
    Remove a file or directory from S3.
```

```
get_key (path)
```

```
put (local_path, destination_s3_path)  
    Put an object stored locally to an S3 path.
```

```
put_string (content, destination_s3_path)  
    Put a string to an S3 path.
```

```
put_multipart (local_path, destination_s3_path, part_size=67108864)  
    Put an object stored locally to an S3 path using S3 multi-part upload (for files > 5GB).
```

Parameters

- **local_path** – Path to source local file
- **destination_s3_path** – URL for target S3 location
- **part_size** – Part size in bytes. Default: 67108864 (64MB), must be >= 5MB and <= 5 GB.

```
copy (source_path, destination_path)  
    Copy an object from one S3 location to another.
```

```
rename (source_path, destination_path)  
    Rename/move an object from one S3 location to another.
```


listdir (*path*)

Get an iterable with S3 folder contents. Iterable contains paths relative to queried path.

list (*path*)

isdir (*path*)

Is the parameter S3 path a directory?

is_dir (*path*)

Is the parameter S3 path a directory?

mkdir (*path*, *parents=True*, *raise_if_exists=False*)

class `luigi.s3.AtomicS3File` (*path*, *s3_client*)

Bases: `luigi.target.AtomicLocalFile`

An S3 file that writes to a temp file and put to S3 on close.

move_to_final_destination ()

class `luigi.s3.ReadableS3File` (*s3_key*)

Bases: `object`

read (*size=0*)

close ()

readable ()

writable ()

seekable ()

class `luigi.s3.S3Target` (*path*, *format=None*, *client=None*)

Bases: `luigi.target.FileSystemTarget`

fs = `None`

open (*mode='r'*)

class `luigi.s3.S3FlagTarget` (*path*, *format=None*, *client=None*, *flag='_SUCCESS'*)

Bases: `luigi.s3.S3Target`

Defines a target directory with a flag-file (defaults to `_SUCCESS`) used to signify job success.

This checks for two things:

- the path exists (just like the `S3Target`)
- the `_SUCCESS` file exists within the directory.

Because Hadoop outputs into a directory and not a single file, the path is assumed to be a directory.

This is meant to be a handy alternative to `AtomicS3File`.

The `AtomicFile` approach can be burdensome for S3 since there are no directories, per se.

If we have 1,000,000 output files, then we have to rename 1,000,000 objects.

Initializes a `S3FlagTarget`.

Parameters

- **path** (*str*) – the directory where the files are stored.
- **client** –
- **flag** (*str*) –

fs = None

exists()

class `luigi.s3.S3EmrTarget(*args, **kwargs)`

Bases: `luigi.s3.S3FlagTarget`

Deprecated. Use `S3FlagTarget`

class `luigi.s3.S3PathTask(*args, **kwargs)`

Bases: `luigi.task.ExternalTask`

A external task that to require existence of a path in S3.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

path = <luigi.parameter.Parameter object>

output()

task_namespace = None

class `luigi.s3.S3EmrTask(*args, **kwargs)`

Bases: `luigi.task.ExternalTask`

An external task that requires the existence of EMR output in S3.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

path = <luigi.parameter.Parameter object>

output()

task_namespace = None

class `luigi.s3.S3FlagTask(*args, **kwargs)`

Bases: `luigi.task.ExternalTask`

An external task that requires the existence of EMR output in S3.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

task_namespace = None

path = <luigi.parameter.Parameter object>

flag = <luigi.parameter.Parameter object>

```
output ()
```

luigi.scalding module

luigi.scalding has moved to luigi.contrib.scalding

luigi.scheduler module

The system for scheduling tasks and executing them in order. Deals with dependencies, priorities, resources, etc. The *Worker* pulls tasks from the scheduler (usually over the REST interface) and executes them. See [Using the Central Scheduler](#) for more info.

```
class luigi.scheduler.Scheduler
```

Bases: object

Abstract base class.

Note that the methods all take string arguments, not Task objects...

```
add_task = NotImplemented
```

```
get_work = NotImplemented
```

```
ping = NotImplemented
```

```
luigi.scheduler.UPSTREAM_SEVERITY_KEY ()
```

T.index(value, [start, [stop]]) -> integer – return first index of value. Raises ValueError if the value is not present.

```
class luigi.scheduler.scheduler (*args, **kwargs)
```

Bases: *luigi.task.Config*

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
retry_delay = <luigi.parameter.FloatParameter object>
```

```
remove_delay = <luigi.parameter.FloatParameter object>
```

```
worker_disconnect_delay = <luigi.parameter.FloatParameter object>
```

```
state_path = <luigi.parameter.Parameter object>
```

```
disable_window = <luigi.parameter.IntParameter object>
```

```
disable_failures = <luigi.parameter.IntParameter object>
```

```
disable_hard_timeout = <luigi.parameter.IntParameter object>
```

```
disable_persist = <luigi.parameter.IntParameter object>
```

```
max_shown_tasks = <luigi.parameter.IntParameter object>
```

```
prune_done_tasks = <luigi.parameter.BoolParameter object>
```

```
record_task_history = <luigi.parameter.BoolParameter object>
```

```
visualization_graph = <luigi.parameter.Parameter object>
```

```
task_namespace = None
```

`luigi.scheduler.fix_time(x)`

class `luigi.scheduler.Failures(window)`

Bases: object

This class tracks the number of failures in a given time window.

Failures added are marked with the current timestamp, and this class counts the number of failures in a sliding time window ending at the present.

Initialize with the given window.

Parameters **window** – how long to track failures for, as a float (number of seconds).

add_failure()

Add a failure event with the current timestamp.

num_failures()

Return the number of failures in the window.

clear()

Clear the failure queue.

class `luigi.scheduler.Task(task_id, status, deps, resources=None, priority=0, family='', module=None, params=None, disable_failures=None, disable_window=None, disable_hard_timeout=None)`

Bases: object

add_failure()

has_excessive_failures()

can_disable()

class `luigi.scheduler.Worker(worker_id, last_active=None)`

Bases: object

Structure for tracking worker activity and keeping their references.

add_info(info)

update(worker_reference)

prune(config)

get_pending_tasks()

is_trivial_worker()

If it's not an assistant having only tasks that are without requirements

assistant

class `luigi.scheduler.SimpleTaskState(state_path)`

Bases: object

Keep track of the current state and handle persistence.

The point of this class is to enable other ways to keep state, eg. by using a database These will be implemented by creating an abstract base class that this and other classes inherit from.

dump()

load()

get_active_tasks(status=None)

get_running_tasks()

```

get_pending_tasks ()
get_task (task_id, default=None, setdefault=None)
has_task (task_id)
re_enable (task, config=None)
set_status (task, new_status, config=None)
prune (task, config, assistants)
inactivate_tasks (delete_tasks)
get_active_workers (last_active_lt=None)
get_assistants (last_active_lt=None)
get_worker_ids ()
get_worker (worker_id)
inactivate_workers (delete_workers)
get_necessary_tasks ()
class luigi.scheduler.CentralPlannerScheduler (config=None, resources=None,
                                              task_history_impl=None, **kwargs)
    Bases: luigi.scheduler.Scheduler
    Async scheduler that can handle multiple workers, etc.
    Can be run locally or on a server (using RemoteScheduler + server.Server).
    Keyword Arguments: :param config: an object of class “scheduler” or None (in which the global instance will
    be used) :param resources: a dict of str->int constraints :param task_history_override: ignore config and use this
    object as the task history
    load ()
    dump ()
    prune ()
    update (worker_id, worker_reference=None)
        Keep track of whenever the worker was last active.
    add_task (task_id=None, status='PENDING', runnable=True, deps=None, new_deps=None,
              expl=None, resources=None, priority=0, family='', module=None, params=None,
              assistant=False, **kwargs)
        •add task identified by task_id if it doesn't exist
        •if deps is not None, update dependency list
        •update status of task
        •add additional workers/stakeholders
        •update priority when needed
    add_worker (worker, info, **kwargs)
    update_resources (**resources)
    get_work (host=None, assistant=False, **kwargs)
    ping (**kwargs)
    graph (**kwargs)

```

dep_graph (*task_id*, ***kwargs*)

task_list (*status*, *upstream_status*, *limit=True*, *search=None*, ***kwargs*)
Query for a subset of tasks by status.

worker_list (*include_running=True*, ***kwargs*)

inverse_dep_graph (*task_id*, ***kwargs*)

task_search (*task_str*, ***kwargs*)
Query for a subset of tasks by task_id.

Parameters **task_str** –

Returns

re_enable_task (*task_id*)

fetch_error (*task_id*, ***kwargs*)

task_history

luigi.server module

Simple REST server that takes commands in a JSON payload Interface to the *CentralPlannerScheduler* class. See [Using the Central Scheduler](#) for more info.

```
class luigi.server.RPCHandler(application, request, **kwargs)
    Bases: tornado.web.RequestHandler

    Handle remote scheduling calls using rpc.RemoteSchedulerResponder.

    initialize(scheduler)

    get(method)

    post(method)

class luigi.server.BaseTaskHistoryHandler(application, request, **kwargs)
    Bases: tornado.web.RequestHandler

    initialize(scheduler)

    get_template_path()

class luigi.server.AllRunHandler(application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler

    get()

class luigi.server.SelectedRunHandler(application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler

    get(name)

luigi.server.from_utc(utcTime, fmt=None)
    convert UTC time string to time.struct_time: change datetime.datetime to time, return time.struct_time type

class luigi.server.RecentRunHandler(application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler

    get()

class luigi.server.ByNameHandler(application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler
```

```

    get (name)
class luigi.server.ByIdHandler (application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler
    get (id)
class luigi.server.ByParamsHandler (application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler
    get (name)
class luigi.server.StaticFileHandler (application, request, **kwargs)
    Bases: tornado.web.RequestHandler
    get (path)
class luigi.server.RootPathHandler (application, request, **kwargs)
    Bases: luigi.server.BaseTaskHistoryHandler
    get ()
luigi.server.app (scheduler)
luigi.server.run (api_port=8082, address=None, scheduler=None, responder=None)
    Runs one instance of the API server.
luigi.server.stop ()

```

luigi.six module

luigi.target module

The abstract *Target* class. It is a central concept of Luigi and represents the state of the workflow.

```
class luigi.target.Target
```

Bases: object

A Target is a resource generated by a *Task*.

For example, a Target might correspond to a file in HDFS or data in a database. The Target interface defines one method that must be overridden: *exists()*, which signifies if the Target has been created or not.

Typically, a *Task* will define one or more Targets as output, and the Task is considered complete if and only if each of its output Targets exist.

```
exists ()
```

Returns True if the *Target* exists and False otherwise.

```
exception luigi.target.FileSystemException
```

Bases: exceptions.Exception

Base class for generic file system exceptions.

```
exception luigi.target.FileAlreadyExists
```

Bases: luigi.target.FileSystemException

Raised when a file system operation can't be performed because a directory exists but is required to not exist.

```
exception luigi.target.MissingParentDirectory
```

Bases: luigi.target.FileSystemException

Raised when a parent directory doesn't exist. (Imagine mkdir without -p)

exception `luigi.target.NotADirectory`Bases: `luigi.target.FileSystemException`

Raised when a file system operation can't be performed because an expected directory is actually a file.

class `luigi.target.FileSystem`Bases: `object`FileSystem abstraction used in conjunction with `FileSystemTarget`.

Typically, a `FileSystem` is associated with instances of a `FileSystemTarget`. The instances of the `py:class:FileSystemTarget` will delegate methods such as `FileSystemTarget.exists()` and `FileSystemTarget.remove()` to the `FileSystem`.

Methods of `FileSystem` raise `FileSystemException` if there is a problem completing the operation.**exists** (*path*)Return `True` if file or directory at *path* exist, `False` otherwise**Parameters** *path* (*str*) – a path within the `FileSystem` to check for existence.**remove** (*path*, *recursive=True*, *skip_trash=True*)Remove file or directory at location *path***Parameters**

- **path** (*str*) – a path within the `FileSystem` to remove.
- **recursive** (*bool*) – if the path is a directory, recursively remove the directory and all of its descendants. Defaults to `True`.

makedirs (*path*, *parents=True*, *raise_if_exists=False*)Create directory at location *path*Creates the directory at *path* and implicitly create parent directories if they do not already exist.**Parameters**

- **path** (*str*) – a path within the `FileSystem` to create as a directory.
- **parents** (*bool*) – Create parent directories when necessary. When *parents=False* and the parent directory doesn't exist, raise `luigi.target.MissingParentDirectory`
- **raise_if_exists** (*bool*) – raise `luigi.target.FileAlreadyExists` if the folder already exists.

Note: This method is optional, not all `FileSystem` subclasses implements it.*Note:* **parents** and **raise_if_exists** were added in August 2014. Some implementations might not support these flags yet.**isdir** (*path*)Return `True` if the location at *path* is a directory. If not, return `False`.**Parameters** *path* (*str*) – a path within the `FileSystem` to check as a directory.*Note:* This method is optional, not all `FileSystem` subclasses implements it.**listdir** (*path*)Return a list of files rooted in *path*.This returns an iterable of the files rooted at *path*. This is intended to be a recursive listing.**Parameters** *path* (*str*) – a path within the `FileSystem` to list.*Note:* This method is optional, not all `FileSystem` subclasses implements it.

class `luigi.target.FileSystemTarget` (*path*)

Bases: `luigi.target.Target`

Base class for FileSystem Targets like `LocalTarget` and `HdfsTarget`.

A `FileSystemTarget` has an associated `FileSystem` to which certain operations can be delegated. By default, `exists()` and `remove()` are delegated to the `FileSystem`, which is determined by the `fs()` property.

Methods of `FileSystemTarget` raise `FileSystemException` if there is a problem completing the operation.

Initializes a `FileSystemTarget` instance.

Parameters `path` (*str*) – the path associated with this `FileSystemTarget`.

fs

The `FileSystem` associated with this `FileSystemTarget`.

open (*mode*)

Open the `FileSystem` target.

This method returns a file-like object which can either be read from or written to depending on the specified mode.

Parameters `mode` (*str*) – the mode *r* opens the `FileSystemTarget` in read-only mode, whereas *w* will open the `FileSystemTarget` in write mode. Subclasses can implement additional options.

exists ()

Returns `True` if the path for this `FileSystemTarget` exists; `False` otherwise.

This method is implemented by using `fs()`.

remove ()

Remove the resource at the path specified by this `FileSystemTarget`.

This method is implemented by using `fs()`.

class `luigi.target.AtomicLocalFile` (*path*)

Bases: `_io.BufferedWriter`

Abstract class to create Target that create a temporary file in the local filesystem before moving it to there final destination

This class is just for the writing part of the Target. See `luigi.file.LocalTarget` for example

close ()

generate_tmp_path (*path*)

move_to_final_destination ()

tmp_path

luigi.task module

The abstract `Task` class. It is a central concept of Luigi and represents the state of the workflow. See [Tasks](#) for an overview.

`luigi.task.namespace` (*namespace=None*)

Call to set namespace of tasks declared after the call.

If called without arguments or with `None` as the namespace, the namespace is reset, which is recommended to do at the end of any file where the namespace is set to avoid unintentionally setting namespace on tasks outside of the scope of the current file.

The namespace of a Task can also be changed by specifying the property `task_namespace`. This solution has the advantage that the namespace doesn't have to be restored.

```
class Task2(luigi.Task):
    task_namespace = 'namespace2'
```

`luigi.task.id_to_name_and_params(task_id)`

exception `luigi.task.BulkCompleteNotImplementedError`

Bases: `exceptions.NotImplementedError`

This is here to trick pylint.

pylint thinks anything raising `NotImplementedError` needs to be implemented in any subclass. `bulk_complete` isn't like that. This tricks pylint into thinking that the default implementation is a valid implementation and not an abstract method.

class `luigi.task.Task(*args, **kwargs)`

Bases: `object`

This is the base class of all Luigi Tasks, the base unit of work in Luigi.

A Luigi Task describes a unit or work.

The key methods of a Task, which must be implemented in a subclass are:

- `run()` - the computation done by this task.
- `requires()` - the list of Tasks that this Task depends on.
- `output()` - the output `Target` that this Task creates.

Parameters to the Task should be declared as members of the class, e.g.:

```
.. code-block:: python
```

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

Each Task exposes a constructor accepting all `Parameter` (and values) as `kwargs`. e.g. `MyTask(count=10)` would instantiate `MyTask`.

In addition to any declared properties and methods, there are a few non-declared properties, which are created by the `Register` metaclass:

Task.task_namespace optional string which is prepended to the task name for the sake of scheduling. If it isn't overridden in a Task, whatever was last declared using `luigi.namespace` will be used.

Task._parameters list of (`parameter_name`, `parameter`) tuples for this task class

Constructor to resolve values for all `Parameters`.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

priority = 0

Priority of the task: the scheduler should favor available tasks with higher priority values first. See [Task priority](#)

disabled = False

resources = {}

Resources used by the task. Should be formatted like {"scp": 1} to indicate that the task requires 1 unit of the scp resource.

worker_timeout = None

Number of seconds after which to time out the run function. No timeout if set to 0. Defaults to 0 or value in client.cfg

use_cmdline_section

Property used by core config such as `-workers` etc. These will be exposed without the class as prefix.

classmethod event_handler (*event*)

Decorator for adding event handlers.

trigger_event (*event, *args, **kwargs*)

Trigger that calls all of the specified events associated with this class.

task_module

Returns what Python module to import to get access to this class.

task_family = 'Task'

classmethod get_params ()

Returns all of the Parameters for this Task.

classmethod get_param_values (*params, args, kwargs*)

Get the values of the parameters from the args and kwargs.

Parameters

- **params** – list of (param_name, Parameter).
- **args** – positional arguments
- **kwargs** – keyword arguments.

Returns list of (*name, value*) tuples, one for each parameter.

initialized ()

Returns `True` if the Task is initialized and `False` otherwise.

classmethod from_str_params (*params_str=None*)

Creates an instance from a str->str hash.

Parameters **params_str** – dict of param name -> value.

to_str_params ()

Convert all parameters to a str->str hash.

clone (*cls=None, **kwargs*)

Creates a new instance from an existing instance where some of the args have changed.

There's at least two scenarios where this is useful (see `test/clone_test.py`):

- remove a lot of boiler plate when you have recursive dependencies and lots of args
- there's task inheritance and some logic is on the base class

Parameters

- **cls** –
- **kwargs** –

Returns

complete()

If the task has any outputs, return `True` if all outputs exists. Otherwise, return `False`.

However, you may freely override this method with custom logic.

classmethod bulk_complete(*parameter_tuples*)

Returns those of *parameter_tuples* for which this Task is complete.

Override (with an efficient implementation) for efficient scheduling with range tools. Keep the logic consistent with that of `complete()`.

output()

The output that this Task produces.

The output of the Task determines if the Task needs to be run—the task is considered finished iff the outputs all exist. Subclasses should override this method to return a single `Target` or a list of `Target` instances.

Implementation note If running multiple workers, the output must be a resource that is accessible by all workers, such as a DFS or database. Otherwise, workers might compute the same output since they don't see the work done by other workers.

See [*Task.output*](#)

requires()

The Tasks that this Task depends on.

A Task will only run if all of the Tasks that it requires are completed. If your Task does not require any other Tasks, then you don't need to override this method. Otherwise, a Subclasses can override this method to return a single Task, a list of Task instances, or a dict whose values are Task instances.

See [*Task.requires*](#)

process_resources()

Override in “template” tasks which provide common resource functionality but allow subclasses to specify additional resources while preserving the name for consistent end-user experience.

input()

Returns the outputs of the Tasks returned by [*requires\(\)*](#)

See [*Task.input*](#)

Returns a list of `Target` objects which are specified as outputs of all required Tasks.

deps()

Internal method used by the scheduler.

Returns the flattened list of requires.

run()

The task run method, to be overridden in a subclass.

See [*Task.run*](#)

on_failure(*exception*)

Override for custom error handling.

This method gets called if an exception is raised in [*run\(\)*](#). Return value of this method is json encoded and sent to the scheduler as the *expl* argument. Its string representation will be used as the body of the error email sent out if any.

Default behavior is to return a string representation of the stack trace.

on_success()

Override for doing custom completion handling for a larger class of tasks

This method gets called when `run()` completes without raising any exceptions.

The returned value is json encoded and sent to the scheduler as the *expl* argument.

Default behavior is to send an None value

task_namespace = None

class `luigi.task.MixinNaiveBulkComplete`

Bases: `object`

Enables a Task to be efficiently scheduled with e.g. range tools, by providing a `bulk_complete` implementation which checks completeness in a loop.

Applicable to tasks whose completeness checking is cheap.

This doesn't exploit output location specific APIs for speed advantage, nevertheless removes redundant scheduler roundtrips.

classmethod `bulk_complete` (*parameter_tuples*)

`luigi.task.externalize` (*task*)

Returns an externalized version of the Task.

See [ExternalTask](#).

class `luigi.task.ExternalTask` (**args, **kwargs*)

Bases: `luigi.task.Task`

Subclass for references to external dependencies.

An ExternalTask's does not have a `run` implementation, which signifies to the framework that this Task's `output()` is generated outside of Luigi.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

run = NotImplemented

task_namespace = None

class `luigi.task WrapperTask` (**args, **kwargs*)

Bases: `luigi.task.Task`

Use for tasks that only wrap other tasks and that by definition are done if all their requirements exist.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

complete()

task_namespace = None

```
class luigi.task.Config(*args, **kwargs)
```

Bases: `luigi.task.Task`

Used for configuration that's not specific to a certain task

TODO: let's refactor Task & Config so that it inherits from a common ParamContainer base class

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

task_namespace = None

```
luigi.task.getpaths(struct)
```

Maps all Tasks in a structured data object to their `.output()`.

```
luigi.task.flatten(struct)
```

Creates a flat list of all all items in structured output (dicts, lists, items):

```
>>> sorted(flatten({'a': 'foo', 'b': 'bar'}))
['bar', 'foo']
>>> sorted(flatten(['foo', ['bar', 'troll']]))
['bar', 'foo', 'troll']
>>> flatten('foo')
['foo']
>>> flatten(42)
[42]
```

```
luigi.task.flatten_output(task)
```

Lists all output targets by recursively walking output-less (wrapper) tasks.

FIXME order consistently.

luigi.task_history module

Abstract class for task history. Currently the only subclass is `DbTaskHistory`.

```
class luigi.task_history.Task(task_id, status, host=None)
```

Bases: object

Interface for methods on TaskHistory

```
class luigi.task_history.TaskHistory
```

Bases: object

Abstract Base Class for updating the run history of a task

task_scheduled(task_id)

task_finished(task_id, successful)

task_started(task_id, worker_host)

```
class luigi.task_history.NopHistory
```

Bases: `luigi.task_history.TaskHistory`

task_scheduled(task_id)

task_finished(task_id, successful)

task_started (*task_id*, *worker_host*)

luigi.task_register module

Define the centralized register of all *Task* classes.

exception `luigi.task_register.TaskClassException`
 Bases: `exceptions.Exception`

class `luigi.task_register.Register`
 Bases: `abc.ABCMeta`

The Metaclass of Task.

Acts as a global registry of Tasks with the following properties:

- 1.Cache instances of objects so that eg. `X(1, 2, 3)` always returns the same object.
- 2.Keep track of all subclasses of Task and expose them.

AMBIGUOUS_CLASS = <object object>

If this value is returned by `_get_reg()` then there is an ambiguous task name (two Task have the same name). This denotes an error.

classmethod `clear_instance_cache()`
 Clear/Reset the instance cache.

classmethod `disable_instance_cache()`
 Disables the instance cache.

task_family
 The task family for the given class.

If `cls.task_namespace` is None then it's the name of the class. Otherwise, `<task_namespace>.` is prefixed to the class name.

classmethod `task_names()`
 List of task names as strings

classmethod `tasks_str()`
 Human-readable register contents dump.

classmethod `get_task_cls(name)`
 Returns an unambiguous class or raises an exception.

classmethod `get_all_params()`
 Compiles and returns all parameters for all Task.

Returns a generator of tuples (TODO: we should make this more elegant)

`luigi.task_register.load_task(module, task_name, params_str)`
 Imports task dynamically given a module and a task name.

luigi.task_status module

Possible values for a Task's status in the Scheduler

luigi.util module

`luigi.util.common_params (task_instance, task_cls)`
Grab all the values in `task_instance` that are found in `task_cls`.

`luigi.util.task_wraps (P)`

class `luigi.util.inherits (task_to_inherit)`
Bases: object
Task inheritance.
Usage:

```
class AnotherTask(luigi.Task):
    n = luigi.IntParameter()
    # ...

@inherits(AnotherTask):
class MyTask(luigi.Task):
    def requires(self):
        return self.clone_parent()

    def run(self):
        print self.n # this will be defined
        # ...
```

class `luigi.util.requires (task_to_require)`
Bases: object
Same as `@inherits`, but also auto-defines the `requires` method.

class `luigi.util.copies (task_to_copy)`
Bases: object
Auto-copies a task.
Usage:

```
@copies(MyTask):
class CopyOfMyTask(luigi.Task):
    def output(self):
        return LocalTarget(self.date.strftime('/var/xyz/report-%Y-%m-%d'))
```

`luigi.util.delegates (task_that_delegates)`
Lets a task call methods on subtask(s).

The way this works is that the subtask is run as a part of the task, but the task itself doesn't have to care about the requirements of the subtasks. The subtask doesn't exist from the scheduler's point of view, and its dependencies are instead required by the main task.

Example:

```
class PowersOfN(luigi.Task):
    n = luigi.IntParameter()
    def f(self, x): return x ** self.n

@delegates
class T(luigi.Task):
    def subtasks(self): return PowersOfN(5)
    def run(self): print self.subtasks().f(42)
```


`luigi.util.previous(task)`

Return a previous Task of the same family.

By default checks if this task family only has one non-global parameter and if it is a `DateParameter`, `DateHourParameter` or `DateIntervalParameter` in which case it returns with the time decremented by 1 (hour, day or interval)

`luigi.util.get_previous_completed(task, max_steps=10)`

luigi.webhdfs module

luigi.webhdfs has moved to luigi.contrib.webhdfs

luigi.worker module

The worker communicates with the scheduler and does two things:

1. Sends all tasks that has to be run
2. Gets tasks from the scheduler that should be run

When running in local mode, the worker talks directly to a `CentralPlannerScheduler` instance. When you run a central server, the worker will talk to the scheduler using a `RemoteScheduler` instance.

exception `luigi.worker.TaskException`

Bases: `exceptions.Exception`

class `luigi.worker.TaskProcess(task, worker_id, result_queue, random_seed=False, worker_timeout=0)`

Bases: `multiprocessing.process.Process`

Wrap all task execution in this class.

Mainly for convenience since this is run in a separate process.

run()

class `luigi.worker.SingleProcessPool`

Bases: `object`

Dummy process pool for using a single processor.

Imitates the api of `multiprocessing.Pool` using single-processor equivalents.

apply_async(function, args)

class `luigi.worker.DequeQueue`

Bases: `collections.deque`

deque wrapper implementing the `Queue` interface.

put()

Add an element to the right side of the deque.

get()

Remove and return the rightmost element.

exception `luigi.worker.AsyncCompletionException(trace)`

Bases: `exceptions.Exception`

Exception indicating that something went wrong with checking complete.

```
class luigi.worker.TracebackWrapper(trace)
    Bases: object
```

Class to wrap tracebacks so we can know they're not just strings.

```
luigi.worker.check_complete(task, out_queue)
    Checks if task is complete, puts the result to out_queue.
```

```
class luigi.worker.worker(*args, **kwargs)
    Bases: luigi.task.Config

    Constructor to resolve values for all Parameters.

    For example, the Task:
```

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
ping_interval = <luigi.parameter.FloatParameter object>
keep_alive = <luigi.parameter.BoolParameter object>
count_uniques = <luigi.parameter.BoolParameter object>
wait_interval = <luigi.parameter.IntParameter object>
max_reschedules = <luigi.parameter.IntParameter object>
timeout = <luigi.parameter.IntParameter object>
task_limit = <luigi.parameter.IntParameter object>
retry_external_tasks = <luigi.parameter.BoolParameter object>
task_namespace = None
```

```
class luigi.worker.KeepAliveThread(scheduler, worker_id, ping_interval)
    Bases: threading.Thread
```

Periodically tell the scheduler that the worker still lives.

```
stop()
run()
```

```
class luigi.worker.Worker(scheduler=None, worker_id=None, worker_processes=1, assistant=False,
                          **kwargs)
    Bases: object
```

Worker object communicates with a scheduler.

Simple class that talks to a scheduler and:

- tells the scheduler what it has to do + its dependencies
- asks for stuff to do (pulls it in a loop and runs it)

```
stop()
    Stop the KeepAliveThread associated with this Worker.
```

This should be called whenever you are done with a worker instance to clean up.

Warning: this should *only* be performed if you are sure this worker is not performing any work or will perform any work after this has been called

TODO: also kill all currently running tasks

TODO (maybe): Worker should be/have a context manager to enforce calling this whenever you stop using a Worker instance

add (*task*, *multiprocess=False*)

Add a Task for the worker to check and possibly schedule and run.

Returns True if task and its dependencies were successfully scheduled or completed before.

run ()

Returns True if all scheduled tasks were executed successfully.

11.1.3 Module contents

Package containing core luigi functionality.

11.2 luigi.contrib package

11.2.1 Subpackages

luigi.contrib.hdfs package

Submodules

luigi.contrib.hdfs.abstract_client module Module containing abstract class about hdfs clients.

class `luigi.contrib.hdfs.abstract_client.HdfsFileSystem`

Bases: `luigi.target.FileSystem`

This client uses Apache 2.x syntax for file system commands, which also matched CDH4.

rename (*path*, *dest*)

Rename or move a file

rename_dont_move (*path*, *dest*)

Override this method with an implementation that uses `rename2`, which is a rename operation that never moves.

For instance, `rename2 a b` never moves `a` into `b` folder.

Currently, the hadoop cli does not support this operation.

We keep the interface simple by just aliasing this to normal `rename` and let individual implementations redefine the method.

`rename2` - <https://github.com/apache/hadoop/blob/ae91b13/hadoop-hdfs-project/hadoop-hdfs/src/main/java/org/apache/hadoop/hdfs/protocol/ClientProtocol.java> (lines 483-523)

remove (*path*, *recursive=True*, *skip_trash=False*)

chmod (*path*, *permissions*, *recursive=False*)

chown (*path*, *owner*, *group*, *recursive=False*)

count (*path*)

Count contents in a directory

copy (*path*, *destination*)

put (*local_path*, *destination*)

```
get (path, local_destination)
makedirs (path, parents=True, raise_if_exists=False)
listdir (path, ignore_directories=False, ignore_files=False, include_size=False, include_type=False,
          include_time=False, recursive=False)
touchz (path)
```

luigi.contrib.hdfs.clients module The implementations of the hdfs clients. The hadoop cli client and the snakebite client.

```
luigi.contrib.hdfs.clients.get_autoconfig_client (show_warnings=True)
    Creates the client as specified in the client.cfg configuration.
```

luigi.contrib.hdfs.config module You can configure what client by setting the “client” config under the “hdfs” section in the configuration, or using the `--hdfs-client` command line option. “hadoopcli” is the slowest, but should work out of the box. “snakebite” is the fastest, but requires Snakebite to be installed.

```
class luigi.contrib.hdfs.config.hdfs (*args, **kwargs)
```

Bases: `luigi.task.Config`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
client_version = <luigi.parameter.IntParameter object>
effective_user = <luigi.parameter.Parameter object>
snakebite_autoconfig = <luigi.parameter.BoolParameter object>
namenode_host = <luigi.parameter.Parameter object>
namenode_port = <luigi.parameter.IntParameter object>
client = <luigi.parameter.Parameter object>
tmp_dir = <luigi.parameter.Parameter object>
task_namespace = None
```

```
class luigi.contrib.hdfs.config.hadoopcli (*args, **kwargs)
```

Bases: `luigi.task.Config`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

```
command = <luigi.parameter.Parameter object>
version = <luigi.parameter.Parameter object>
task_namespace = None
```

```
luigi.contrib.hdfs.config.load_hadoop_cmd()
```

```
luigi.contrib.hdfs.config.get_configured_hadoop_version()
```

CDH4 (hadoop 2+) has a slightly different syntax for interacting with hdfs via the command line.

The default version is CDH4, but one can override this setting with “cdh3” or “apache1” in the hadoop section of the config in order to use the old syntax.

```
luigi.contrib.hdfs.config.get_configured_hdfs_client(show_warnings=True)
```

This is a helper that fetches the configuration value for ‘client’ in the [hdfs] section. It will return the client that retains backwards compatibility when ‘client’ isn’t configured.

```
luigi.contrib.hdfs.config.tmp_path(path=None, include_unix_username=True)
```

@param path: target path for which it is needed to generate temporary location @type path: str @type include_unix_username: bool @rtype: str

Note that include_unix_username might work on windows too.

luigi.contrib.hdfs.error module The implementations of the hdfs clients. The hadoop cli client and the snakebite client.

```
exception luigi.contrib.hdfs.error.HDFSClientError(command, returncode, stdout, stderr)
```

Bases: `exceptions.Exception`

luigi.contrib.hdfs.format module

```
class luigi.contrib.hdfs.format.HdfsReadPipe(path)
```

Bases: `luigi.format.InputPipeProcessWrapper`

```
class luigi.contrib.hdfs.format.HdfsAtomicWritePipe(path)
```

Bases: `luigi.format.OutputPipeProcessWrapper`

File like object for writing to HDFS

The referenced file is first written to a temporary location and then renamed to final location on `close()`. If `close()` isn’t called the temporary file will be cleaned up when this object is garbage collected

TODO: if this is buggy, change it so it first writes to a local temporary file and then uploads it on completion

```
abort()
```

```
close()
```

```
class luigi.contrib.hdfs.format.HdfsAtomicWriteDirPipe(path, data_extension='')
```

Bases: `luigi.format.OutputPipeProcessWrapper`

Writes a data<data_extension> file to a directory at <path>.

```
abort()
```

```
close()
```

```
class luigi.contrib.hdfs.format.PlainFormat
```

Bases: `luigi.format.Format`

```
input = 'bytes'
```

```
output = 'hdfs'
```

```
hdfs_writer(path)
```

```
hdfs_reader(path)
```

```
pipe_reader(path)
```

```
pipe_writer(output_pipe)
```

```
class luigi.contrib.hdfs.format.PlainDirFormat
```

```
    Bases: luigi.format.Format
```

```
    input = 'bytes'
```

```
    output = 'hdfs'
```

```
    hdfs_writer (path)
```

```
    hdfs_reader (path)
```

```
    pipe_reader (path)
```

```
    pipe_writer (path)
```

```
class luigi.contrib.hdfs.format.CompatibleHdfsFormat (writer, reader, input=None)
```

```
    Bases: luigi.format.Format
```

```
    output = 'hdfs'
```

```
    pipe_writer (output)
```

```
    pipe_reader (input)
```

```
    hdfs_writer (output)
```

```
    hdfs_reader (input)
```

luigi.contrib.hdfs.hadoopcli_clients module The implementations of the hdfs clients. The hadoop cli client and the snakebite client.

```
luigi.contrib.hdfs.hadoopcli_clients.create_hadoopcli_client()
```

Given that we want one of the hadoop cli clients (unlike snakebite), this one will return the right one.

```
class luigi.contrib.hdfs.hadoopcli_clients.HdfsClient
```

```
    Bases: luigi.contrib.hdfs.abstract_client.HdfsFileSystem
```

This client uses Apache 2.x syntax for file system commands, which also matched CDH4.

```
    recursive_listdir_cmd = ['-ls', '-R']
```

```
    static call_check (command)
```

```
    exists (path)
```

Use `hadoop fs -stat` to check file existence.

```
    rename (path, dest)
```

```
    remove (path, recursive=True, skip_trash=False)
```

```
    chmod (path, permissions, recursive=False)
```

```
    chown (path, owner, group, recursive=False)
```

```
    count (path)
```

```
    copy (path, destination)
```

```
    put (local_path, destination)
```

```
    get (path, local_destination)
```

```
    getmerge (path, local_destination, new_line=False)
```

```
    mkdir (path, parents=True, raise_if_exists=False)
```

```
    listdir (path, ignore_directories=False, ignore_files=False, include_size=False, include_type=False,
              include_time=False, recursive=False)
```

touchz (*path*)

class `luigi.contrib.hdfs.hadoopcli_clients.HdfsClientCdh3`
 Bases: `luigi.contrib.hdfs.hadoopcli_clients.HdfsClient`

This client uses CDH3 syntax for file system commands.

mkdir (*path*)

No -p switch, so this will fail creating ancestors.

remove (*path*, *recursive=True*, *skip_trash=False*)

class `luigi.contrib.hdfs.hadoopcli_clients.HdfsClientApache1`
 Bases: `luigi.contrib.hdfs.hadoopcli_clients.HdfsClientCdh3`

This client uses Apache 1.x syntax for file system commands, which are similar to CDH3 except for the file existence check.

recursive_listdir_cmd = ['-lsr']

exists (*path*)

luigi.contrib.hdfs.snakebite_client module A luigi file system client that wraps around snakebite

Originally written by Alan Brenner <alan@magnetic.com> github.com/alanbbr

class `luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient`
 Bases: `luigi.contrib.hdfs.abstract_client.HdfsFileSystem`

A hdfs client using snakebite. Since Snakebite has a python API, it'll be about 100 times faster than the hadoop cli client, which does shell out to a java program on each file system operation.

static list_path (*path*)

get_bite ()

If Luigi has forked, we have a different PID, and need to reconnect.

exists (*path*)

Use snakebite.test to check file existence.

Parameters *path* (*string*) – path to test

Returns boolean, True if path exists in HDFS

rename (*path*, *dest*)

Use snakebite.rename, if available.

Parameters

- **path** (*either a string or sequence of strings*) – source file(s)
- **dest** (*string*) – destination file (single input) or directory (multiple)

Returns list of renamed items

rename_dont_move (*path*, *dest*)

Use snakebite.rename_dont_move, if available.

Parameters

- **path** (*string*) – source path (single input)
- **dest** (*string*) – destination path

Returns True if succeeded

Raises `snakebite.errors.FileAlreadyExistsException`

remove (*path*, *recursive=True*, *skip_trash=False*)

Use snakebite.delete, if available.

Parameters

- **path** (*either a string or a sequence of strings*) – delete-able file(s) or directory(ies)
- **recursive** (*boolean, default is True*) – delete directories trees like *nix: rm -r
- **skip_trash** (*boolean, default is False (use trash)*) – do or don't move deleted items into the trash first

Returns list of deleted items

chmod (*path*, *permissions*, *recursive=False*)

Use snakebite.chmod, if available.

Parameters

- **path** (*either a string or sequence of strings*) – update-able file(s)
- **permissions** (*octal*) – *nix style permission number
- **recursive** (*boolean, default is False*) – change just listed entry(ies) or all in directories

Returns list of all changed items

chown (*path*, *owner*, *group*, *recursive=False*)

Use snakebite.chown/chgrp, if available.

One of owner or group must be set. Just setting group calls chgrp.

Parameters

- **path** (*either a string or sequence of strings*) – update-able file(s)
- **owner** (*string*) – new owner, can be blank
- **group** (*string*) – new group, can be blank
- **recursive** (*boolean, default is False*) – change just listed entry(ies) or all in directories

Returns list of all changed items

count (*path*)

Use snakebite.count, if available.

Parameters **path** (*string*) – directory to count the contents of

Returns dictionary with content_size, dir_count and file_count keys

copy (*path*, *destination*)

Raise a NotImplementedError exception.

put (*local_path*, *destination*)

Raise a NotImplementedError exception.

get (*path*, *local_destination*)

Use snakebite.copyToLocal, if available.

Parameters

- **path** (*string*) – HDFS file
- **local_destination** (*string*) – path on the system running Luigi

mkdir (*path*, *parents=True*, *mode=493*, *raise_if_exists=False*)

Use snakebite.mkdir, if available.

Snakebite's mkdir method allows control over full path creation, so by default, tell it to build a full path to work like `hadoop fs -mkdir`.

Parameters

- **path** (*string*) – HDFS path to create
- **parents** (*boolean*, *default is True*) – create any missing parent directories
- **mode** (*octal*, *default 0755*) – *nix style owner/group/other permissions

lsdir (*path*, *ignore_directories=False*, *ignore_files=False*, *include_size=False*, *include_type=False*, *include_time=False*, *recursive=False*)

Use snakebite.ls to get the list of items in a directory.

Parameters

- **path** (*string*) – the directory to list
- **ignore_directories** (*boolean*, *default is False*) – if True, do not yield directory entries
- **ignore_files** (*boolean*, *default is False*) – if True, do not yield file entries
- **include_size** (*boolean*, *default is False (do not include)*) – include the size in bytes of the current item
- **include_type** (*boolean*, *default is False (do not include)*) – include the type (d or f) of the current item
- **include_time** (*boolean*, *default is False (do not include)*) – include the last modification time of the current item
- **recursive** (*boolean*, *default is False (do not recurse)*) – list subdirectory contents

Returns yield with a string, or if any of the include_* settings are true, a tuple starting with the path, and include_* items in order

touchz (*path*)

Raise a NotImplementedError exception.

luigi.contrib.hdfs.target module Provides access to HDFS using the *HdfsTarget*, a subclass of *Target*.

class luigi.contrib.hdfs.target.**HdfsTarget** (*path=None*, *format=None*, *is_tmp=False*, *fs=None*)

Bases: *luigi.target.FileSystemTarget*

fs

glob_exists (*expected_files*)

open (*mode='r'*)

remove (*skip_trash=False*)

rename (*path*, *raise_if_exists=False*)

Rename does not change self.path, so be careful with assumptions.

Not recommended for directories. Use move_dir. [spotify/luigi#522](#)

move (*path*, *raise_if_exists=False*)

Move does not change self.path, so be careful with assumptions.

Not recommended for directories. Use move_dir. [spotify/luigi#522](#)

move_dir (*path*)

Rename a directory.

The implementation uses *rename_dont_move*, which on some clients is just a normal *mv* operation, which can cause nested directories.

One could argue that the implementation should use the mkdir+raise_if_exists approach, but we at Spotify have had more trouble with that over just using plain mv. See [spotify/luigi#557](#)

is_writable ()

Currently only works with hadoopcli

Module contents

Provides access to HDFS using the `HdfsTarget`, a subclass of `Target`. You can configure what client by setting the “client” config under the “hdfs” section in the configuration, or using the `--hdfs-client` command line option. “hadoopcli” is the slowest, but should work out of the box. “snakebite” is the fastest, but requires Snakebite to be installed.

Currently (4th May) the `luigi.contrib.hdfs` module is under reorganization. We recommend importing the reexports from `luigi.contrib.hdfs` instead of the sub-modules, as we’re not yet sure how the final structure of the sub-modules will be. Eventually this module will be empty and you’ll have to import directly from the sub-modules like `luigi.contrib.hdfs.config`.

11.2.2 Submodules

luigi.contrib.bigquery module

class `luigi.contrib.bigquery.CreateDisposition`

Bases: `object`

CREATE_IF_NEEDED = ‘CREATE_IF_NEEDED’

CREATE_NEVER = ‘CREATE_NEVER’

class `luigi.contrib.bigquery.WriteDisposition`

Bases: `object`

WRITE_TRUNCATE = ‘WRITE_TRUNCATE’

WRITE_APPEND = ‘WRITE_APPEND’

WRITE_EMPTY = ‘WRITE_EMPTY’

class `luigi.contrib.bigquery.QueryMode`

Bases: `object`

INTERACTIVE = ‘INTERACTIVE’

BATCH = ‘BATCH’

class `luigi.contrib.bigquery.SourceFormat`

Bases: `object`

CSV = ‘CSV’

DATASTORE_BACKUP = 'DATASTORE_BACKUP'

NEWLINE_DELIMITED_JSON = 'NEWLINE_DELIMITED_JSON'

class `luigi.contrib.bigquery.BQDataset` (*project_id*, *dataset_id*)
 Bases: tuple

dataset_id
 Alias for field number 1

project_id
 Alias for field number 0

class `luigi.contrib.bigquery.BQTable`
 Bases: `luigi.contrib.bigquery.BQDataset`

dataset

class `luigi.contrib.bigquery.BigqueryClient` (*oauth_credentials=None*, *descriptor=''*,
http_=None)
 Bases: object

A client for Google BigQuery.

For details of how authentication and the descriptor work, see the documentation for the GCS client. The descriptor URL for BigQuery is <https://www.googleapis.com/discovery/v1/apis/bigquery/v2/rest>

dataset_exists (*dataset*)
 Returns whether the given dataset exists.

Parameters *dataset* (`BQDataset`) –

table_exists (*table*)
 Returns whether the given table exists.

Parameters *table* (`BQTable`) –

make_dataset (*dataset*, *raise_if_exists=False*, *body={}*)
 Creates a new dataset with the default permissions.

Parameters

- **dataset** (`BQDataset`) –
- **raise_if_exists** – whether to raise an exception if the dataset already exists.

Raises `luigi.target.FileAlreadyExists` if *raise_if_exists=True* and the dataset exists

delete_dataset (*dataset*, *delete_nonempty=True*)
 Deletes a dataset (and optionally any tables in it), if it exists.

Parameters

- **dataset** (`BQDataset`) –
- **delete_nonempty** – if true, will delete any tables before deleting the dataset

delete_table (*table*)
 Deletes a table, if it exists.

Parameters *table* (`BQTable`) –

list_datasets (*project_id*)
 Returns the list of datasets in a given project.

Parameters *project_id* (*str*) –

list_tables (*dataset*)

Returns the list of tables in a given dataset.

Parameters *dataset* ([BQDataset](#)) –

run_job (*project_id*, *body*, *dataset=None*)

Runs a bigquery “job”. See the documentation for the format of body.

Note: You probably don’t need to use this directly. Use the tasks defined below.

Parameters *dataset* ([BQDataset](#)) –

copy (*source_table*, *dest_table*, *create_disposition='CREATE_IF_NEEDED'*,
write_disposition='WRITE_TRUNCATE')
Copies (or appends) a table to another table.

Parameters

- **source_table** ([BQTable](#)) –
- **dest_table** ([BQTable](#)) –
- **create_disposition** ([CreateDisposition](#)) – whether to create the table if needed
- **write_disposition** ([WriteDisposition](#)) – whether to append/truncate/fail if the table exists

class `luigi.contrib.bigquery.BigqueryTarget` (*project_id*, *dataset_id*, *table_id*, *client=None*)

Bases: [luigi.target.Target](#)

classmethod `from_bqtable` (*table*, *client=None*)

A constructor that takes a [BQTable](#).

Parameters *table* ([BQTable](#)) –

exists ()

class `luigi.contrib.bigquery.BigqueryLoadTask` (**args*, ***kwargs*)

Bases: [luigi.task.Task](#)

Load data into bigquery from GCS.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

source_format

The source format to use (see [SourceFormat](#)).

write_disposition

What to do if the table already exists. By default this will fail the job.

See [WriteDisposition](#)

schema

Schema in the format defined at <https://cloud.google.com/bigquery/docs/reference/v2/jobs#configuration.load.schema>.

If the value is falsy, it is omitted and inferred by bigquery, which only works for CSV inputs.

max_bad_records

source_uris

Source data which should be in GCS.

run()

task_namespace = None

class `luigi.contrib.bigquery.BigqueryRunQueryTask(*args, **kwargs)`

Bases: `luigi.task.Task`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

write_disposition

What to do if the table already exists. By default this will fail the job.

See `WriteDisposition`

create_disposition

Whether to create the table or not. See `CreateDisposition`

query

The query, in text form.

query_mode

The query mode. See `QueryMode`.

run()

task_namespace = None

luigi.contrib.esindex module

Support for Elasticsearch (1.0.0 or newer).

Provides an `ElasticsearchTarget` and a `CopyToIndex` template task.

Modeled after `luigi.contrib.rdbms.CopyToTable`.

A minimal example (assuming elasticsearch is running on localhost:9200):

```
class ExampleIndex(CopyToIndex):
    index = 'example'

    def docs(self):
        return [{'_id': 1, 'title': 'An example document.'}]

if __name__ == '__main__':
    task = ExampleIndex()
    luigi.build([task], local_scheduler=True)
```

All options:

```
class ExampleIndex(CopyToIndex):
    host = 'localhost'
    port = 9200
    index = 'example'
```

```
doc_type = 'default'
purge_existing_index = True
marker_index_hist_size = 1

def docs(self):
    return [{'_id': 1, 'title': 'An example document.'}]

if __name__ == '__main__':
    task = ExampleIndex()
    luigi.build([task], local_scheduler=True)
```

Host, port, index, doc_type parameters are standard elasticsearch.

purge_existing_index will delete the index, whenever an update is required. This is useful, when one deals with “dumps” that represent the whole data, not just updates.

marker_index_hist_size sets the maximum number of entries in the ‘marker’ index:

- 0 (default) keeps all updates,
- 1 to only remember the most recent update to the index.

This can be useful, if an index needs to be recreated, even though the corresponding indexing task has been run sometime in the past - but a later indexing task might have altered the index in the meantime.

There are two luigi *client.cfg* configuration options:

```
[elasticsearch]

marker-index = update_log
marker-doc-type = entry
```

```
class luigi.contrib.esindex.ElasticsearchTarget(host, port, index, doc_type, up-
                                             date_id, marker_index_hist_size=0,
                                             http_auth=None, timeout=10, ex-
                                             tra_elasticsearch_args={})
```

Bases: *luigi.target.Target*

Target for a resource in Elasticsearch.

Parameters

- **host** (*str*) – Elasticsearch server host
- **port** (*int*) – Elasticsearch server port
- **index** (*str*) – index name
- **doc_type** (*str*) – doctype name
- **update_id** (*str*) – an identifier for this data set
- **marker_index_hist_size** (*int*) – list of changes to the index to remember
- **timeout** (*int*) – Elasticsearch connection timeout
- **extra_elasticsearch_args** – extra args for Elasticsearch

marker_index = ‘update_log’

marker_doc_type = ‘entry’

marker_index_document_id()

Generate an id for the indicator document.

touch()

Mark this update as complete.

The document id would be sufficient but, for documentation, we index the parameters *update_id*, *target_index*, *target_doc_type* and *date* as well.

exists()

Test, if this task has been run.

create_marker_index()

Create the index that will keep track of the tasks if necessary.

ensure_hist_size()

Shrink the history of updates for a *index/doc_type* combination down to *self.marker_index_hist_size*.

class `luigi.contrib.esindex.CopyToIndex(*args, **kwargs)`

Bases: `luigi.task.Task`

Template task for inserting a data set into Elasticsearch.

Usage:

1. Subclass and override the required *index* attribute.
2. Implement a custom *docs* method, that returns an iterable over the documents. A document can be a JSON string, e.g. from a newline-delimited JSON (ldj) file (default implementation) or some dictionary.

Optional attributes:

- *doc_type* (default),
- *host* (localhost),
- *port* (9200),
- *settings* ({'settings': {}})
- *mapping* (None),
- *chunk_size* (2000),
- *raise_on_error* (True),
- *purge_existing_index* (False),
- *marker_index_hist_size* (0)

If settings are defined, they are only applied at index creation time.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

host

ES hostname.

port

ES port.

http_auth

ES optional http auth information as either ':' separated string or a tuple, e.g. ('user', 'pass') or "user:pass".

index

The target index.

May exist or not.

doc_type

The target doc_type.

mapping

Dictionary with custom mapping or *None*.

settings

Settings to be used at index creation time.

chunk_size

Single API call for this number of docs.

raise_on_error

Whether to fail fast.

purge_existing_index

Whether to delete the *index* completely before any indexing.

marker_index_hist_size

Number of event log entries in the marker index. 0: unlimited.

timeout

Timeout.

extra_elasticsearch_args

Extra arguments to pass to the Elasticsearch constructor

docs ()

Return the documents to be indexed.

Beside the user defined fields, the document may contain an *_index*, *_type* and *_id*.

create_index ()

Override to provide code for creating the target index.

By default it will be created without any special settings or mappings.

delete_index ()

Delete the index, if it exists.

update_id ()

This id will be a unique identifier for this indexing task.

output ()

Returns a ElasticsearchTarget representing the inserted dataset.

Normally you don't override this.

run ()

Run task, namely:

- purge existing index, if requested (*purge_existing_index*),
- create the index, if missing,
- apply mappings, if given,
- set refresh interval to -1 (disable) for performance reasons,
- bulk index in batches of size *chunk_size* (2000),

- set refresh interval to 1s,
- refresh Elasticsearch,
- create entry in marker index.

task_namespace = None

luigi.contrib.ftp module

This library is a wrapper of ftplib. It is convenient to move data from/to FTP.

There is an example on how to use it (example/ftp_experiment_outputs.py)

You can also find unittest for each class.

Be aware that normal ftp do not provide secure communication.

class luigi.contrib.ftp.**RemoteFileSystem**(*host, username=None, password=None, port=21, tls=False*)

Bases: *luigi.target.FileSystem*

exists (*path, mtime=None*)

Return *True* if file or directory at *path* exist, *False* otherwise.

Additional check on modified time when *mtime* is passed in.

Return *False* if the file's modified time is older *mtime*.

remove (*path, recursive=True*)

Remove file or directory at location *path*.

Parameters

- **path** (*str*) – a path within the FileSystem to remove.
- **recursive** (*bool*) – if the path is a directory, recursively remove the directory and all of its descendants. Defaults to *True*.

put (*local_path, path*)

get (*path, local_path*)

class luigi.contrib.ftp.**AtomicFtpFile** (*fs, path*)

Bases: *luigi.target.AtomicLocalFile*

Simple class that writes to a temp file and upload to ftp on close().

Also cleans up the temp file if close is not invoked.

Initializes an AtomicFtpfile instance. :param fs: :param path: :type path: str

move_to_final_destination ()

fs

class luigi.contrib.ftp.**RemoteTarget** (*path, host, format=None, username=None, password=None, port=21, mtime=None, tls=False*)

Bases: *luigi.target.FileSystemTarget*

Target used for reading from remote files.

The target is implemented using ssh commands streaming data over the network.

fs

open (*mode*)

Open the FileSystem target.

This method returns a file-like object which can either be read from or written to depending on the specified mode.

Parameters *mode* (*str*) – the mode *r* opens the FileSystemTarget in read-only mode, whereas *w* will open the FileSystemTarget in write mode. Subclasses can implement additional options.

exists ()

put (*local_path*)

get (*local_path*)

luigi.contrib.gcs module

luigi bindings for Google Cloud Storage

exception `luigi.contrib.gcs.InvalidDeleteException`

Bases: `luigi.target.FileSystemException`

class `luigi.contrib.gcs.GCSClient` (*oauth_credentials=None*, *descriptor=''*, *http_=None*)

Bases: `luigi.target.FileSystem`

An implementation of a FileSystem over Google Cloud Storage.

There are several ways to use this class. By default it will use the app default credentials, as described at <https://developers.google.com/identity/protocols/application-default-credentials>. Alternatively, you may pass an oauth2client credentials object. e.g. to use a service account:

```
credentials = oauth2client.client.SignedJwtAssertionCredentials(
    '012345678912-ThisIsARandomServiceAccountEmail@developer.gserviceaccount.com',
    'These are the contents of the p12 file that came with the service account',
    scope='https://www.googleapis.com/auth/devstorage.read_write')
client = GCSClient(oauth_credentials=credentials)
```

Warning: By default this class will use “automated service discovery” which will require a connection to the web. The google api client downloads a JSON file to “create” the library interface on the fly. If you want a more hermetic build, you can pass the contents of this file (currently found at <https://www.googleapis.com/discovery/v1/apis/storage/v1/rest>) as the *descriptor* argument.

exists (*path*)

isdir (*path*)

remove (*path*, *recursive=True*)

put (*filename*, *dest_path*, *mimetype=None*)

put_string (*contents*, *dest_path*, *mimetype=None*)

mkdir (*path*, *parents=True*, *raise_if_exists=False*)

copy (*source_path*, *destination_path*)

rename (*source_path*, *destination_path*)

Rename/move an object from one S3 location to another.

listdir (*path*)

Get an iterable with S3 folder contents. Iterable contains paths relative to queried path.

download (*path*)

```
class luigi.contrib.gcs.AtomicGCSFile (path, gcs_client)
```

Bases: `luigi.target.AtomicLocalFile`

A GCS file that writes to a temp file and put to GCS on close.

```
move_to_final_destination ()
```

```
class luigi.contrib.gcs.GCSTarget (path, format=None, client=None)
```

Bases: `luigi.target.FileSystemTarget`

```
fs = None
```

```
open (mode='r')
```

```
class luigi.contrib.gcs.GCSFlagTarget (path, format=None, client=None, flag='_SUCCESS')
```

Bases: `luigi.contrib.gcs.GCSTarget`

Defines a target directory with a flag-file (defaults to `_SUCCESS`) used to signify job success.

This checks for two things:

- the path exists (just like the `GCSTarget`)
- the `_SUCCESS` file exists within the directory.

Because Hadoop outputs into a directory and not a single file, the path is assumed to be a directory.

This is meant to be a handy alternative to `AtomicGCSFile`.

The `AtomicFile` approach can be burdensome for GCS since there are no directories, per se.

If we have 1,000,000 output files, then we have to rename 1,000,000 objects.

Initializes a `S3FlagTarget`.

Parameters

- **path** (*str*) – the directory where the files are stored.
- **client** –
- **flag** (*str*) –

```
fs = None
```

```
exists ()
```

luigi.contrib.hadoop module

Run Hadoop Mapreduce jobs using Hadoop Streaming. To run a job, you need to subclass `luigi.contrib.hadoop.JobTask` and implement a `mapper` and `reducer` methods. See [Example – Top Artists](#) for an example of how to run a Hadoop job.

```
class luigi.contrib.hadoop.BaseHadoopJobTask (*args, **kwargs)
```

Bases: `luigi.task.Task`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask (luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask (count=10)`.

```
batch_counter_default = 1
```

```
deps ()
final_combiner = NotImplemented
final_mapper = NotImplemented
final_reducer = NotImplemented
init_hadoop ()
init_local ()
    Implement any work to setup any internal datastructure etc here.
    You can add extra input using the requires_local/input_local methods.
    Anything you set on the object will be pickled and available on the Hadoop nodes.
input_hadoop ()
input_local ()
job_runner ()
jobconfs ()
mr_priority = NotImplemented
on_failure (exception)
pool = <luigi.parameter.Parameter object>
requires_hadoop ()
requires_local ()
    Default impl - override this method if you need any local input to be accessible in init().
run ()
task_id = None
task_namespace = None

class luigi.contrib.hadoop.DefaultHadoopJobRunner
    Bases: luigi.contrib.hadoop.HadoopJobRunner
    The default job runner just reads from config and sets stuff.

exception luigi.contrib.hadoop.HadoopJobError (message, out=None, err=None)
    Bases: exceptions.RuntimeError

class luigi.contrib.hadoop.HadoopJobRunner (streaming_jar,      modules=None,      stream-
                                             ing_args=None,      libjars=None,      lib-
                                             jars_in_hdfs=None,      jobconfs=None,      in-
                                             put_format=None,      output_format=None,
                                             end_job_with_atomic_move_dir=True)
    Bases: luigi.contrib.hadoop.JobRunner
    Takes care of uploading & executing a Hadoop job using Hadoop streaming.
    TODO: add code to support Elastic Mapreduce (using boto) and local execution.

    finish ()
    run_job (job)

class luigi.contrib.hadoop.HadoopRunContext
    Bases: object
    kill_job (captured_signal=None, stack_frame=None)
```

class `luigi.contrib.hadoop.JobRunner`

Bases: `object`

run_job = `NotImplemented`

class `luigi.contrib.hadoop.JobTask(*args, **kwargs)`

Bases: `luigi.contrib.hadoop.BaseHadoopJobTask`

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

add_link (*src, dst*)

combiner = `NotImplemented`

data_interchange_format = `'python'`

deserialize

dump (*directory*=`''`)

Dump instance to file.

extra_files ()

Can be overridden in subclass.

Each element is either a string, or a pair of two strings (*src, dst*).

- *src* can be a directory (in which case everything will be copied recursively).

- *dst* can include subdirectories (foo/bar/baz.txt etc)

Uses Hadoop's -files option so that the same file is reused across tasks.

extra_modules ()

incr_counter (**args, **kwargs*)

Increments a Hadoop counter.

Since counters can be a bit slow to update, this batches the updates.

init_combiner ()

init_mapper ()

init_reducer ()

internal_reader (*input_stream*)

Reader which uses python eval on each part of a tab separated string. Yields a tuple of python objects.

internal_serialize

internal_writer (*outputs, stdout*)

Writer which outputs the python repr for each item.

job_runner ()

Get the MapReduce runner for this job.

If all outputs are `HdfsTargets`, the `DefaultHadoopJobRunner` will be used. Otherwise, the `LocalJobRunner` which streams all data through the local machine will be used (great for testing).

jobconfs ()

mapper (*item*)

Re-define to process an input item (usually a line of input data).

Defaults to identity mapper that sends all lines to the same reducer.

n_reduce_tasks = 25

reader (*input_stream*)

Reader is a method which iterates over input lines and outputs records.

The default implementation yields one argument containing the line for each line in the input.

reducer = **NotImplemented**

run_combiner (*stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>*)

run_mapper (*stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>*)

Run the mapper on the hadoop node.

run_reducer (*stdin=<open file '<stdin>', mode 'r'>, stdout=<open file '<stdout>', mode 'w'>*)

Run the reducer on the hadoop node.

serialize

task_namespace = **None**

writer (*outputs, stdout, stderr=<open file '<stderr>', mode 'w'>*)

Writer format is a method which iterates over the output records from the reducer and formats them for output.

The default implementation outputs tab separated items.

class `luigi.contrib.hadoop.LocalJobRunner` (*samplelines=None*)

Bases: `luigi.contrib.hadoop.JobRunner`

Will run the job locally.

This is useful for debugging and also unit testing. Tries to mimic Hadoop Streaming.

TODO: integrate with JobTask

group (*input_stream*)

run_job (*job*)

sample (*input_stream, n, output*)

`luigi.contrib.hadoop.attach` (**packages*)

Attach a python package to hadoop map reduce tarballs to make those packages available on the hadoop cluster.

`luigi.contrib.hadoop.create_packages_archive` (*packages, filename*)

Create a tar archive which will contain the files for the packages listed in packages.

`luigi.contrib.hadoop.dereference` (*f*)

`luigi.contrib.hadoop.fetch_task_failures` (*tracking_url*)

Uses mechanize to fetch the actual task logs from the task tracker.

This is highly opportunistic, and we might not succeed. So we set a low timeout and hope it works. If it does not, it's not the end of the world.

TODO: Yarn has a REST API that we should probably use instead:
<http://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/WebServicesIntro.html>

`luigi.contrib.hadoop.flatten` (*sequence*)

A simple generator which flattens a sequence.

Only one level is flattened.

```
(1, (2, 3), 4) -> (1, 2, 3, 4)
```

```
luigi.contrib.hadoop.get_extra_files(extra_files)
```

```
class luigi.contrib.hadoop.hadoop(*args, **kwargs)
    Bases: luigi.task.Config
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

pool = <luigi.parameter.Parameter object>

task_namespace = None

```
luigi.contrib.hadoop.run_and_track_hadoop_job(arglist, tracking_url_callback=None,
                                              env=None)
```

Runs the job by invoking the command from the given arglist. Finds tracking urls from the output and attempts to fetch errors using those urls if the job fails. Throws `HadoopJobError` with information about the error (including stdout and stderr from the process) on failure and returns normally otherwise.

Parameters

- **arglist** –
- **tracking_url_callback** –
- **env** –

Returns

luigi.contrib.hadoop_jar module

Provides functionality to run a Hadoop job using a Jar

```
luigi.contrib.hadoop_jar.fix_paths(job)
```

Coerce input arguments to use temporary files when used for output.

Return a list of temporary file pairs (tmpfile, destination path) and a list of arguments.

Converts each `HdfsTarget` to a string for the path.

```
exception luigi.contrib.hadoop_jar.HadoopJarJobError
    Bases: exceptions.Exception
```

```
class luigi.contrib.hadoop_jar.HadoopJarJobRunner
    Bases: luigi.contrib.hadoop.JobRunner
```

JobRunner for *hadoop jar* commands. Used to run a `HadoopJarJobTask`.

run_job (job)

```
class luigi.contrib.hadoop_jar.HadoopJarJobTask(*args, **kwargs)
    Bases: luigi.contrib.hadoop.BaseHadoopJobTask
```

A job task for *hadoop jar* commands that define a jar and (optional) main method.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

jar()

Path to the jar for this Hadoop Job.

main()

optional main method for this Hadoop Job.

job_runner()

atomic_output()

If True, then rewrite output arguments to be temp locations and atomically move them into place after the job finishes.

ssh()

Set this to run hadoop command remotely via ssh. It needs to be a dict that looks like {"host": "myhost", "key_file": None, "username": None, ["no_host_key_check": False]}

args()

Returns an array of args to pass to the job (after `hadoop jar <jar> <main>`).

task_namespace = None

luigi.contrib.hive module

exception `luigi.contrib.hive.HiveCommandError` (*message*, *out=None*, *err=None*)

Bases: `exceptions.RuntimeError`

`luigi.contrib.hive.load_hive_cmd()`

`luigi.contrib.hive.get_hive_syntax()`

`luigi.contrib.hive.run_hive` (*args*, *check_return_code=True*)

Runs the *hive* from the command line, passing in the given args, and returning stdout.

With the apache release of Hive, so of the table existence checks (which are done using DESCRIBE do not exit with a return code of 0 so we need an option to ignore the return code and just return stdout for parsing

`luigi.contrib.hive.run_hive_cmd` (*hivecmd*, *check_return_code=True*)

Runs the given hive query and returns stdout.

`luigi.contrib.hive.run_hive_script` (*script*)

Runs the contents of the given script in hive and returns stdout.

class `luigi.contrib.hive.HiveClient`

Bases: `object`

table_location (*table*, *database='default'*, *partition=None*)

Returns location of db.table (or db.table.partition). partition is a dict of partition key to value.

table_schema (*table*, *database='default'*)

Returns list of [(name, type)] for each column in database.table.

table_exists (*table*, *database='default'*, *partition=None*)

Returns true if db.table (or db.table.partition) exists. partition is a dict of partition key to value.

partition_spec (*partition*)

Turn a dict into a string partition specification


```
class luigi.contrib.hive.HiveCommandClient
```

Bases: `luigi.contrib.hive.HiveClient`

Uses *hive* invocations to find information.

table_location (*table*, *database*='default', *partition*=None)

table_exists (*table*, *database*='default', *partition*=None)

table_schema (*table*, *database*='default')

partition_spec (*partition*)

Turns a dict into the a Hive partition specification string.

```
class luigi.contrib.hive.ApacheHiveCommandClient
```

Bases: `luigi.contrib.hive.HiveCommandClient`

A subclass for the HiveCommandClient to (in some cases) ignore the return code from the hive command so that we can just parse the output.

table_schema (*table*, *database*='default')

```
class luigi.contrib.hive.MetastoreClient
```

Bases: `luigi.contrib.hive.HiveClient`

table_location (*table*, *database*='default', *partition*=None)

table_exists (*table*, *database*='default', *partition*=None)

table_schema (*table*, *database*='default')

partition_spec (*partition*)

```
class luigi.contrib.hive.HiveThriftContext
```

Bases: `object`

Context manager for hive metastore client.

```
luigi.contrib.hive.get_default_client()
```

```
class luigi.contrib.hive.HiveQueryTask(*args, **kwargs)
```

Bases: `luigi.contrib.hadoop.BaseHadoopJobTask`

Task to run a hive query.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

n_reduce_tasks = None

bytes_per_reducer = None

reducers_max = None

query ()

Text of query to run in hive

hiverc ()

Location of an rc file to run before the query if hiverc-location key is specified in client.cfg, will default to the value there otherwise returns None.

Returning a list of rc files will load all of them in order.

hiveconfs ()

Returns an dict of key=value settings to be passed along to the hive command line via `-hiveconf`. By default, sets `mapred.job.name` to `task_id` and if not `None`, sets:

- `mapred.reduce.tasks` (`n_reduce_tasks`)
- `mapred.fairscheduler.pool` (`pool`) or `mapred.job.queue.name` (`pool`)
- `hive.exec.reducers.bytes.per.reducer` (`bytes_per_reducer`)
- `hive.exec.reducers.max` (`reducers_max`)

job_runner ()

task_namespace = `None`

class `luigi.contrib.hive.HiveQueryRunner`

Bases: `luigi.contrib.hadoop.JobRunner`

Runs a `HiveQueryTask` by shelling out to `hive`.

prepare_outputs (*job*)

Called before job is started.

If output is a `FileSystemTarget`, create parent directories so the `hive` command won't fail

run_job (*job*)

class `luigi.contrib.hive.HiveTableTarget` (*table*, *database*=`'default'`, *client*=`None`)

Bases: `luigi.target.Target`

`exists` returns true if the table exists.

exists ()

path

Returns the path to this table in HDFS.

open (*mode*)

class `luigi.contrib.hive.HivePartitionTarget` (*table*, *partition*, *database*=`'default'`, *fail_missing_table*=`True`, *client*=`None`)

Bases: `luigi.target.Target`

`exists` returns true if the table's partition exists.

exists ()

path

Returns the path for this `HiveTablePartitionTarget`'s data.

open (*mode*)

class `luigi.contrib.hive.ExternalHiveTask` (**args*, ***kwargs*)

Bases: `luigi.task.ExternalTask`

External task that depends on a Hive table/partition.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

database = `<luigi.parameter.Parameter object>`

```

table = <luigi.parameter.Parameter object>
partition = <luigi.parameter.Parameter object>
task_namespace = None
output ()

```

luigi.contrib.mysql module

class `luigi.contrib.mysql.MySqlTarget` (*host, database, user, password, table, update_id*)
 Bases: `luigi.target.Target`

Target for a resource in MySQL.

Initializes a MySqlTarget instance.

Parameters

- **host** (*str*) – MySQL server address. Possibly a host:port string.
- **database** (*str*) – database name.
- **user** (*str*) – database user
- **password** (*str*) – password for specified user.
- **update_id** (*str*) – an identifier for this data set.

marker_table = 'table_updates'

touch (*connection=None*)

Mark this update as complete.

IMPORTANT, If the marker table doesn't exist, the connection transaction will be aborted and the connection reset. Then the marker table will be created.

exists (*connection=None*)

connect (*autocommit=False*)

create_marker_table ()

Create marker table if it doesn't exist.

Using a separate connection since the transaction might have to be reset.

luigi.contrib.pig module

Apache Pig support. Example configuration section in client.cfg:

```

[pig]
# pig home directory
home: /usr/share/pig

```

class `luigi.contrib.pig.PigJobTask` (**args, **kwargs*)
 Bases: `luigi.task.Task`

Constructor to resolve values for all Parameters.

For example, the Task:

```

class MyTask(luigi.Task):
    count = luigi.IntParameter()

```

can be instantiated as `MyTask(count=10)`.

pig_home()

pig_command_path()

pig_env_vars()

Dictionary of environment variables that should be set when running Pig.

Ex:: return { 'PIG_CLASSPATH': '/your/path' }

pig_properties()

Dictionary of properties that should be set when running Pig.

Example:

```
return { 'pig.additional.jars': '/path/to/your/jar' }
```

pig_parameters()

Dictionary of parameters that should be set for the Pig job.

Example:

```
return { 'YOUR_PARAM_NAME': 'Your param value' }
```

pig_options()

List of options that will be appended to the Pig command.

Example:

```
return ['-x', 'local']
```

output()

pig_script_path()

Return the path to the Pig script to be run.

run()

track_and_progress(cmd)

task_namespace = None

class `luigi.contrib.pig.PigRunContext`

Bases: `object`

kill_job (*captured_signal=None, stack_frame=None*)

exception `luigi.contrib.pig.PigJobError` (*message, out=None, err=None*)

Bases: `exceptions.RuntimeError`

luigi.contrib.pyspark_runner module

The pyspark program.

This module will be run by spark-submit for PySparkTask jobs.

The first argument is a path to the pickled instance of the PySparkTask, other arguments are the ones returned by `PySparkTask.app_options()`

class `luigi.contrib.pyspark_runner.PySparkRunner` (*job, *args*)

Bases: `object`

run()

luigi.contrib.rdbms module

A common module for posgres like databases, such as postgres or redshift

class `luigi.contrib.rdbms.CopyToTable(*args, **kwargs)`
 Bases: `luigi.task.MixinNaiveBulkComplete`, `luigi.task.Task`

An abstract task for inserting a data set into RDBMS.

Usage:

Subclass and override the following attributes:

- host*,
- database*,
- user*,
- password*,
- table*
- columns*

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

host

database

user

password

table

columns = []

null_values = (None,)

column_separator = '\t'

create_table (*connection*)

Override to provide code for creating the target table.

By default it will be created using types (optionally) specified in columns.

If overridden, use the provided connection object for setting up the table in order to create the table and insert data using the same transaction.

update_id ()

This update id will be a unique identifier for this insert on this table.

output ()

init_copy (*connection*)

Override to perform custom queries.

Any code here will be formed in the same transaction as the main copy, just prior to copying data. Example use cases include truncating the table or removing all data older than X in the database to keep a rolling window of data available in the table.

copy (*cursor*, *file*)

task_namespace = None

luigi.contrib.redis_store module

class luigi.contrib.redis_store.**RedisTarget** (*host*, *port*, *db*, *update_id*, *password*=None, *socket_timeout*=None, *expire*=None)

Bases: *luigi.target.Target*

Target for a resource in Redis.

Parameters

- **host** (*str*) – Redis server host
- **port** (*int*) – Redis server port
- **db** (*int*) – database index
- **update_id** (*str*) – an identifier for this data hash
- **password** (*str*) – a password to connect to the redis server
- **socket_timeout** (*int*) – client socket timeout
- **expire** (*int*) – timeout before the target is deleted

marker_prefix = <luigi.parameter.Parameter object>

marker_key ()

Generate a key for the indicator hash.

touch ()

Mark this update as complete.

We index the parameters *update_id* and *date*.

exists ()

Test, if this task has been run.

luigi.contrib.redshift module

class luigi.contrib.redshift.**RedshiftTarget** (*host*, *database*, *user*, *password*, *table*, *update_id*)

Bases: *luigi.postgres.PostgresTarget*

Target for a resource in Redshift.

Redshift is similar to postgres with a few adjustments required by redshift.

Args: *host* (*str*): Postgres server address. Possibly a host:port string. *database* (*str*): Database name *user* (*str*): Database user *password* (*str*): Password for specified user *update_id* (*str*): An identifier for this data set

marker_table = 'table_updates'

use_db_timestamps = False

class `luigi.contrib.redshift.S3CopyToTable(*args, **kwargs)`

Bases: `luigi.contrib.rdbms.CopyToTable`

Template task for inserting a data set into Redshift from s3.

Usage:

- Subclass and override the required attributes: ** host, * database, * user, * password, * table, * columns, * aws_access_key_id, * aws_secret_access_key, * s3_load_path.*

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

s3_load_path

Override to return the load path.

aws_access_key_id

Override to return the key id.

aws_secret_access_key

Override to return the secret access key.

copy_options

Add extra copy options, for example:

- TIMEFORMAT 'auto'
- IGNOREHEADER 1
- TRUNCATECOLUMNS
- IGNOREBLANKLINES

table_attributes()

Add extra table attributes, for example: `DISTSTYLE KEY DISTKEY (MY_FIELD) SORTKEY (MY_FIELD_2, MY_FIELD_3)`

do_truncate_table()

Return True if table should be truncated before copying new data in.

truncate_table(connection)

create_table(connection)

Override to provide code for creating the target table.

By default it will be created using types (optionally) specified in columns.

If overridden, use the provided connection object for setting up the table in order to create the table and insert data using the same transaction.

run()

If the target table doesn't exist, `self.create_table` will be called to attempt to create the table.

copy(cursor, f)

Defines copying from s3 into redshift.

output()

Returns a RedshiftTarget representing the inserted dataset.

Normally you don't override this.

does_table_exist (*connection*)

Determine whether the table already exists.

task_namespace = None

class `luigi.contrib.redshift.S3CopyJSONToTable` (**args, **kwargs*)

Bases: `luigi.contrib.redshift.S3CopyToTable`

Template task for inserting a JSON data set into Redshift from s3.

Usage:

- Subclass and override the required attributes:

- `-host`,
- `-database`,
- `-user`,
- `-password`,
- `-table`,
- `-columns`,
- `-aws_access_key_id`,
- `-aws_secret_access_key`,
- `-s3_load_path`,
- `-jsonpath`,
- `-copy_json_options`.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

jsonpath

Override the jsonpath schema location for the table.

copy_json_options

Add extra copy options, for example:

- GZIP
- LZOP

copy (*cursor, f*)

Defines copying JSON from s3 into redshift.

task_namespace = None

class `luigi.contrib.redshift.RedshiftManifestTask` (**args, **kwargs*)

Bases: `luigi.s3.S3PathTask`

Generic task to generate a manifest file that can be used in `S3CopyToTable` in order to copy multiple files from your s3 folder into a redshift table at once.

For full description on how to use the manifest file see <http://docs.aws.amazon.com/redshift/latest/dg/loading-data-files-using-manifest.html>

Usage:

- requires parameters

- **path** - s3 path to the generated manifest file, including the name of the generated file to be copied into a redshift table
- **folder_paths** - s3 paths to the folders containing files you wish to be copied

Output:

- generated manifest file

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

folder_paths = <luigi.parameter.Parameter object>

text_target = True

run()

task_namespace = None

class `luigi.contrib.redshift.KillOpenRedshiftSessions(*args, **kwargs)`

Bases: `luigi.task.Task`

An task for killing any open Redshift sessions in a given database. This is necessary to prevent open user sessions with transactions against the table from blocking drop or truncate table commands.

Usage:

Subclass and override the required *host*, *database*, *user*, and *password* attributes.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

connection_reset_wait_seconds = <luigi.parameter.IntParameter object>

host

database

user

password

update_id()

This update id will be a unique identifier for this insert on this table.

output()

Returns a RedshiftTarget representing the inserted dataset.

Normally you don't override this.

```
run ()
    Kill any open Redshift sessions for the given database.

task_namespace = None
```

luigi.contrib.scalding module

luigi.contrib.scalding.**logger** = <logging.Logger object>
Scalding support for Luigi.

Example configuration section in client.cfg:

```
[scalding]
# scala home directory, which should include a lib subdir with scala jars.
scala-home: /usr/share/scala

# scalding home directory, which should include a lib subdir with
# scalding-* assembly-* jars as built from the official Twitter build script.
scalding-home: /usr/share/scalding

# provided dependencies, e.g. jars required for compiling but not executing
# scalding jobs. Currently required jars:
# org.apache.hadoop/hadoop-core/0.20.2
# org.slf4j/slf4j-log4j12/1.6.6
# log4j/log4j/1.2.15
# commons-httpclient/commons-httpclient/3.1
# commons-cli/commons-cli/1.2
# org.apache.zookeeper/zookeeper/3.3.4
scalding-provided: /usr/share/scalding/provided

# additional jars required.
scalding-libjars: /usr/share/scalding/libjars
```

class luigi.contrib.scalding.**ScaldingJobRunner**
Bases: [luigi.contrib.hadoop.JobRunner](#)
JobRunner for *pyscald* commands. Used to run a ScaldingJobTask.

```
get_scala_jars (include_compiler=False)
get_scalding_jars ()
get_scalding_core ()
get_provided_jars ()
get_libjars ()
get_tmp_job_jar (source)
get_build_dir (source)
get_job_class (source)
build_job_jar (job)
run_job (job)
```

class luigi.contrib.scalding.**ScaldingJobTask** (*args, **kwargs)
Bases: [luigi.contrib.hadoop.BaseHadoopJobTask](#)
A job task for Scalding that define a scala source and (optional) main method.

requires() should return a dictionary where the keys are Scalding argument names and values are sub tasks or lists of subtasks.

For example:

```
{'input1': A, 'input2': C} => --input1 <Aoutput> --input2 <Coutput>
{'input1': [A, B], 'input2': [C]} => --input1 <Aoutput> <Boutput> --input2 <Coutput>
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

relpath (*current_file*, *rel_path*)

Compute path given current file and relative path.

source ()

Path to the scala source for this Scalding Job

Either one of source() or jar() must be specified.

jar ()

Path to the jar file for this Scalding Job

Either one of source() or jar() must be specified.

extra_jars ()

Extra jars for building and running this Scalding Job.

job_class ()

optional main job class for this Scalding Job.

job_runner ()

atomic_output ()

If True, then rewrite output arguments to be temp locations and atomically move them into place after the job finishes.

requires ()

job_args ()

Extra arguments to pass to the Scalding job.

task_namespace = None

args ()

Returns an array of args to pass to the job.

luigi.contrib.spark module

class `luigi.contrib.spark.SparkRunContext` (*proc*)

Bases: `object`

kill_job (*captured_signal=None*, *stack_frame=None*)

exception `luigi.contrib.spark.SparkJobError` (*message*, *out=None*, *err=None*)

Bases: `exceptions.RuntimeError`

class `luigi.contrib.spark.SparkSubmitTask` (*args, **kwargs)

Bases: `luigi.task.Task`

Template task for running a Spark job

Supports running jobs on Spark local, standalone, Mesos or Yarn

See <http://spark.apache.org/docs/latest/submitting-applications.html> for more information

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

name = None

entry_class = None

app = None

app_options ()

Subclass this method to map your task parameters to the app's arguments

spark_submit

master

deploy_mode

jars

py_files

files

conf

properties_file

driver_memory

driver_java_options

driver_library_path

driver_class_path

executor_memory

driver_cores

supervise

total_executor_cores

executor_cores

queue

num_executors

archives

hadoop_conf_dir

get_environment ()

```

spark_command()
app_command()
run()
task_namespace = None

```

```

class luigi.contrib.spark.PySparkTask(*args, **kwargs)
    Bases: luigi.contrib.spark.SparkSubmitTask

```

Template task for running an inline PySpark job

Simply implement the `main` method in your subclass

You can optionally define package names to be distributed to the cluster with `py_packages` (uses luigi's global py-packages configuration by default)

Constructor to resolve values for all Parameters.

For example, the Task:

```

class MyTask(luigi.Task):
    count = luigi.IntParameter()

```

can be instantiated as `MyTask(count=10)`.

```

app = '/home/docs/checkouts/readthedocs.org/user_builds/luigi/checkouts/latest/luigi/contrib/pyspark_runner.py'

```

```

deploy_mode = 'client'

```

```

name

```

```

py_packages

```

```

setup(conf)

```

Called by the `pyspark_runner` with a `SparkConf` instance that will be used to instantiate the `SparkContext`

Parameters `conf` – `SparkConf`

```

setup_remote(sc)

```

```

main(sc, *args)

```

Called by the `pyspark_runner` with a `SparkContext` and any arguments returned by `app_options()`

Parameters

- **sc** – `SparkContext`
- **args** – arguments list

```

app_command()

```

```

run()

```

```

task_namespace = None

```

```

class luigi.contrib.spark.SparkJob(*args, **kwargs)
    Bases: luigi.task.Task

```

Deprecated since version 1.1.1: Use `SparkSubmitTask` or `PySparkTask` instead.

Constructor to resolve values for all Parameters.

For example, the Task:

```

class MyTask(luigi.Task):
    count = luigi.IntParameter()

```

can be instantiated as `MyTask(count=10)`.

spark_workers = None

spark_master_memory = None

spark_worker_memory = None

queue = <luigi.parameter.Parameter object>

temp_hadoop_output_file = None

requires_local()

Default impl - override this method if you need any local input to be accessible in `init()`.

requires_hadoop()

input_local()

input()

deps()

jar()

job_class()

job_args()

output()

run()

track_progress(*proc*)

task_namespace = None

class `luigi.contrib.spark.Spark1xBackwardCompat(*args, **kwargs)`

Bases: `luigi.contrib.spark.SparkSubmitTask`

Adapts `SparkSubmitTask` interface to (Py)Spark1xJob interface

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

master

output()

spark_options()

dependency_jars()

job_args()

jars

app_options()

spark_command()

task_namespace = None

class `luigi.contrib.spark.Spark1xJob(*args, **kwargs)`
 Bases: `luigi.contrib.spark.Spark1xBackwardCompat`

Deprecated since version 1.1.1: Use `SparkSubmitTask` or `PySparkTask` instead.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

job_class()

jar()

entry_class

app

run()

task_namespace = None

class `luigi.contrib.spark.PySpark1xJob(*args, **kwargs)`
 Bases: `luigi.contrib.spark.Spark1xBackwardCompat`

Deprecated since version 1.1.1: Use `SparkSubmitTask` or `PySparkTask` instead.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

program()

app

run()

task_namespace = None

luigi.contrib.sparkey module

class `luigi.contrib.sparkey.SparkeyExportTask(*args, **kwargs)`
 Bases: `luigi.task.Task`

A luigi task that writes to a local sparkey log file.

Subclasses should implement the `requires` and `output` methods. The output must be a `luigi.LocalTarget`.

The resulting sparkey log file will contain one entry for every line in the input, mapping from the first value to a tab-separated list of the rest of the line.

To generate a simple key-value index, yield “key”, “value” pairs from the input(s) to this task.

separator = ‘\t’

run()

task_namespace = None

luigi.contrib.sqla module

Support for SQLAlchemy. Provides SQLAlchemyTarget for storing in databases supported by SQLAlchemy. The user would be responsible for installing the required database driver to connect using SQLAlchemy.

Minimal example of a job to copy data to database using SQLAlchemy is as shown below:

```
from sqlalchemy import String
import luigi
from luigi.contrib import sqla

class SQLATask(sqla.CopyToTable):
    # columns defines the table schema, with each element corresponding
    # to a column in the format (args, kwargs) which will be sent to
    # the sqlalchemy.Column(*args, **kwargs)
    columns = [
        (["item", String(64)], {"primary_key": True}),
        (["property", String(64)], {})
    ]
    connection_string = "sqlite://" # in memory SQLite database
    table = "item_property" # name of the table to store data

    def rows(self):
        for row in [("item1", "property1"), ("item2", "property2")]:
            yield row

if __name__ == '__main__':
    task = SQLATask()
    luigi.build([task], local_scheduler=True)
```

If the target table where the data needs to be copied already exists, then the column schema definition can be skipped and instead the reflect flag can be set as True. Here is a modified version of the above example:

```
from sqlalchemy import String
import luigi
from luigi.contrib import sqla

class SQLATask(sqla.CopyToTable):
    # If database table is already created, then the schema can be loaded
    # by setting the reflect flag to True
    reflect = True
    connection_string = "sqlite://" # in memory SQLite database
    table = "item_property" # name of the table to store data

    def rows(self):
        for row in [("item1", "property1"), ("item2", "property2")]:
            yield row

if __name__ == '__main__':
    task = SQLATask()
    luigi.build([task], local_scheduler=True)
```

In the above examples, the data that needs to be copied was directly provided by overriding the rows method. Alternately, if the data comes from another task, the modified example would look as shown below:

```
from sqlalchemy import String
import luigi
from luigi.contrib import sqla
from luigi.mock import MockFile
```



```

class BaseTask(luigi.Task):
    def output(self):
        return MockFile("BaseTask")

    def run(self):
        out = self.output().open("w")
        TASK_LIST = ["item%d\tproperty%d\n" % (i, i) for i in range(10)]
        for task in TASK_LIST:
            out.write(task)
        out.close()

class SQLATask(sqla.CopyToTable):
    # columns defines the table schema, with each element corresponding
    # to a column in the format (args, kwargs) which will be sent to
    # the sqlalchemy.Column(*args, **kwargs)
    columns = [
        ("item", String(64)), {"primary_key": True},
        ("property", String(64)), {}
    ]
    connection_string = "sqlite://" # in memory SQLite database
    table = "item_property" # name of the table to store data

    def requires(self):
        return BaseTask()

if __name__ == '__main__':
    task1, task2 = SQLATask(), BaseTask()
    luigi.build([task1, task2], local_scheduler=True)

```

In the above example, the output from *BaseTask* is copied into the database. Here we did not have to implement the *rows* method because by default *rows* implementation assumes every line is a row with column values separated by a tab. One can define *column_separator* option for the task if the values are say comma separated instead of tab separated.

You can pass in database specific connection arguments by setting the *connect_args* dictionary. The options will be passed directly to the DBAPI's connect method as keyword arguments.

The other option to *sqla.CopyToTable* that can be of help with performance aspect is the *chunk_size*. The default is 5000. This is the number of rows that will be inserted in a transaction at a time. Depending on the size of the inserts, this value can be tuned for performance.

See here for a [tutorial on building task pipelines using luigi](#) and using [SQLAlchemy in workflow pipelines](#).

Author: Gouthaman Balaraman Date: 01/02/2015

```

class luigi.contrib.sqla.SQLAlchemyTarget(connection_string, target_table, update_id,
                                         echo=False, connect_args=None)

```

Bases: *luigi.target.Target*

Database target using SQLAlchemy.

This will rarely have to be directly instantiated by the user.

Typical usage would be to override *luigi.contrib.sqla.CopyToTable* class to create a task to write to the database.

Constructor for the SQLAlchemyTarget.

Parameters

- **connection_string** (*str*) – SQLAlchemy connection string
- **target_table** (*str*) – The table name for the data

- **update_id** (*str*) – An identifier for this data set
- **echo** (*bool*) – Flag to setup SQLAlchemy logging
- **connect_args** (*dict*) – A dictionary of connection arguments

Returns

marker_table = None

class Connection (*engine, pid*)

Bases: tuple

engine

Alias for field number 0

pid

Alias for field number 1

SQLAlchemyTarget.**connect_args** = {}

SQLAlchemyTarget.**engine**

Return an engine instance, creating it if it doesn't exist.

Recreate the engine connection if it wasn't originally created by the current process.

SQLAlchemyTarget.**touch**()

Mark this update as complete.

SQLAlchemyTarget.**exists**()

SQLAlchemyTarget.**create_marker_table**()

Create marker table if it doesn't exist.

Using a separate connection since the transaction might have to be reset.

SQLAlchemyTarget.**open**(*mode*)

class luigi.contrib.sqla.**CopyToTable** (**args, **kwargs*)

Bases: *luigi.task.Task*

An abstract task for inserting a data set into SQLAlchemy RDBMS

Usage:

- subclass and override the required *connection_string*, *table* and *columns* attributes.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

echo = False

connect_args = {}

connection_string()

table

columns = []

column_separator = '\t'

chunk_size = 5000

reflect = False

create_table (*engine*)

Override to provide code for creating the target table.

By default it will be created using types specified in columns. If the table exists, then it binds to the existing table.

If overridden, use the provided connection object for setting up the table in order to create the table and insert data using the same transaction. :param engine: The sqlalchemy engine instance :type engine: object

update_id ()

This update id will be a unique identifier for this insert on this table.

output ()

rows ()

Return/yield tuples or lists corresponding to each row to be inserted.

This method can be overridden for custom file types or formats.

run ()

copy (*conn, ins_rows, table_bound*)

This method does the actual insertion of the rows of data given by ins_rows into the database. A task that needs row updates instead of insertions should overload this method. :param conn: The sqlalchemy connection object :param ins_rows: The dictionary of rows with the keys in the format _<column_name>. For example if you have a table with a column name “property”, then the key in the dictionary would be “_property”. This format is consistent with the bindparam usage in sqlalchemy. :param table_bound: The object referring to the table :return:

task_namespace = None

luigi.contrib.ssh module

Light-weight remote execution library and utilities.

There are some examples in the unittest, but I added another more luigi-specific in the examples directory (examples/ssh_remote_execution.py)

RemoteContext is meant to provide functionality similar to that of the standard library subprocess module, but where the commands executed are run on a remote machine instead, without the user having to think about prefixing everything with “ssh” and credentials etc.

Using this mini library (which is just a convenience wrapper for subprocess), *RemoteTarget* is created to let you stream data from a remotely stored file using the luigi *FileSystemTarget* semantics.

As a bonus, *RemoteContext* also provides a really cool feature that let’s you set up ssh tunnels super easily using a python context manager (there is an example in the integration part of unittests).

This can be super convenient when you want secure communication using a non-secure protocol or circumvent firewalls (as long as they are open for ssh traffic).

class luigi.contrib.ssh.**RemoteContext** (*host, **kwargs*)

Bases: object

Popen (*cmd, **kwargs*)

Remote Popen.

check_output (*cmd*)

Execute a shell command remotely and return the output.

Simplified version of Popen when you only want the output as a string and detect any errors.

tunnel (*args, **kws)

Open a tunnel between localhost:local_port and remote_host:remote_port via the host specified by this context.

Remember to close() the returned “tunnel” object in order to clean up after yourself when you are done with the tunnel.

class `luigi.contrib.ssh.RemoteFileSystem` (host, **kwargs)

Bases: `luigi.target.FileSystem`

exists (path)

Return *True* if file or directory at *path* exist, False otherwise.

listdir (path)

isdir (path)

Return *True* if directory at *path* exist, False otherwise.

remove (path, recursive=True)

Remove file or directory at location *path*.

mkdir (path, parents=True, raise_if_exists=False)

put (local_path, path)

get (path, local_path)

class `luigi.contrib.ssh.AtomicRemoteFileWriter` (fs, path)

Bases: `luigi.format.OutputPipeProcessWrapper`

close ()

tmp_path

fs

class `luigi.contrib.ssh.RemoteTarget` (path, host, format=None, **kwargs)

Bases: `luigi.target.FileSystemTarget`

Target used for reading from remote files.

The target is implemented using ssh commands streaming data over the network.

fs

open (mode='r')

put (local_path)

get (local_path)

luigi.contrib.target module

class `luigi.contrib.target.CascadingClient` (clients, method_names=None)

Bases: `object`

A FilesystemClient that will cascade failing function calls through a list of clients.

Which clients are used are specified at time of construction.

ALL_METHOD_NAMES = ['exists', 'rename', 'remove', 'chmod', 'chown', 'count', 'copy', 'get', 'put', 'mkdir', 'list', 'listdir']

luigi.contrib.webhdfs module

Provides a *WebHdfsTarget* and *WebHdfsClient* using the Python hdfs

```

class luigi.contrib.webhdfs.WebHdfsTarget (path, client=None, format=None)
    Bases: luigi.target.FileSystemTarget

    fs = None

    open (mode='r')

class luigi.contrib.webhdfs.ReadableWebHdfsFile (path, client)
    Bases: object

    read ()

    readlines (char='\n')

    close ()

class luigi.contrib.webhdfs.AtomicWebHdfsFile (path, client)
    Bases: luigi.target.AtomicLocalFile

    An Hdfs file that writes to a temp file and put to WebHdfs on close.

    move_to_final_destination ()

class luigi.contrib.webhdfs.WebHdfsClient (host=None, port=None, user=None)
    Bases: object

    get_config (key)

    walk (path, depth=1)

    exists (path)
        Returns true if the path exists and false otherwise.

    upload (hdfs_path, local_path, overwrite=False)

    download (hdfs_path, local_path, overwrite=False, n_threads=-1)

    remove (hdfs_path, recursive=False)

    read (hdfs_path, offset=0, length=None, buffer_size=None, chunk_size=1024, buffer_char=None)

```

11.2.3 Module contents

Package containing optional and-on functionality.

11.3 luigi.tools package

11.3.1 Submodules

luigi.tools.deps module

```

luigi.tools.deps.get_task_requires (task)

luigi.tools.deps.dfs_paths (start_task, goal_task_family, path=None)

```

class `luigi.tools.deps.upstream(*args, **kwargs)`

Bases: `luigi.task.Config`

Used to provide the parameter upstream-task-family

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

family = `<luigi.parameter.Parameter object>`

task_namespace = `None`

`luigi.tools.deps.find_deps(task, upstream_task_family)`

Finds all dependencies that start with the given task and have a path to `upstream_task_family`

Returns all deps on all paths between task and upstream

`luigi.tools.deps.find_deps_cli()`

Finds all tasks on all paths from provided CLI task

`luigi.tools.deps.main()`

luigi.tools.luigi_grep module

class `luigi.tools.luigi_grep.LuigiGrep(host, port)`

Bases: `object`

graph_url

prefix_search(job_name_prefix)

searches for jobs matching the given `job_name_prefix`.

status_search(status)

searches for jobs matching the given status

`luigi.tools.luigi_grep.main()`

luigi.tools.parse_task module

`luigi.tools.parse_task.id_to_name_and_params(task_id)`

Turn a `task_id` into a (`task_family`, {`params`}) tuple.

E.g. calling with `Foo(bar=bar, baz=baz)` returns `('Foo', {'bar': 'bar', 'baz': 'baz'})`.

luigi.tools.range module

Produces contiguous completed ranges of recurring tasks.

See `RangeDaily` and `RangeHourly` for basic usage.

Caveat - if gaps accumulate, their causes (e.g. missing dependencies) going unmonitored/unmitigated, then this will eventually keep retrying the same gaps over and over and make no progress to more recent times. (See `'task_limit'` and `'reverse'` parameters.) TODO foolproof against that kind of misuse?

class `luigi.tools.range.RangeEvent`

Bases: `luigi.event.Event`

Events communicating useful metrics.

COMPLETE_COUNT would normally be nondecreasing, and its derivative would describe performance (how many instances complete invocation-over-invocation).

COMPLETE_FRACTION reaching 1 would be a telling event in case of a backfill with defined start and stop. Would not be strikingly useful for a typical recurring task without stop defined, fluctuating close to 1.

DELAY is measured from the first found missing datehour till (current time + hours_forward), or till stop if it is defined. In hours for Hourly. TBD different units for other frequencies? TODO any different for reverse mode? From first missing till last missing? From last gap till stop?

COMPLETE_COUNT = 'event.tools.range.complete.count'

COMPLETE_FRACTION = 'event.tools.range.complete.fraction'

DELAY = 'event.tools.range.delay'

class `luigi.tools.range.RangeBase (*args, **kwargs)`

Bases: `luigi.task.WrapperTask`

Produces a contiguous completed range of a recurring task.

Made for the common use case where a task is parameterized by e.g. DateParameter, and assurance is needed that any gaps arising from downtime are eventually filled.

Emits events that one can use to monitor gaps and delays.

At least one of start and stop needs to be specified.

(This is quite an abstract base class for subclasses with different datetime parameter class, e.g. DateParameter, DateHourParameter, ..., and different parameter naming, e.g. days_back/forward, hours_back/forward, ..., as well as different documentation wording, for good user experience.)

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

of = <luigi.parameter.Parameter object>

start = <luigi.parameter.Parameter object>

stop = <luigi.parameter.Parameter object>

reverse = <luigi.parameter.BoolParameter object>

task_limit = <luigi.parameter.IntParameter object>

now = <luigi.parameter.IntParameter object>

datetime_to_parameter (dt)

parameter_to_datetime (p)

moving_start (now)

Returns a datetime from which to ensure contiguousness in the case when start is None or unfeasibly far back.

moving_stop (*now*)

Returns a datetime till which to ensure contiguousness in the case when stop is None or unfeasibly far forward.

finite_datetimes (*finite_start*, *finite_stop*)

Returns the individual datetimes in interval [*finite_start*, *finite_stop*) for which task completeness should be required, as a sorted list.

requires ()

missing_datetimes (*task_cls*, *finite_datetimes*)

Override in subclasses to do bulk checks.

Returns a sorted list.

This is a conservative base implementation that brutally checks completeness, instance by instance.

Inadvisable as it may be slow.

task_namespace = None

class `luigi.tools.range.RangeDailyBase` (**args*, ***kwargs*)

Bases: `luigi.tools.range.RangeBase`

Produces a contiguous completed range of a daily recurring task.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

start = <luigi.parameter.DateParameter object>

stop = <luigi.parameter.DateParameter object>

days_back = <luigi.parameter.IntParameter object>

days_forward = <luigi.parameter.IntParameter object>

datetime_to_parameter (*dt*)

parameter_to_datetime (*p*)

moving_start (*now*)

moving_stop (*now*)

finite_datetimes (*finite_start*, *finite_stop*)

Simply returns the points in time that correspond to turn of day.

task_namespace = None

class `luigi.tools.range.RangeHourlyBase` (**args*, ***kwargs*)

Bases: `luigi.tools.range.RangeBase`

Produces a contiguous completed range of an hourly recurring task.

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```


can be instantiated as `MyTask(count=10)`.

start = <luigi.parameter.DateHourParameter object>

stop = <luigi.parameter.DateHourParameter object>

hours_back = <luigi.parameter.IntParameter object>

hours_forward = <luigi.parameter.IntParameter object>

datetime_to_parameter (*dt*)

parameter_to_datetime (*p*)

moving_start (*now*)

moving_stop (*now*)

finite_datetimes (*finite_start*, *finite_stop*)

Simply returns the points in time that correspond to whole hours.

task_namespace = None

`luigi.tools.range.most_common` (*items*)

Wanted functionality from Counters (new in Python 2.7).

`luigi.tools.range.infer_bulk_complete_from_fs` (*datetimes*, *datetime_to_task*, *datetime_to_re*)

Efficiently determines missing datetimes by filesystem listing.

The current implementation works for the common case of a task writing output to a `FileSystemTarget` whose path is built using `strftime` with format like `'...%Y...%m...%d...%H...'`, without custom `complete()` or `exists()`.

(Eventually Luigi could have ranges of completion as first-class citizens. Then this listing business could be factored away/be provided for explicitly in target API or some kind of a history server.)

class `luigi.tools.range.RangeDaily` (**args*, ***kwargs*)

Bases: `luigi.tools.range.RangeDailyBase`

Efficiently produces a contiguous completed range of a daily recurring task that takes a single `DateParameter`.

Falls back to infer it from output filesystem listing to facilitate the common case usage.

Convenient to use even from command line, like:

```
luigi --module your.module RangeDaily --of YourActualTask --start 2014-01-01
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

missing_datetimes (*task_cls*, *finite_datetimes*)

task_namespace = None

class `luigi.tools.range.RangeHourly` (**args*, ***kwargs*)

Bases: `luigi.tools.range.RangeHourlyBase`

Efficiently produces a contiguous completed range of an hourly recurring task that takes a single `DateHourParameter`.

Benefits from `bulk_complete` information to efficiently cover gaps.

Falls back to infer it from output filesystem listing to facilitate the common case usage.

Convenient to use even from command line, like:

```
luigi --module your.module RangeHourly --of YourActualTask --start 2014-01-01T00
```

Constructor to resolve values for all Parameters.

For example, the Task:

```
class MyTask(luigi.Task):  
    count = luigi.IntParameter()
```

can be instantiated as `MyTask(count=10)`.

missing_datetimes (*task_cls, finite_datetimes*)

task_namespace = None

11.3.2 Module contents

Sort of a standard library for doing stuff with Tasks at a somewhat abstract level.

Submodule introduced to stop growing util.py unstructured.

11.4 Indices and tables

- [genindex](#)
- [modindex](#)
- [search](#)

I

`luigi`, 135
`luigi.cmdline`, 97
`luigi.configuration`, 97
`luigi.contrib`, 177
`luigi.contrib.bigquery`, 142
`luigi.contrib.esindex`, 145
`luigi.contrib.ftp`, 149
`luigi.contrib.gcs`, 150
`luigi.contrib.hadoop`, 151
`luigi.contrib.hadoop_jar`, 155
`luigi.contrib.hdfs`, 142
`luigi.contrib.hdfs.abstract_client`, 135
`luigi.contrib.hdfs.clients`, 136
`luigi.contrib.hdfs.config`, 136
`luigi.contrib.hdfs.error`, 137
`luigi.contrib.hdfs.format`, 137
`luigi.contrib.hdfs.hadoopcli_clients`, 138
`luigi.contrib.hdfs.snakebite_client`, 139
`luigi.contrib.hdfs.target`, 141
`luigi.contrib.hive`, 156
`luigi.contrib.mysqlldb`, 159
`luigi.contrib.pig`, 159
`luigi.contrib.pyspark_runner`, 160
`luigi.contrib.rdbms`, 161
`luigi.contrib.redis_store`, 162
`luigi.contrib.redshift`, 162
`luigi.contrib.scalding`, 166
`luigi.contrib.spark`, 167
`luigi.contrib.sparkey`, 171
`luigi.contrib.sqlla`, 172
`luigi.contrib.ssh`, 175
`luigi.contrib.target`, 176
`luigi.contrib.webhdfs`, 177
`luigi.date_interval`, 98
`luigi.db_task_history`, 100
`luigi.deprecate_kwarg`, 101
`luigi.event`, 102
`luigi.file`, 102
`luigi.format`, 103
`luigi.hadoop`, 105
`luigi.hadoop_jar`, 105
`luigi.hdfs`, 105
`luigi.hive`, 105
`luigi.interface`, 105
`luigi.lock`, 108
`luigi.mock`, 108
`luigi.mrrunner`, 109
`luigi.notifications`, 109
`luigi.parameter`, 110
`luigi.postgres`, 113
`luigi.process`, 115
`luigi.rpc`, 115
`luigi.s3`, 116
`luigi.scalding`, 119
`luigi.scheduler`, 119
`luigi.server`, 122
`luigi.target`, 123
`luigi.task`, 125
`luigi.task_history`, 130
`luigi.task_register`, 131
`luigi.task_status`, 131
`luigi.tools`, 182
`luigi.tools.deps`, 177
`luigi.tools.luigi_grep`, 178
`luigi.tools.parse_task`, 178
`luigi.tools.range`, 178
`luigi.util`, 132
`luigi.webhdfs`, 133
`luigi.worker`, 133

A

- `abort()` (`luigi.contrib.hdfs.format.HdfsAtomicWriteDirPipe` method), 53, 137
- `abort()` (`luigi.contrib.hdfs.format.HdfsAtomicWritePipe` method), 53, 137
- `abort()` (`luigi.format.OutputPipeProcessWrapper` method), 103
- `acquire_for()` (in module `luigi.lock`), 108
- `add()` (`luigi.worker.Worker` method), 135
- `add_config_path()` (`luigi.configuration.LuigiConfigParser` class method), 98
- `add_failure()` (`luigi.scheduler.Failures` method), 120
- `add_failure()` (`luigi.scheduler.Task` method), 120
- `add_global_parameters()` (in module `luigi.interface`), 106
- `add_info()` (`luigi.scheduler.Worker` method), 120
- `add_link()` (`luigi.contrib.hadoop.JobTask` method), 69, 153
- `add_task` (`luigi.scheduler.Scheduler` attribute), 119
- `add_task()` (`luigi.rpc.RemoteScheduler` method), 115
- `add_task()` (`luigi.scheduler.CentralPlannerScheduler` method), 121
- `add_task_parameters()` (in module `luigi.interface`), 106
- `add_to_cmdline_parser()` (`luigi.parameter.Parameter` method), 112
- `add_worker()` (`luigi.rpc.RemoteScheduler` method), 116
- `add_worker()` (`luigi.scheduler.CentralPlannerScheduler` method), 121
- `ALL_METHOD_NAMES` (`luigi.contrib.target.CascadingClient` attribute), 92, 176
- `AllRunHandler` (class in `luigi.server`), 122
- `AMBIGUOUS_CLASS` (`luigi.task_register.Register` attribute), 131
- `ApacheHiveCommandClient` (class in `luigi.contrib.hive`), 73, 157
- `app` (`luigi.contrib.spark.PySpark1xJob` attribute), 87, 171
- `app` (`luigi.contrib.spark.PySparkTask` attribute), 84, 169
- `app` (`luigi.contrib.spark.Spark1xJob` attribute), 86, 171
- `app` (`luigi.contrib.spark.SparkSubmitTask` attribute), 83, 168
- `app()` (in module `luigi.server`), 123
- `app_command()` (`luigi.contrib.spark.PySparkTask` method), 85, 169
- `app_command()` (`luigi.contrib.spark.SparkSubmitTask` method), 84, 169
- `app_options()` (`luigi.contrib.spark.Spark1xBackwardCompat` method), 86, 170
- `app_options()` (`luigi.contrib.spark.SparkSubmitTask` method), 83, 168
- `apply_async()` (`luigi.worker.SingleProcessPool` method), 133
- `archives` (`luigi.contrib.spark.SparkSubmitTask` attribute), 84, 168
- `ArgParseInterface` (class in `luigi.interface`), 106
- `args()` (`luigi.contrib.hadoop_jar.HadoopJarJobTask` method), 72, 156
- `args()` (`luigi.contrib.scalding.ScaldingJobTask` method), 83, 167
- `assistant` (`luigi.interface.core` attribute), 106
- `assistant` (`luigi.scheduler.Worker` attribute), 120
- `AsyncCompletionException`, 133
- `atomic_file` (class in `luigi.file`), 102
- `atomic_output()` (`luigi.contrib.hadoop_jar.HadoopJarJobTask` method), 72, 156
- `atomic_output()` (`luigi.contrib.scalding.ScaldingJobTask` method), 83, 167
- `AtomicFtpFile` (class in `luigi.contrib.ftp`), 65, 149
- `AtomicGCSFile` (class in `luigi.contrib.gcs`), 66, 151
- `AtomicLocalFile` (class in `luigi.target`), 125
- `AtomicRemoteFileWriter` (class in `luigi.contrib.ssh`), 91, 176
- `AtomicS3File` (class in `luigi.s3`), 117
- `AtomicWebHdfsFile` (class in `luigi.contrib.webhdfs`), 92, 177
- `attach()` (in module `luigi.contrib.hadoop`), 70, 154
- `aws_access_key_id` (`luigi.contrib.redshift.S3CopyToTable` attribute), 78, 163
- `aws_secret_access_key` (`luigi.contrib.redshift.S3CopyToTable` attribute), 79, 163

B

BaseHadoopJobTask (class in `luigi.contrib.hadoop`), 67, 151

BaseTaskHistoryHandler (class in `luigi.server`), 122

BaseWrapper (class in `luigi.format`), 103

BATCH (`luigi.contrib.bigquery.QueryMode` attribute), 58, 142

batch_counter_default (`luigi.contrib.hadoop.BaseHadoopJobTask` attribute), 67, 151

BigqueryClient (class in `luigi.contrib.bigquery`), 59, 143

BigqueryLoadTask (class in `luigi.contrib.bigquery`), 60, 144

BigqueryRunQueryTask (class in `luigi.contrib.bigquery`), 61, 145

BigqueryTarget (class in `luigi.contrib.bigquery`), 60, 144

BooleanParameter (class in `luigi.parameter`), 113

BoolParameter (class in `luigi.parameter`), 113

BQDataset (class in `luigi.contrib.bigquery`), 59, 143

BQTable (class in `luigi.contrib.bigquery`), 59, 143

BROKEN_TASK (`luigi.event.Event` attribute), 102

build() (in module `luigi.interface`), 108

build_job_jar() (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166

bulk_complete() (`luigi.task.MixinNaiveBulkComplete` class method), 129

bulk_complete() (`luigi.task.Task` class method), 128

BulkCompleteNotImplementedError, 126

ByIdHandler (class in `luigi.server`), 123

ByNameHandler (class in `luigi.server`), 122

ByParamsHandler (class in `luigi.server`), 123

bytes_per_reducer (`luigi.contrib.hive.HiveQueryTask` attribute), 73, 157

Bzip2Format (class in `luigi.format`), 105

C

call_check() (`luigi.contrib.hdfs.hadoopcli_clients.HdfsClient` static method), 54, 138

can_disable() (`luigi.scheduler.Task` method), 120

CascadingClient (class in `luigi.contrib.target`), 92, 176

CentralPlannerScheduler (class in `luigi.scheduler`), 121

ChainFormat (class in `luigi.format`), 104

check_complete() (in module `luigi.worker`), 134

check_output() (`luigi.contrib.ssh.RemoteContext` method), 91, 175

check_pid() (in module `luigi.process`), 115

chmod() (`luigi.contrib.hdfs.abstract_client.HdfsFileSystem` method), 52, 135

chmod() (`luigi.contrib.hdfs.hadoopcli_clients.HdfsClient` method), 54, 138

chmod() (`luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient` method), 56, 140

chown() (`luigi.contrib.hdfs.abstract_client.HdfsFileSystem` method), 52, 135

chown() (`luigi.contrib.hdfs.hadoopcli_clients.HdfsClient` method), 54, 138

chown() (`luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient` method), 56, 140

chunk_size (`luigi.contrib.esindex.CopyToIndex` attribute), 64, 148

chunk_size (`luigi.contrib.sqla.CopyToTable` attribute), 90, 174

clear() (`luigi.mock.MockFileSystem` method), 109

clear() (`luigi.scheduler.Failures` method), 120

clear_instance_cache() (`luigi.task_register.Register` class method), 131

client (`luigi.contrib.hdfs.config.hdfs` attribute), 52, 136

client_version (`luigi.contrib.hdfs.config.hdfs` attribute), 52, 136

clone() (`luigi.task.Task` method), 127

close() (`luigi.contrib.hdfs.format.HdfsAtomicWriteDirPipe` method), 53, 137

close() (`luigi.contrib.hdfs.format.HdfsAtomicWritePipe` method), 53, 137

close() (`luigi.contrib.ssh.AtomicRemoteFileWriter` method), 91, 176

close() (`luigi.contrib.webhdfs.ReadableWebHdfsFile` method), 92, 177

close() (`luigi.format.InputPipeProcessWrapper` method), 103

close() (`luigi.format.OutputPipeProcessWrapper` method), 103

close() (`luigi.s3.ReadableS3File` method), 117

close() (`luigi.target.AtomicLocalFile` method), 125

column_separator (`luigi.contrib.rdbms.CopyToTable` attribute), 77, 161

column_separator (`luigi.contrib.sqla.CopyToTable` attribute), 90, 174

columns (`luigi.contrib.rdbms.CopyToTable` attribute), 77, 161

columns (`luigi.contrib.sqla.CopyToTable` attribute), 90, 174

combiner (`luigi.contrib.hadoop.JobTask` attribute), 69, 153

command (`luigi.contrib.hdfs.config.hadoopcli` attribute), 53, 136

common_params() (in module `luigi.util`), 132

CompatibleHdfsFormat (class in `luigi.contrib.hdfs.format`), 54, 138

complete() (`luigi.task.Task` method), 127

complete() (`luigi.task WrapperTask` method), 129

COMPLETE_COUNT (`luigi.tools.range.RangeEvent` attribute), 94, 179

COMPLETE_FRACTION (`luigi.tools.range.RangeEvent` attribute), 94, 179

conf (`luigi.contrib.spark.SparkSubmitTask` attribute), 84, 168

Config (class in `luigi.task`), 129

[connect\(\)](#) ([luigi.contrib.mysql.MySqlTarget](#) method), [75](#), [159](#)
[connect\(\)](#) ([luigi.postgres.PostgresTarget](#) method), [114](#)
[connect_args](#) ([luigi.contrib.sqla.CopyToTable](#) attribute), [90](#), [174](#)
[connect_args](#) ([luigi.contrib.sqla.SQLAlchemyTarget](#) attribute), [89](#), [174](#)
[connection_reset_wait_seconds](#) ([luigi.contrib.redshift.KillOpenRedshiftSessions](#) attribute), [81](#), [165](#)
[connection_string\(\)](#) ([luigi.contrib.sqla.CopyToTable](#) method), [90](#), [174](#)
[copies](#) (class in [luigi.util](#)), [132](#)
[copy\(\)](#) ([luigi.contrib.bigquery.BigqueryClient](#) method), [60](#), [144](#)
[copy\(\)](#) ([luigi.contrib.gcs.GCSClient](#) method), [66](#), [150](#)
[copy\(\)](#) ([luigi.contrib.hdfs.abstract_client.HdfsFileSystem](#) method), [52](#), [135](#)
[copy\(\)](#) ([luigi.contrib.hdfs.hadoopcli_clients.HdfsClient](#) method), [54](#), [138](#)
[copy\(\)](#) ([luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient](#) method), [56](#), [140](#)
[copy\(\)](#) ([luigi.contrib.rdbms.CopyToTable](#) method), [77](#), [162](#)
[copy\(\)](#) ([luigi.contrib.redshift.S3CopyJSONToTable](#) method), [80](#), [164](#)
[copy\(\)](#) ([luigi.contrib.redshift.S3CopyToTable](#) method), [79](#), [163](#)
[copy\(\)](#) ([luigi.contrib.sqla.CopyToTable](#) method), [90](#), [175](#)
[copy\(\)](#) ([luigi.file.LocalTarget](#) method), [103](#)
[copy\(\)](#) ([luigi.postgres.CopyToTable](#) method), [115](#)
[copy\(\)](#) ([luigi.s3.S3Client](#) method), [116](#)
[copy_json_options](#) ([luigi.contrib.redshift.S3CopyJSONToTable](#) attribute), [80](#), [164](#)
[copy_options](#) ([luigi.contrib.redshift.S3CopyToTable](#) attribute), [79](#), [163](#)
[CopyToIndex](#) (class in [luigi.contrib.esindex](#)), [63](#), [147](#)
[CopyToTable](#) (class in [luigi.contrib.rdbms](#)), [76](#), [161](#)
[CopyToTable](#) (class in [luigi.contrib.sqla](#)), [90](#), [174](#)
[CopyToTable](#) (class in [luigi.postgres](#)), [114](#)
[core](#) (class in [luigi.interface](#)), [105](#)
[count\(\)](#) ([luigi.contrib.hdfs.abstract_client.HdfsFileSystem](#) method), [52](#), [135](#)
[count\(\)](#) ([luigi.contrib.hdfs.hadoopcli_clients.HdfsClient](#) method), [54](#), [138](#)
[count\(\)](#) ([luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient](#) method), [56](#), [140](#)
[count_uniques](#) ([luigi.worker.worker](#) attribute), [134](#)
[counter](#) ([luigi.parameter.Parameter](#) attribute), [111](#)
[create_disposition](#) ([luigi.contrib.bigquery.BigqueryRunQueryTask](#) attribute), [61](#), [145](#)
[create_hadoopcli_client\(\)](#) (in module [luigi.contrib.hdfs.hadoopcli_clients](#)), [54](#), [138](#)
[CREATE_IF_NEEDED](#) ([luigi.contrib.bigquery.CreateDisposition](#) attribute), [58](#), [142](#)
[create_index\(\)](#) ([luigi.contrib.esindex.CopyToIndex](#) method), [64](#), [148](#)
[create_local_scheduler\(\)](#) ([luigi.interface.WorkerSchedulerFactory](#) method), [106](#)
[create_marker_index\(\)](#) ([luigi.contrib.esindex.ElasticsearchTarget](#) method), [63](#), [147](#)
[create_marker_table\(\)](#) ([luigi.contrib.mysql.MySqlTarget](#) method), [75](#), [159](#)
[create_marker_table\(\)](#) ([luigi.contrib.sqla.SQLAlchemyTarget](#) method), [89](#), [174](#)
[create_marker_table\(\)](#) ([luigi.postgres.PostgresTarget](#) method), [114](#)
[CREATE_NEVER](#) ([luigi.contrib.bigquery.CreateDisposition](#) attribute), [58](#), [142](#)
[create_packages_archive\(\)](#) (in module [luigi.contrib.hadoop](#)), [70](#), [154](#)
[create_remote_scheduler\(\)](#) ([luigi.interface.WorkerSchedulerFactory](#) method), [106](#)
[create_subprocess\(\)](#) ([luigi.format.InputPipeProcessWrapper](#) method), [103](#)
[create_table\(\)](#) ([luigi.contrib.rdbms.CopyToTable](#) method), [77](#), [161](#)
[create_table\(\)](#) ([luigi.contrib.redshift.S3CopyToTable](#) method), [79](#), [163](#)
[create_table\(\)](#) ([luigi.contrib.sqla.CopyToTable](#) method), [90](#), [175](#)
[create_worker\(\)](#) ([luigi.interface.WorkerSchedulerFactory](#) method), [106](#)
[CreateDisposition](#) (class in [luigi.contrib.bigquery](#)), [58](#), [142](#)
[CSV](#) ([luigi.contrib.bigquery.SourceFormat](#) attribute), [58](#), [142](#)
[Custom](#) (class in [luigi.date_interval](#)), [100](#)

D

[daemonize\(\)](#) (in module [luigi.process](#)), [115](#)
[data_interchange_format](#) ([luigi.contrib.hadoop.JobTask](#) attribute), [69](#), [153](#)
[database](#) ([luigi.contrib.hive.ExternalHiveTask](#) attribute), [74](#), [158](#)
[database](#) ([luigi.contrib.rdbms.CopyToTable](#) attribute), [77](#), [161](#)
[database](#) ([luigi.contrib.redshift.KillOpenRedshiftSessions](#) attribute), [81](#), [165](#)
[dataset](#) ([luigi.contrib.bigquery.BQTable](#) attribute), [59](#), [143](#)
[dataset_exists\(\)](#) ([luigi.contrib.bigquery.BigqueryClient](#) method), [59](#), [143](#)
[dataset_id](#) ([luigi.contrib.bigquery.BQDataset](#) attribute), [59](#), [143](#)

- DATASTORE_BACKUP (luigi.contrib.bigquery.SourceFormat attribute), 58, 142
- Date (class in luigi.date_interval), 99
- date_format (luigi.parameter.DateHourParameter attribute), 112
- date_format (luigi.parameter.DateMinuteParameter attribute), 112
- DateHourParameter (class in luigi.parameter), 112
- DateInterval (class in luigi.date_interval), 98
- DateIntervalParameter (class in luigi.parameter), 113
- DateMinuteParameter (class in luigi.parameter), 112
- DateParameter (class in luigi.parameter), 112
- dates() (luigi.date_interval.DateInterval method), 99
- datetime_to_parameter() (luigi.tools.range.RangeBase method), 95, 179
- datetime_to_parameter() (luigi.tools.range.RangeDailyBase method), 95, 180
- datetime_to_parameter() (luigi.tools.range.RangeHourlyBase method), 96, 181
- days_back (luigi.tools.range.RangeDailyBase attribute), 95, 180
- days_forward (luigi.tools.range.RangeDailyBase attribute), 95, 180
- DbTaskHistory (class in luigi.db_task_history), 100
- DefaultHadoopJobRunner (class in luigi.contrib.hadoop), 68, 152
- DELAY (luigi.tools.range.RangeEvent attribute), 94, 179
- delegates() (in module luigi.util), 132
- delete_dataset() (luigi.contrib.bigquery.BigqueryClient method), 59, 143
- delete_index() (luigi.contrib.esindex.CopyToIndex method), 64, 148
- delete_table() (luigi.contrib.bigquery.BigqueryClient method), 59, 143
- dep_graph() (luigi.rpc.RemoteScheduler method), 115
- dep_graph() (luigi.scheduler.CentralPlannerScheduler method), 121
- DEPENDENCY_DISCOVERED (luigi.event.Event attribute), 102
- dependency_jars() (luigi.contrib.spark.Spark1xBackwardCompatibility method), 86, 170
- DEPENDENCY_MISSING (luigi.event.Event attribute), 102
- DEPENDENCY_PRESENT (luigi.event.Event attribute), 102
- deploy_mode (luigi.contrib.spark.PySparkTask attribute), 84, 169
- deploy_mode (luigi.contrib.spark.SparkSubmitTask attribute), 83, 168
- deprecate_kwarg() (in module luigi.deprecate_kwarg), 101
- deps() (luigi.contrib.hadoop.BaseHadoopJobTask method), 67, 151
- deps() (luigi.contrib.spark.SparkJob method), 85, 170
- deps() (luigi.task.Task method), 128
- DequeQueue (class in luigi.worker), 133
- dereference() (in module luigi.contrib.hadoop), 70, 154
- deserialize (luigi.contrib.hadoop.JobTask attribute), 69, 153
- dfs_paths() (in module luigi.tools.deps), 93, 177
- disable_failures (luigi.scheduler.scheduler attribute), 119
- disable_hard_timeout (luigi.scheduler.scheduler attribute), 119
- disable_instance_cache() (luigi.task_register.Register class method), 131
- disable_persist (luigi.scheduler.scheduler attribute), 119
- disable_window (luigi.scheduler.scheduler attribute), 119
- disabled (luigi.task.Task attribute), 126
- do_truncate_table() (luigi.contrib.redshift.S3CopyToTable method), 79, 163
- doc_type (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
- docs() (luigi.contrib.esindex.CopyToIndex method), 64, 148
- does_table_exist() (luigi.contrib.redshift.S3CopyToTable method), 79, 163
- download() (luigi.contrib.gcs.GCSCClient method), 66, 150
- download() (luigi.contrib.webhdfs.WebHdfsClient method), 93, 177
- driver_class_path (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- driver_cores (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- driver_java_options (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- driver_library_path (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- driver_memory (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- dump() (luigi.contrib.hadoop.JobTask method), 69, 153
- dump() (luigi.scheduler.CentralPlannerScheduler method), 121
- dump() (luigi.scheduler.SimpleTaskState method), 120
- DuplicateParameterException, 110
- DynamicArgParseInterface (class in luigi.interface), 107

E

- echo (luigi.contrib.sqla.CopyToTable attribute), 90, 174
- effective_user (luigi.contrib.hdfs.config.hdfs attribute), 52, 136
- ElasticsearchTarget (class in luigi.contrib.esindex), 62, 146
- email_type() (in module luigi.notifications), 109
- engine (luigi.contrib.sqla.SQLAlchemyTarget attribute), 89, 174

- engine (luigi.contrib.sqla.SQLAlchemyTarget.Connection attribute), 89, 174
- ensure_hist_size() (luigi.contrib.esindex.ElasticsearchTarget method), 63, 147
- entry_class (luigi.contrib.spark.Spark1xJob attribute), 86, 171
- entry_class (luigi.contrib.spark.SparkSubmitTask attribute), 83, 168
- error_task_names() (in module luigi.interface), 106
- Event (class in luigi.event), 102
- event_handler() (luigi.task.Task class method), 127
- event_name (luigi.db_task_history.TaskEvent attribute), 101
- events (luigi.db_task_history.TaskRecord attribute), 101
- executor_cores (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- executor_memory (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- exists() (luigi.contrib.bigquery.BigqueryTarget method), 60, 144
- exists() (luigi.contrib.esindex.ElasticsearchTarget method), 63, 147
- exists() (luigi.contrib.ftp.RemoteFileSystem method), 65, 149
- exists() (luigi.contrib.ftp.RemoteTarget method), 66, 150
- exists() (luigi.contrib.gcs.GCSClient method), 66, 150
- exists() (luigi.contrib.gcs.GCSFlagTarget method), 67, 151
- exists() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient method), 54, 138
- exists() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClientApache1 method), 55, 139
- exists() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 55, 139
- exists() (luigi.contrib.hive.HivePartitionTarget method), 74, 158
- exists() (luigi.contrib.hive.HiveTableTarget method), 74, 158
- exists() (luigi.contrib.mysql.db.MySqlTarget method), 75, 159
- exists() (luigi.contrib.redis_store.RedisTarget method), 78, 162
- exists() (luigi.contrib.sqla.SQLAlchemyTarget method), 89, 174
- exists() (luigi.contrib.ssh.RemoteFileSystem method), 91, 176
- exists() (luigi.contrib.webhdfs.WebHdfsClient method), 92, 177
- exists() (luigi.file.LocalFileSystem method), 102
- exists() (luigi.mock.MockFileSystem method), 108
- exists() (luigi.mock.MockTarget method), 109
- exists() (luigi.postgres.PostgresTarget method), 114
- exists() (luigi.s3.S3Client method), 116
- exists() (luigi.s3.S3FlagTarget method), 118
- exists() (luigi.target.FileSystem method), 124
- exists() (luigi.target.FileSystemTarget method), 125
- exists() (luigi.target.Target method), 123
- ExternalHiveTask (class in luigi.contrib.hive), 74, 158
- externalize() (in module luigi.task), 129
- ExternalTask (class in luigi.task), 129
- extra_elasticsearch_args (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
- extra_files() (luigi.contrib.hadoop.JobTask method), 69, 153
- extra_jars() (luigi.contrib.scalding.ScaldingJobTask method), 83, 167
- extra_modules() (luigi.contrib.hadoop.JobTask method), 69, 153
- extract_packages_archive() (luigi.mrrunner.Runner method), 109
- ## F
- FAILURE (luigi.event.Event attribute), 102
- Failures (class in luigi.scheduler), 120
- family (luigi.tools.deps.upstream attribute), 93, 178
- fetch_error() (luigi.rpc.RemoteScheduler method), 116
- fetch_error() (luigi.scheduler.CentralPlannerScheduler method), 122
- fetch_task_failures() (in module luigi.contrib.hadoop), 70, 154
- File (class in luigi.file), 103
- FileAlreadyExists, 123
- FileNotFoundException, 116
- files (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- FileSystem (class in luigi.target), 124
- FileSystemException, 123
- FileSystemTarget (class in luigi.target), 124
- FileWrapper (class in luigi.format), 103
- final_combiner (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 67, 152
- final_mapper (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 67, 152
- final_reducer (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 67, 152
- find_all_by_name() (luigi.db_task_history.DbTaskHistory method), 100
- find_all_by_parameters() (luigi.db_task_history.DbTaskHistory method), 100
- find_all_events() (luigi.db_task_history.DbTaskHistory method), 100
- find_all_runs() (luigi.db_task_history.DbTaskHistory method), 100
- find_deps() (in module luigi.tools.deps), 93, 178
- find_deps_cli() (in module luigi.tools.deps), 93, 178
- find_latest_runs() (luigi.db_task_history.DbTaskHistory method), 100

`find_task_by_id()` (`luigi.db_task_history.DbTaskHistory` method), 100

`finish()` (`luigi.contrib.hadoop.HadoopJobRunner` method), 68, 152

`finite_datetimes()` (`luigi.tools.range.RangeBase` method), 95, 180

`finite_datetimes()` (`luigi.tools.range.RangeDailyBase` method), 95, 180

`finite_datetimes()` (`luigi.tools.range.RangeHourlyBase` method), 96, 181

`fix_paths()` (in module `luigi.contrib.hadoop_jar`), 71, 155

`fix_time()` (in module `luigi.scheduler`), 119

`flag` (`luigi.s3.S3FlagTask` attribute), 118

`flatten()` (in module `luigi.contrib.hadoop`), 70, 154

`flatten()` (in module `luigi.task`), 130

`flatten_output()` (in module `luigi.task`), 130

`FloatParameter` (class in `luigi.parameter`), 112

`fn` (`luigi.file.LocalTarget` attribute), 103

`folder_paths` (`luigi.contrib.redshift.RedshiftManifestTask` attribute), 80, 165

`Format` (class in `luigi.format`), 104

`from_bqtable()` (`luigi.contrib.bigquery.BigqueryTarget` class method), 60, 144

`from_date()` (`luigi.date_interval.Date` class method), 99

`from_date()` (`luigi.date_interval.DateInterval` class method), 99

`from_date()` (`luigi.date_interval.Month` class method), 99

`from_date()` (`luigi.date_interval.Week` class method), 99

`from_date()` (`luigi.date_interval.Year` class method), 100

`from_str_params()` (`luigi.task.Task` class method), 127

`from_utc()` (in module `luigi.server`), 122

`fs` (`luigi.contrib.ftp.AtomicFtpFile` attribute), 65, 149

`fs` (`luigi.contrib.ftp.RemoteTarget` attribute), 65, 149

`fs` (`luigi.contrib.gcs.GCSFlagTarget` attribute), 67, 151

`fs` (`luigi.contrib.gcs.GCSTarget` attribute), 67, 151

`fs` (`luigi.contrib.hdfs.target.HdfsTarget` attribute), 57, 141

`fs` (`luigi.contrib.ssh.AtomicRemoteFileWriter` attribute), 92, 176

`fs` (`luigi.contrib.ssh.RemoteTarget` attribute), 92, 176

`fs` (`luigi.contrib.webhdfs.WebHdfsTarget` attribute), 92, 177

`fs` (`luigi.file.LocalTarget` attribute), 102

`fs` (`luigi.mock.MockTarget` attribute), 109

`fs` (`luigi.s3.S3FlagTarget` attribute), 117

`fs` (`luigi.s3.S3Target` attribute), 117

`fs` (`luigi.target.FileSystemTarget` attribute), 125

G

`GCSCClient` (class in `luigi.contrib.gcs`), 66, 150

`GCSFlagTarget` (class in `luigi.contrib.gcs`), 67, 151

`GCSTarget` (class in `luigi.contrib.gcs`), 67, 151

`generate_email()` (in module `luigi.notifications`), 109

`generate_tmp_path()` (`luigi.file.atomic_file` method), 102

`generate_tmp_path()` (`luigi.target.AtomicLocalFile` method), 125

`get()` (`luigi.configuration.LuigiConfigParser` method), 98

`get()` (`luigi.contrib.ftp.RemoteFileSystem` method), 65, 149

`get()` (`luigi.contrib.ftp.RemoteTarget` method), 66, 150

`get()` (`luigi.contrib.hdfs.abstract_client.HdfsFileSystem` method), 52, 135

`get()` (`luigi.contrib.hdfs.hadoopcli_clients.HdfsClient` method), 55, 138

`get()` (`luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient` method), 56, 140

`get()` (`luigi.contrib.ssh.RemoteFileSystem` method), 91, 176

`get()` (`luigi.contrib.ssh.RemoteTarget` method), 92, 176

`get()` (`luigi.server.AllRunHandler` method), 122

`get()` (`luigi.server.ByIdHandler` method), 123

`get()` (`luigi.server.ByNameHandler` method), 122

`get()` (`luigi.server.ByParamsHandler` method), 123

`get()` (`luigi.server.RecentRunHandler` method), 122

`get()` (`luigi.server.RootPathHandler` method), 123

`get()` (`luigi.server.RPCHandler` method), 122

`get()` (`luigi.server.SelectedRunHandler` method), 122

`get()` (`luigi.server.StaticFileHandler` method), 123

`get()` (`luigi.worker.DequeQueue` method), 133

`get_active_tasks()` (`luigi.scheduler.SimpleTaskState` method), 120

`get_active_workers()` (`luigi.scheduler.SimpleTaskState` method), 121

`get_all_data()` (`luigi.mock.MockFileSystem` method), 108

`get_all_params()` (`luigi.task_register.Register` class method), 131

`get_assistants()` (`luigi.scheduler.SimpleTaskState` method), 121

`get_autoconfig_client()` (in module `luigi.contrib.hdfs.clients`), 52, 136

`get_bite()` (`luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient` method), 55, 139

`get_build_dir()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166

`get_config()` (in module `luigi.configuration`), 98

`get_config()` (`luigi.contrib.webhdfs.WebHdfsClient` method), 92, 177

`get_configured_hadoop_version()` (in module `luigi.contrib.hdfs.config`), 53, 137

`get_configured_hdfs_client()` (in module `luigi.contrib.hdfs.config`), 53, 137

`get_data()` (`luigi.mock.MockFileSystem` method), 108

`get_default_client()` (in module `luigi.contrib.hive`), 73, 157

`get_default_format()` (in module `luigi.format`), 105

`get_environment()` (`luigi.contrib.spark.SparkSubmitTask` method), 84, 168

`get_extra_files()` (in module `luigi.contrib.hadoop`), 70,

- 155
- `get_global_parameters()` (in module `luigi.interface`), 106
- `get_hive_syntax()` (in module `luigi.contrib.hive`), 72, 156
- `get_info()` (in module `luigi.lock`), 108
- `get_job_class()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_key()` (`luigi.s3.S3Client` method), 116
- `get_libjars()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_log_format()` (in module `luigi.process`), 115
- `get_necessary_tasks()` (`luigi.scheduler.SimpleTaskState` method), 121
- `get_param_values()` (`luigi.task.Task` class method), 127
- `get_params()` (`luigi.task.Task` class method), 127
- `get_pending_tasks()` (`luigi.scheduler.SimpleTaskState` method), 120
- `get_pending_tasks()` (`luigi.scheduler.Worker` method), 120
- `get_previous_completed()` (in module `luigi.util`), 133
- `get_provided_jars()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_running_tasks()` (`luigi.scheduler.SimpleTaskState` method), 120
- `get_scala_jars()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_scalding_core()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_scalding_jars()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_spool_handler()` (in module `luigi.process`), 115
- `get_task()` (`luigi.scheduler.SimpleTaskState` method), 121
- `get_task_cls()` (`luigi.task_register.Register` class method), 131
- `get_task_parameters()` (in module `luigi.interface`), 106
- `get_task_requires()` (in module `luigi.tools.deps`), 93, 177
- `get_template_path()` (`luigi.server.BaseTaskHistoryHandler` method), 122
- `get_tmp_job_jar()` (`luigi.contrib.scalding.ScaldingJobRunner` method), 82, 166
- `get_work` (`luigi.scheduler.Scheduler` attribute), 119
- `get_work()` (`luigi.rpc.RemoteScheduler` method), 115
- `get_work()` (`luigi.scheduler.CentralPlannerScheduler` method), 121
- `get_worker()` (`luigi.scheduler.SimpleTaskState` method), 121
- `get_worker_ids()` (`luigi.scheduler.SimpleTaskState` method), 121
- `getboolean()` (`luigi.configuration.LuigiConfigParser` method), 98
- `getfloat()` (`luigi.configuration.LuigiConfigParser` method), 98
- `getint()` (`luigi.configuration.LuigiConfigParser` method), 98
- `getintdict()` (`luigi.configuration.LuigiConfigParser` method), 98
- `getmerge()` (`luigi.contrib.hdfs.hadoopcli_clients.HdfsClient` method), 55, 138
- `getpaths()` (in module `luigi.task`), 130
- `getpcmd()` (in module `luigi.lock`), 108
- `glob_exists()` (`luigi.contrib.hdfs.target.HdfsTarget` method), 57, 141
- `graph()` (`luigi.rpc.RemoteScheduler` method), 115
- `graph()` (`luigi.scheduler.CentralPlannerScheduler` method), 121
- `graph_url` (`luigi.tools.luigi_grep.LuigiGrep` attribute), 93, 178
- `group()` (`luigi.contrib.hadoop.LocalJobRunner` method), 70, 154
- `GzipFormat` (class in `luigi.format`), 104
- ## H
- `hadoop` (class in `luigi.contrib.hadoop`), 70, 155
- `hadoop_conf_dir` (`luigi.contrib.spark.SparkSubmitTask` attribute), 84, 168
- `hadoopcli` (class in `luigi.contrib.hdfs.config`), 52, 136
- `HadoopJarJobError`, 71, 155
- `HadoopJarJobRunner` (class in `luigi.contrib.hadoop_jar`), 71, 155
- `HadoopJarJobTask` (class in `luigi.contrib.hadoop_jar`), 71, 155
- `HadoopJobError`, 68, 152
- `HadoopJobRunner` (class in `luigi.contrib.hadoop`), 68, 152
- `HadoopRunContext` (class in `luigi.contrib.hadoop`), 68, 152
- `has_excessive_failures()` (`luigi.scheduler.Task` method), 120
- `has_task()` (`luigi.scheduler.SimpleTaskState` method), 121
- `has_task_value()` (`luigi.parameter.Parameter` method), 111
- `has_value` (`luigi.parameter.Parameter` attribute), 111
- `hdfs` (class in `luigi.contrib.hdfs.config`), 52, 136
- `hdfs_reader()` (`luigi.contrib.hdfs.format.CompatibleHdfsFormat` method), 54, 138
- `hdfs_reader()` (`luigi.contrib.hdfs.format.PlainDirFormat` method), 54, 138
- `hdfs_reader()` (`luigi.contrib.hdfs.format.PlainFormat` method), 54, 137
- `hdfs_writer()` (`luigi.contrib.hdfs.format.CompatibleHdfsFormat` method), 54, 138
- `hdfs_writer()` (`luigi.contrib.hdfs.format.PlainDirFormat` method), 54, 138
- `hdfs_writer()` (`luigi.contrib.hdfs.format.PlainFormat` method), 54, 137
- `HdfsAtomicWriteDirPipe` (class in `luigi.contrib.hdfs.format`), 53, 137
- `HdfsAtomicWritePipe` (class in `luigi.contrib.hdfs.format`), 53, 137

- HdfsClient (class in luigi.contrib.hdfs.hadoopcli_clients), 54, 138
- HdfsClientApache1 (class in luigi.contrib.hdfs.hadoopcli_clients), 139
- HdfsClientCdh3 (class in luigi.contrib.hdfs.hadoopcli_clients), 139
- HDFSClientError, 53, 137
- HdfsFileSystem (class in luigi.contrib.hdfs.abstract_client), 51, 135
- HdfsReadPipe (class in luigi.contrib.hdfs.format), 53, 137
- HdfsTarget (class in luigi.contrib.hdfs.target), 57, 141
- HiveClient (class in luigi.contrib.hive), 72, 156
- HiveCommandClient (class in luigi.contrib.hive), 72, 156
- HiveCommandError, 72, 156
- hiveconfs() (luigi.contrib.hive.HiveQueryTask method), 73, 158
- HivePartitionTarget (class in luigi.contrib.hive), 74, 158
- HiveQueryRunner (class in luigi.contrib.hive), 74, 158
- HiveQueryTask (class in luigi.contrib.hive), 73, 157
- hiverc() (luigi.contrib.hive.HiveQueryTask method), 73, 157
- HiveTableTarget (class in luigi.contrib.hive), 74, 158
- HiveThriftContext (class in luigi.contrib.hive), 73, 157
- host (luigi.contrib.esindex.CopyToIndex attribute), 63, 147
- host (luigi.contrib.rdbms.CopyToTable attribute), 77, 161
- host (luigi.contrib.redshift.KillOpenRedshiftSessions attribute), 81, 165
- host (luigi.db_task_history.TaskRecord attribute), 101
- hours() (luigi.date_interval.DateInterval method), 99
- hours_back (luigi.tools.range.RangeHourlyBase attribute), 96, 181
- hours_forward (luigi.tools.range.RangeHourlyBase attribute), 96, 181
- http_auth (luigi.contrib.esindex.CopyToIndex attribute), 63, 147
- |
- id (luigi.db_task_history.TaskEvent attribute), 101
- id (luigi.db_task_history.TaskRecord attribute), 101
- id_to_name_and_params() (in module luigi.task), 126
- id_to_name_and_params() (in module luigi.tools.parse_task), 94, 178
- inactivate_tasks() (luigi.scheduler.SimpleTaskState method), 121
- inactivate_workers() (luigi.scheduler.SimpleTaskState method), 121
- incr_counter() (luigi.contrib.hadoop.JobTask method), 69, 153
- index (luigi.contrib.esindex.CopyToIndex attribute), 63, 147
- infer_bulk_complete_from_fs() (in module luigi.tools.range), 96, 181
- inherits (class in luigi.util), 132
- init_combiner() (luigi.contrib.hadoop.JobTask method), 69, 153
- init_copy() (luigi.contrib.rdbms.CopyToTable method), 77, 161
- init_hadoop() (luigi.contrib.hadoop.BaseHadoopJobTask method), 67, 152
- init_local() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
- init_mapper() (luigi.contrib.hadoop.JobTask method), 69, 153
- init_reducer() (luigi.contrib.hadoop.JobTask method), 69, 153
- initialize() (luigi.server.BaseTaskHistoryHandler method), 122
- initialize() (luigi.server.RPCHandler method), 122
- initialized() (luigi.task.Task method), 127
- input (luigi.contrib.hdfs.format.PlainDirFormat attribute), 54, 138
- input (luigi.contrib.hdfs.format.PlainFormat attribute), 53, 137
- input (luigi.format.Bzip2Format attribute), 105
- input (luigi.format.GzipFormat attribute), 105
- input (luigi.format.NewlineFormat attribute), 104
- input (luigi.format.TextFormat attribute), 104
- input() (luigi.contrib.spark.SparkJob method), 85, 170
- input() (luigi.task.Task method), 128
- input_hadoop() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
- input_local() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
- input_local() (luigi.contrib.spark.SparkJob method), 85, 170
- InputPipeProcessWrapper (class in luigi.format), 103
- instance() (luigi.configuration.LuigiConfigParser class method), 98
- INTERACTIVE (luigi.contrib.bigquery.QueryMode attribute), 58, 142
- Interface (class in luigi.interface), 106
- internal_reader() (luigi.contrib.hadoop.JobTask method), 69, 153
- internal_serialize (luigi.contrib.hadoop.JobTask attribute), 69, 153
- internal_writer() (luigi.contrib.hadoop.JobTask method), 69, 153
- IntParameter (class in luigi.parameter), 112
- InvalidDeleteException, 66, 116, 150
- inverse_dep_graph() (luigi.rpc.RemoteScheduler method), 115
- inverse_dep_graph() (luigi.scheduler.CentralPlannerScheduler method), 122
- is_dir() (luigi.s3.S3Client method), 117

is_trivial_worker() (luigi.scheduler.Worker method), 120
 is_writable() (luigi.contrib.hdfs.target.HdfsTarget method), 58, 142
 isdir() (luigi.contrib.gcs.GCSClient method), 66, 150
 isdir() (luigi.contrib.ssh.RemoteFileSystem method), 91, 176
 isdir() (luigi.file.LocalFileSystem method), 102
 isdir() (luigi.mock.MockFileSystem method), 109
 isdir() (luigi.s3.S3Client method), 117
 isdir() (luigi.target.FileSystem method), 124

J

jar() (luigi.contrib.hadoop_jar.HadoopJarJobTask method), 71, 156
 jar() (luigi.contrib.scalding.ScaldingJobTask method), 82, 167
 jar() (luigi.contrib.spark.Spark1xJob method), 86, 171
 jar() (luigi.contrib.spark.SparkJob method), 85, 170
 jars (luigi.contrib.spark.Spark1xBackwardCompat attribute), 86, 170
 jars (luigi.contrib.spark.SparkSubmitTask attribute), 83, 168
 job_args() (luigi.contrib.scalding.ScaldingJobTask method), 83, 167
 job_args() (luigi.contrib.spark.Spark1xBackwardCompat method), 86, 170
 job_args() (luigi.contrib.spark.SparkJob method), 85, 170
 job_class() (luigi.contrib.scalding.ScaldingJobTask method), 83, 167
 job_class() (luigi.contrib.spark.Spark1xJob method), 86, 171
 job_class() (luigi.contrib.spark.SparkJob method), 85, 170
 job_runner() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
 job_runner() (luigi.contrib.hadoop.JobTask method), 69, 153
 job_runner() (luigi.contrib.hadoop_jar.HadoopJarJobTask method), 72, 156
 job_runner() (luigi.contrib.hive.HiveQueryTask method), 74, 158
 job_runner() (luigi.contrib.scalding.ScaldingJobTask method), 83, 167
 jobconfs() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
 jobconfs() (luigi.contrib.hadoop.JobTask method), 69, 153
 JobRunner (class in luigi.contrib.hadoop), 68, 153
 JobTask (class in luigi.contrib.hadoop), 68, 153
 jsonpath (luigi.contrib.redshift.S3CopyJSONToTable attribute), 80, 164

K

keep_alive (luigi.worker.worker attribute), 134

KeepAliveThread (class in luigi.worker), 134
 kill_job() (luigi.contrib.hadoop.HadoopRunContext method), 68, 152
 kill_job() (luigi.contrib.pig.PigRunContext method), 76, 160
 kill_job() (luigi.contrib.spark.SparkRunContext method), 83, 167
 KillOpenRedshiftSessions (class in luigi.contrib.redshift), 81, 165

L

list() (luigi.s3.S3Client method), 117
 list_datasets() (luigi.contrib.bigquery.BigqueryClient method), 59, 143
 list_path() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient static method), 55, 139
 list_tables() (luigi.contrib.bigquery.BigqueryClient method), 59, 143
 listdir() (luigi.contrib.gcs.GCSClient method), 66, 150
 listdir() (luigi.contrib.hdfs.abstract_client.HdfsFileSystem method), 52, 136
 listdir() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient method), 55, 138
 listdir() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 57, 141
 listdir() (luigi.contrib.ssh.RemoteFileSystem method), 91, 176
 listdir() (luigi.file.LocalFileSystem method), 102
 listdir() (luigi.mock.MockFileSystem method), 108
 listdir() (luigi.s3.S3Client method), 116
 listdir() (luigi.target.FileSystem method), 124
 load() (luigi.scheduler.CentralPlannerScheduler method), 121
 load() (luigi.scheduler.SimpleTaskState method), 120
 load_hadoop_cmd() (in module luigi.contrib.hdfs.config), 53, 136
 load_hive_cmd() (in module luigi.contrib.hive), 72, 156
 load_task() (in module luigi.task_register), 131
 local_scheduler (luigi.interface.core attribute), 106
 LocalFileSystem (class in luigi.file), 102
 LocalJobRunner (class in luigi.contrib.hadoop), 70, 154
 LocalTarget (class in luigi.file), 102
 lock_pid_dir (luigi.interface.core attribute), 106
 lock_size (luigi.interface.core attribute), 106
 logger (in module luigi.contrib.scalding), 81, 166
 logging_conf_file (luigi.interface.core attribute), 106
 luigi (module), 135
 luigi.cmdline (module), 97
 luigi.configuration (module), 97
 luigi.contrib (module), 93, 177
 luigi.contrib.bigquery (module), 58, 142
 luigi.contrib.esindex (module), 61, 145
 luigi.contrib.ftp (module), 65, 149
 luigi.contrib.gcs (module), 66, 150

luigi.contrib.hadoop (module), 67, 151
luigi.contrib.hadoop_jar (module), 71, 155
luigi.contrib.hdfs (module), 58, 142
luigi.contrib.hdfs.abstract_client (module), 51, 135
luigi.contrib.hdfs.clients (module), 52, 136
luigi.contrib.hdfs.config (module), 52, 136
luigi.contrib.hdfs.error (module), 53, 137
luigi.contrib.hdfs.format (module), 53, 137
luigi.contrib.hdfs.hadoopcli_clients (module), 54, 138
luigi.contrib.hdfs.snakebite_client (module), 55, 139
luigi.contrib.hdfs.target (module), 57, 141
luigi.contrib.hive (module), 72, 156
luigi.contrib.mysqlldb (module), 75, 159
luigi.contrib.pig (module), 75, 159
luigi.contrib.pyspark_runner (module), 76, 160
luigi.contrib.rdbms (module), 76, 161
luigi.contrib.redis_store (module), 77, 162
luigi.contrib.redshift (module), 78, 162
luigi.contrib.scalding (module), 81, 166
luigi.contrib.spark (module), 83, 167
luigi.contrib.sparkey (module), 87, 171
luigi.contrib.sqla (module), 87, 172
luigi.contrib.ssh (module), 91, 175
luigi.contrib.target (module), 92, 176
luigi.contrib.webhdfs (module), 92, 177
luigi.date_interval (module), 98
luigi.db_task_history (module), 100
luigi.deprecate_kwarg (module), 101
luigi.event (module), 102
luigi.file (module), 102
luigi.format (module), 103
luigi.hadoop (module), 105
luigi.hadoop_jar (module), 105
luigi.hdfs (module), 105
luigi.hive (module), 105
luigi.interface (module), 105
luigi.lock (module), 108
luigi.mock (module), 108
luigi.mrrunner (module), 109
luigi.notifications (module), 109
luigi.parameter (module), 110
luigi.postgres (module), 113
luigi.process (module), 115
luigi.rpc (module), 115
luigi.s3 (module), 116
luigi.scalding (module), 119
luigi.scheduler (module), 119
luigi.server (module), 122
luigi.target (module), 123
luigi.task (module), 125
luigi.task_history (module), 130
luigi.task_register (module), 131
luigi.task_status (module), 131
luigi.tools (module), 97, 182

luigi.tools.deps (module), 93, 177
luigi.tools.luigi_grep (module), 93, 178
luigi.tools.parse_task (module), 94, 178
luigi.tools.range (module), 94, 178
luigi.util (module), 132
luigi.webhdfs (module), 133
luigi.worker (module), 133
luigi_run() (in module luigi.cmdline), 97
LuigiConfigParser (class in luigi.configuration), 98
luigid() (in module luigi.cmdline), 97
LuigiGrep (class in luigi.tools.luigi_grep), 93, 178

M

main() (in module luigi.mrrunner), 109
main() (in module luigi.tools.deps), 93, 178
main() (in module luigi.tools.luigi_grep), 93, 178
main() (luigi.contrib.hadoop_jar.HadoopJarJobTask method), 72, 156
main() (luigi.contrib.spark.PySparkTask method), 85, 169
make_dataset() (luigi.contrib.bigquery.BigqueryClient method), 59, 143
makedirs() (luigi.file.LocalTarget method), 102
map_column() (luigi.postgres.CopyToTable method), 115
mapper() (luigi.contrib.hadoop.JobTask method), 69, 153
mapping (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
marker_doc_type (luigi.contrib.esindex.ElasticsearchTarget attribute), 62, 146
marker_index (luigi.contrib.esindex.ElasticsearchTarget attribute), 62, 146
marker_index_document_id() (luigi.contrib.esindex.ElasticsearchTarget method), 62, 146
marker_index_hist_size (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
marker_key() (luigi.contrib.redis_store.RedisTarget method), 78, 162
marker_prefix (luigi.contrib.redis_store.RedisTarget attribute), 78, 162
marker_table (luigi.contrib.mysqlldb.MySqlTarget attribute), 75, 159
marker_table (luigi.contrib.redshift.RedshiftTarget attribute), 78, 162
marker_table (luigi.contrib.sqla.SQLAlchemyTarget attribute), 89, 174
marker_table (luigi.postgres.PostgresTarget attribute), 114
master (luigi.contrib.spark.Spark1xBackwardCompat attribute), 86, 170
master (luigi.contrib.spark.SparkSubmitTask attribute), 83, 168
max_bad_records (luigi.contrib.bigquery.BigqueryLoadTask attribute), 60, 144
max_reschedules (luigi.worker.worker attribute), 134

- max_shown_tasks (luigi.scheduler.scheduler attribute), 119
 - MetastoreClient (class in luigi.contrib.hive), 73, 157
 - missing_datetimes() (luigi.tools.range.RangeBase method), 95, 180
 - missing_datetimes() (luigi.tools.range.RangeDaily method), 97, 181
 - missing_datetimes() (luigi.tools.range.RangeHourly method), 97, 182
 - MissingParameterException, 110
 - MissingParentDirectory, 123
 - MixedUnicodeBytesFormat (class in luigi.format), 104
 - MixedUnicodeBytesWrapper (class in luigi.format), 103
 - MixinNaiveBulkComplete (class in luigi.task), 129
 - mkdir() (luigi.contrib.gcs.GCSClient method), 66, 150
 - mkdir() (luigi.contrib.hdfs.abstract_client.HdfsFileSystem method), 52, 136
 - mkdir() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient method), 55, 138
 - mkdir() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClientCdh method), 55, 139
 - mkdir() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 57, 140
 - mkdir() (luigi.contrib.ssh.RemoteFileSystem method), 91, 176
 - mkdir() (luigi.file.LocalFileSystem method), 102
 - mkdir() (luigi.mock.MockFileSystem method), 109
 - mkdir() (luigi.s3.S3Client method), 117
 - mkdir() (luigi.target.FileSystem method), 124
 - MockFile (class in luigi.mock), 109
 - MockFileSystem (class in luigi.mock), 108
 - MockTarget (class in luigi.mock), 109
 - module (luigi.interface.core attribute), 106
 - Month (class in luigi.date_interval), 99
 - most_common() (in module luigi.tools.range), 96, 181
 - move() (luigi.contrib.hdfs.target.HdfsTarget method), 58, 141
 - move() (luigi.file.LocalTarget method), 102
 - move_dir() (luigi.contrib.hdfs.target.HdfsTarget method), 58, 142
 - move_dir() (luigi.file.LocalTarget method), 102
 - move_to_final_destination() (luigi.contrib.ftp.AtomicFtpFile method), 65, 149
 - move_to_final_destination() (luigi.contrib.gcs.AtomicGCSFile method), 66, 151
 - move_to_final_destination() (luigi.contrib.webhdfs.AtomicWebHdfsFile method), 92, 177
 - move_to_final_destination() (luigi.file.atomic_file method), 102
 - move_to_final_destination() (luigi.s3.AtomicS3File method), 117
 - move_to_final_destination() (luigi.target.AtomicLocalFile method), 125
 - moving_start() (luigi.tools.range.RangeBase method), 95, 179
 - moving_start() (luigi.tools.range.RangeDailyBase method), 95, 180
 - moving_start() (luigi.tools.range.RangeHourlyBase method), 96, 181
 - moving_stop() (luigi.tools.range.RangeBase method), 95, 179
 - moving_stop() (luigi.tools.range.RangeDailyBase method), 95, 180
 - moving_stop() (luigi.tools.range.RangeHourlyBase method), 96, 181
 - mr_priority (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 68, 152
 - MultiReplacer (class in luigi.postgres), 113
 - MySQLTarget (class in luigi.contrib.mysql), 75, 159
- ## N
- n_reduce_tasks (luigi.contrib.hadoop.JobTask attribute), 69, 154
 - n_reduce_tasks (luigi.contrib.hive.HiveQueryTask attribute), 73, 157
 - name (luigi.contrib.spark.PySparkTask attribute), 84, 169
 - name (luigi.contrib.spark.SparkSubmitTask attribute), 83, 168
 - name (luigi.db_task_history.TaskParameter attribute), 101
 - name (luigi.db_task_history.TaskRecord attribute), 101
 - namenode_host (luigi.contrib.hdfs.config.hdfs attribute), 52, 136
 - namenode_port (luigi.contrib.hdfs.config.hdfs attribute), 52, 136
 - namespace() (in module luigi.task), 125
 - NEWLINE_DELIMITED_JSON (luigi.contrib.bigquery.SourceFormat attribute), 59, 143
 - NewlineFormat (class in luigi.format), 104
 - NewlineWrapper (class in luigi.format), 103
 - next() (luigi.date_interval.DateInterval method), 99
 - NO_DEFAULT (luigi.configuration.LuigiConfigParser attribute), 98
 - no_lock (luigi.interface.core attribute), 106
 - NopFormat (class in luigi.format), 104
 - NopHistory (class in luigi.task_history), 130
 - NotADirectory, 123
 - now (luigi.tools.range.RangeBase attribute), 95, 179
 - null_values (luigi.contrib.rdbms.CopyToTable attribute), 77, 161
 - num_executors (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
 - num_failures() (luigi.scheduler.Failures method), 120

O

of (luigi.tools.range.RangeBase attribute), 94, 179
on_failure() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
on_failure() (luigi.task.Task method), 128
on_success() (luigi.task.Task method), 128
open() (luigi.contrib.ftp.RemoteTarget method), 65, 149
open() (luigi.contrib.gcs.GCSTarget method), 67, 151
open() (luigi.contrib.hdfs.target.HdfsTarget method), 57, 141
open() (luigi.contrib.hive.HivePartitionTarget method), 74, 158
open() (luigi.contrib.hive.HiveTableTarget method), 74, 158
open() (luigi.contrib.sqla.SQLAlchemyTarget method), 90, 174
open() (luigi.contrib.ssh.RemoteTarget method), 92, 176
open() (luigi.contrib.webhdfs.WebHdfsTarget method), 92, 177
open() (luigi.file.LocalTarget method), 102
open() (luigi.mock.MockTarget method), 109
open() (luigi.postgres.PostgresTarget method), 114
open() (luigi.s3.S3Target method), 117
open() (luigi.target.FileSystemTarget method), 125
OptParseInterface (class in luigi.interface), 107
output (luigi.contrib.hdfs.format.CompatibleHdfsFormat attribute), 54, 138
output (luigi.contrib.hdfs.format.PlainDirFormat attribute), 54, 138
output (luigi.contrib.hdfs.format.PlainFormat attribute), 54, 137
output (luigi.format.Bzip2Format attribute), 105
output (luigi.format.GzipFormat attribute), 105
output (luigi.format.MixedUnicodeBytesFormat attribute), 104
output (luigi.format.NewlineFormat attribute), 104
output (luigi.format.TextFormat attribute), 104
output() (luigi.contrib.esindex.CopyToIndex method), 64, 148
output() (luigi.contrib.hive.ExternalHiveTask method), 74, 159
output() (luigi.contrib.pig.PigJobTask method), 76, 160
output() (luigi.contrib.rdbms.CopyToTable method), 77, 161
output() (luigi.contrib.redshift.KillOpenRedshiftSessions method), 81, 165
output() (luigi.contrib.redshift.S3CopyToTable method), 79, 163
output() (luigi.contrib.spark.Spark1xBackwardCompat method), 86, 170
output() (luigi.contrib.spark.SparkJob method), 85, 170
output() (luigi.contrib.sqla.CopyToTable method), 90, 175
output() (luigi.postgres.CopyToTable method), 115

output() (luigi.s3.S3EmrTask method), 118
output() (luigi.s3.S3FlagTask method), 118
output() (luigi.s3.S3PathTask method), 118
output() (luigi.task.Task method), 128
OutputPipeProcessWrapper (class in luigi.format), 103

P

parallel_scheduling (luigi.interface.core attribute), 106
Parameter (class in luigi.parameter), 110
parameter_to_datetime() (luigi.tools.range.RangeBase method), 95, 179
parameter_to_datetime() (luigi.tools.range.RangeDailyBase method), 95, 180
parameter_to_datetime() (luigi.tools.range.RangeHourlyBase method), 96, 181
ParameterException, 110
parameters (luigi.db_task_history.TaskRecord attribute), 101
parse() (luigi.date_interval.Custom class method), 100
parse() (luigi.date_interval.Date class method), 99
parse() (luigi.date_interval.DateInterval class method), 99
parse() (luigi.date_interval.Month class method), 99
parse() (luigi.date_interval.Week class method), 99
parse() (luigi.date_interval.Year class method), 100
parse() (luigi.interface.ArgParseInterface method), 107
parse() (luigi.interface.DynamicArgParseInterface method), 107
parse() (luigi.interface.Interface method), 106
parse() (luigi.interface.OptParseInterface method), 107
parse() (luigi.parameter.BoolParameter method), 113
parse() (luigi.parameter.DateHourParameter method), 112
parse() (luigi.parameter.DateIntervalParameter method), 113
parse() (luigi.parameter.DateParameter method), 112
parse() (luigi.parameter.FloatParameter method), 112
parse() (luigi.parameter.IntParameter method), 112
parse() (luigi.parameter.Parameter method), 111
parse() (luigi.parameter.TimeDeltaParameter method), 113
parse_from_args() (luigi.parameter.Parameter method), 112
parse_from_input() (luigi.parameter.Parameter method), 111
parse_task() (luigi.interface.ArgParseInterface method), 107
parser_dest() (luigi.parameter.Parameter method), 112
partition (luigi.contrib.hive.ExternalHiveTask attribute), 74, 159
partition_spec() (luigi.contrib.hive.HiveClient method), 72, 156
partition_spec() (luigi.contrib.hive.HiveCommandClient method), 73, 157

partition_spec() (luigi.contrib.hive.MetastoreClient method), 73, 157
 PassThroughOptionParser (class in luigi.interface), 107
 password (luigi.contrib.rdbms.CopyToTable attribute), 77, 161
 password (luigi.contrib.redshift.KillOpenRedshiftSessions attribute), 81, 165
 path (luigi.contrib.hive.HivePartitionTarget attribute), 74, 158
 path (luigi.contrib.hive.HiveTableTarget attribute), 74, 158
 path (luigi.mock.MockTarget attribute), 109
 path (luigi.s3.S3EmrTask attribute), 118
 path (luigi.s3.S3FlagTask attribute), 118
 path (luigi.s3.S3PathTask attribute), 118
 pid (luigi.contrib.sqla.SQLAlchemyTarget.Connection attribute), 89, 174
 pig_command_path() (luigi.contrib.pig.PigJobTask method), 75, 160
 pig_env_vars() (luigi.contrib.pig.PigJobTask method), 75, 160
 pig_home() (luigi.contrib.pig.PigJobTask method), 75, 160
 pig_options() (luigi.contrib.pig.PigJobTask method), 76, 160
 pig_parameters() (luigi.contrib.pig.PigJobTask method), 76, 160
 pig_properties() (luigi.contrib.pig.PigJobTask method), 75, 160
 pig_script_path() (luigi.contrib.pig.PigJobTask method), 76, 160
 PigJobError, 76, 160
 PigJobTask (class in luigi.contrib.pig), 75, 159
 PigRunContext (class in luigi.contrib.pig), 76, 160
 ping (luigi.scheduler.Scheduler attribute), 119
 ping() (luigi.rpc.RemoteScheduler method), 115
 ping() (luigi.scheduler.CentralPlannerScheduler method), 121
 ping_interval (luigi.worker.worker attribute), 134
 pipe_reader() (luigi.contrib.hdfs.format.CompatibleHdfsFormat method), 54, 138
 pipe_reader() (luigi.contrib.hdfs.format.PlainDirFormat method), 54, 138
 pipe_reader() (luigi.contrib.hdfs.format.PlainFormat method), 54, 137
 pipe_reader() (luigi.format.Bzip2Format method), 105
 pipe_reader() (luigi.format.ChainFormat method), 104
 pipe_reader() (luigi.format.Format class method), 104
 pipe_reader() (luigi.format.GzipFormat method), 105
 pipe_reader() (luigi.format.NopFormat method), 104
 pipe_reader() (luigi.format.WrappedFormat method), 104
 pipe_writer() (luigi.contrib.hdfs.format.CompatibleHdfsFormat method), 54, 138
 pipe_writer() (luigi.contrib.hdfs.format.PlainDirFormat method), 54, 138
 pipe_writer() (luigi.contrib.hdfs.format.PlainFormat method), 54, 137
 pipe_writer() (luigi.format.Bzip2Format method), 105
 pipe_writer() (luigi.format.ChainFormat method), 104
 pipe_writer() (luigi.format.Format class method), 104
 pipe_writer() (luigi.format.GzipFormat method), 105
 pipe_writer() (luigi.format.NopFormat method), 104
 pipe_writer() (luigi.format.WrappedFormat method), 104
 PlainDirFormat (class in luigi.contrib.hdfs.format), 54, 138
 PlainFormat (class in luigi.contrib.hdfs.format), 53, 137
 pool (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 68, 152
 pool (luigi.contrib.hadoop.hadoop attribute), 71, 155
 Popen() (luigi.contrib.ssh.RemoteContext method), 91, 175
 port (luigi.contrib.esindex.CopyToIndex attribute), 63, 147
 post() (luigi.server.RPCHandler method), 122
 PostgresTarget (class in luigi.postgres), 114
 prefix_search() (luigi.tools.luigi_grep.LuigiGrep method), 93, 178
 prepare_outputs() (luigi.contrib.hive.HiveQueryRunner method), 74, 158
 prev() (luigi.date_interval.DateInterval method), 99
 previous() (in module luigi.util), 132
 print_exception() (in module luigi.mrrunner), 109
 priority (luigi.task.Task attribute), 126
 process_resources() (luigi.task.Task method), 128
 PROCESSING_TIME (luigi.event.Event attribute), 102
 program() (luigi.contrib.spark.PySpark1xJob method), 87, 171
 project_id (luigi.contrib.bigquery.BQDataset attribute), 59, 143
 properties_file (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
 prune() (luigi.rpc.RemoteScheduler method), 116
 prune() (luigi.scheduler.CentralPlannerScheduler method), 121
 prune() (luigi.scheduler.SimpleTaskState method), 121
 prune() (luigi.scheduler.Worker method), 120
 prune_done_tasks (luigi.scheduler.scheduler attribute), 119
 purge_existing_index (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
 put() (luigi.contrib.ftp.RemoteFileSystem method), 65, 149
 put() (luigi.contrib.ftp.RemoteTarget method), 66, 150
 put() (luigi.contrib.gcs.GCSClient method), 66, 150
 put() (luigi.contrib.hdfs.abstract_client.HdfsFileSystem method), 52, 135
 put() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient method), 55, 138

- put() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 56, 140
- put() (luigi.contrib.ssh.RemoteFileSystem method), 91, 176
- put() (luigi.contrib.ssh.RemoteTarget method), 92, 176
- put() (luigi.s3.S3Client method), 116
- put() (luigi.worker.DequeueQueue method), 133
- put_multipart() (luigi.s3.S3Client method), 116
- put_string() (luigi.contrib.gcs.GCSClient method), 66, 150
- put_string() (luigi.s3.S3Client method), 116
- py_files (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- py_packages (luigi.contrib.spark.PySparkTask attribute), 84, 169
- PySpark1xJob (class in luigi.contrib.spark), 86, 171
- PySparkRunner (class in luigi.contrib.pyspark_runner), 76, 160
- PySparkTask (class in luigi.contrib.spark), 84, 169
- ## Q
- query (luigi.contrib.bigquery.BigqueryRunQueryTask attribute), 61, 145
- query() (luigi.contrib.hive.HiveQueryTask method), 73, 157
- query_mode (luigi.contrib.bigquery.BigqueryRunQueryTask attribute), 61, 145
- QueryMode (class in luigi.contrib.bigquery), 58, 142
- queue (luigi.contrib.spark.SparkJob attribute), 85, 170
- queue (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168
- ## R
- raise_on_error (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
- RangeBase (class in luigi.tools.range), 94, 179
- RangeDaily (class in luigi.tools.range), 96, 181
- RangeDailyBase (class in luigi.tools.range), 95, 180
- RangeEvent (class in luigi.tools.range), 94, 178
- RangeHourly (class in luigi.tools.range), 97, 181
- RangeHourlyBase (class in luigi.tools.range), 96, 180
- re_enable() (luigi.scheduler.SimpleTaskState method), 121
- re_enable_task() (luigi.rpc.RemoteScheduler method), 116
- re_enable_task() (luigi.scheduler.CentralPlannerScheduler method), 122
- read() (luigi.contrib.webhdfs.ReadableWebHdfsFile method), 92, 177
- read() (luigi.contrib.webhdfs.WebHdfsClient method), 93, 177
- read() (luigi.format.NewlineWrapper method), 103
- read() (luigi.s3.ReadableS3File method), 117
- readable() (luigi.format.InputPipeProcessWrapper method), 103
- readable() (luigi.format.OutputPipeProcessWrapper method), 103
- readable() (luigi.s3.ReadableS3File method), 117
- ReadableS3File (class in luigi.s3), 117
- ReadableWebHdfsFile (class in luigi.contrib.webhdfs), 92, 177
- reader() (luigi.contrib.hadoop.JobTask method), 70, 154
- readlines() (luigi.contrib.webhdfs.ReadableWebHdfsFile method), 92, 177
- RecentRunHandler (class in luigi.server), 122
- record_task_history (luigi.scheduler.scheduler attribute), 119
- recursive_listdir_cmd (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient attribute), 54, 138
- recursive_listdir_cmd (luigi.contrib.hdfs.hadoopcli_clients.HdfsClientApache attribute), 55, 139
- RedisTarget (class in luigi.contrib.redis_store), 77, 162
- RedshiftManifestTask (class in luigi.contrib.redshift), 80, 164
- RedshiftTarget (class in luigi.contrib.redshift), 78, 162
- reducer (luigi.contrib.hadoop.JobTask attribute), 70, 154
- reducers_max (luigi.contrib.hive.HiveQueryTask attribute), 73, 157
- reflect (luigi.contrib.sqla.CopyToTable attribute), 90, 175
- Register (class in luigi.task_register), 131
- reload() (luigi.configuration.LuigiConfigParser class method), 98
- relpath() (luigi.contrib.scalding.ScaldingJobTask method), 82, 167
- RemoteContext (class in luigi.contrib.ssh), 91, 175
- RemoteFileSystem (class in luigi.contrib.ftp), 65, 149
- RemoteFileSystem (class in luigi.contrib.ssh), 91, 176
- RemoteScheduler (class in luigi.rpc), 115
- RemoteTarget (class in luigi.contrib.ftp), 65, 149
- RemoteTarget (class in luigi.contrib.ssh), 92, 176
- remove() (luigi.contrib.ftp.RemoteFileSystem method), 65, 149
- remove() (luigi.contrib.gcs.GCSClient method), 66, 150
- remove() (luigi.contrib.hdfs.abstract_client.HdfsFileSystem method), 51, 135
- remove() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient method), 54, 138
- remove() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClientCdh3 method), 55, 139
- remove() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 56, 139
- remove() (luigi.contrib.hdfs.target.HdfsTarget method), 57, 141
- remove() (luigi.contrib.ssh.RemoteFileSystem method), 91, 176
- remove() (luigi.contrib.webhdfs.WebHdfsClient method), 93, 177

remove() (luigi.file.LocalFileSystem method), 102
 remove() (luigi.file.LocalTarget method), 103
 remove() (luigi.mock.MockFileSystem method), 108
 remove() (luigi.s3.S3Client method), 116
 remove() (luigi.target.FileSystem method), 124
 remove() (luigi.target.FileSystemTarget method), 125
 remove_delay (luigi.scheduler.scheduler attribute), 119
 rename() (luigi.contrib.gcs.GCSClient method), 66, 150
 rename() (luigi.contrib.hdfs.abstract_client.HdfsFileSystem method), 51, 135
 rename() (luigi.contrib.hdfs.hadoopcli_clients.HdfsClient method), 54, 138
 rename() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 55, 139
 rename() (luigi.contrib.hdfs.target.HdfsTarget method), 57, 141
 rename() (luigi.mock.MockTarget method), 109
 rename() (luigi.s3.S3Client method), 116
 rename_dont_move() (luigi.contrib.hdfs.abstract_client.HdfsFileSystem method), 51, 135
 rename_dont_move() (luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient method), 55, 139
 requires (class in luigi.util), 132
 requires() (luigi.contrib.scalding.ScaldingJobTask method), 83, 167
 requires() (luigi.task.Task method), 128
 requires() (luigi.tools.range.RangeBase method), 95, 180
 requires_hadoop() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
 requires_hadoop() (luigi.contrib.spark.SparkJob method), 85, 170
 requires_local() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
 requires_local() (luigi.contrib.spark.SparkJob method), 85, 170
 reset_global() (luigi.parameter.Parameter method), 111
 resources (luigi.task.Task attribute), 126
 retry_delay (luigi.scheduler.scheduler attribute), 119
 retry_external_tasks (luigi.worker.worker attribute), 134
 reverse (luigi.tools.range.RangeBase attribute), 95, 179
 RootPathHandler (class in luigi.server), 123
 rows() (luigi.contrib.sqla.CopyToTable method), 90, 175
 rows() (luigi.postgres.CopyToTable method), 114
 RPCError, 115
 RPCHandler (class in luigi.server), 122
 run (luigi.task.ExternalTask attribute), 129
 run() (in module luigi.interface), 107
 run() (in module luigi.server), 123
 run() (luigi.contrib.bigquery.BigqueryLoadTask method), 61, 145
 run() (luigi.contrib.bigquery.BigqueryRunQueryTask method), 61, 145
 run() (luigi.contrib.esindex.CopyToIndex method), 64, 148
 run() (luigi.contrib.hadoop.BaseHadoopJobTask method), 68, 152
 run() (luigi.contrib.pig.PigJobTask method), 76, 160
 run() (luigi.contrib.pyspark_runner.PySparkRunner method), 76, 160
 run() (luigi.contrib.redshift.KillOpenRedshiftSessions method), 81, 165
 run() (luigi.contrib.redshift.RedshiftManifestTask method), 80, 165
 run() (luigi.contrib.redshift.S3CopyToTable method), 79, 163
 run() (luigi.contrib.spark.PySpark1xJob method), 87, 171
 run() (luigi.contrib.spark.PySparkTask method), 85, 169
 run() (luigi.contrib.spark.Spark1xJob method), 86, 171
 run() (luigi.contrib.spark.SparkJob method), 85, 170
 run() (luigi.contrib.spark.SparkSubmitTask method), 84, 169
 run() (luigi.contrib.sparkey.SparkeyExportTask method), 87, 171
 run() (luigi.contrib.sqla.CopyToTable method), 90, 175
 run() (luigi.contrib.sqlalchemy.Interface static method), 106
 run() (luigi.mrrunner.Runner method), 109
 run() (luigi.postgres.CopyToTable method), 115
 run() (luigi.task.Task method), 128
 run() (luigi.worker.KeepAliveThread method), 134
 run() (luigi.worker.TaskProcess method), 133
 run() (luigi.worker.Worker method), 135
 run_and_track_hadoop_job() (in module luigi.contrib.hadoop), 71, 155
 run_combiner() (luigi.contrib.hadoop.JobTask method), 70, 154
 run_hive() (in module luigi.contrib.hive), 72, 156
 run_hive_cmd() (in module luigi.contrib.hive), 72, 156
 run_hive_script() (in module luigi.contrib.hive), 72, 156
 run_job (luigi.contrib.hadoop.JobRunner attribute), 68, 153
 run_job() (luigi.contrib.bigquery.BigqueryClient method), 60, 144
 run_job() (luigi.contrib.hadoop.HadoopJobRunner method), 68, 152
 run_job() (luigi.contrib.hadoop.LocalJobRunner method), 70, 154
 run_job() (luigi.contrib.hadoop_jar.HadoopJarJobRunner method), 71, 155
 run_job() (luigi.contrib.hive.HiveQueryRunner method), 74, 158
 run_job() (luigi.contrib.scalding.ScaldingJobRunner method), 82, 166
 run_mapper() (luigi.contrib.hadoop.JobTask method), 70, 154
 run_reducer() (luigi.contrib.hadoop.JobTask method), 70, 154
 Runner (class in luigi.mrrunner), 109

S

- s3_load_path (luigi.contrib.redshift.S3CopyToTable attribute), 78, 163
- S3Client (class in luigi.s3), 116
- S3CopyJSONToTable (class in luigi.contrib.redshift), 79, 164
- S3CopyToTable (class in luigi.contrib.redshift), 78, 162
- S3EmrTarget (class in luigi.s3), 118
- S3EmrTask (class in luigi.s3), 118
- S3FlagTarget (class in luigi.s3), 117
- S3FlagTask (class in luigi.s3), 118
- S3PathTask (class in luigi.s3), 118
- S3Target (class in luigi.s3), 117
- sample() (luigi.contrib.hadoop.LocalJobRunner method), 70, 154
- ScaldingJobRunner (class in luigi.contrib.scalding), 82, 166
- ScaldingJobTask (class in luigi.contrib.scalding), 82, 166
- Scheduler (class in luigi.scheduler), 119
- scheduler (class in luigi.scheduler), 119
- scheduler_host (luigi.interface.core attribute), 106
- scheduler_port (luigi.interface.core attribute), 106
- schema (luigi.contrib.bigquery.BigqueryLoadTask attribute), 60, 144
- seekable() (luigi.format.InputPipeProcessWrapper method), 103
- seekable() (luigi.format.OutputPipeProcessWrapper method), 103
- seekable() (luigi.s3.ReadableS3File method), 117
- SelectedRunHandler (class in luigi.server), 122
- send_email() (in module luigi.notifications), 110
- send_email_sendgrid() (in module luigi.notifications), 110
- send_email_ses() (in module luigi.notifications), 110
- send_email_smtp() (in module luigi.notifications), 110
- send_error_email() (in module luigi.notifications), 110
- separator (luigi.contrib.sparkey.SparkeyExportTask attribute), 87, 171
- serialize (luigi.contrib.hadoop.JobTask attribute), 70, 154
- serialize() (luigi.parameter.DateHourParameter method), 112
- serialize() (luigi.parameter.Parameter method), 111
- serialize_to_input() (luigi.parameter.Parameter method), 112
- set() (luigi.configuration.LuigiConfigParser method), 98
- set_global() (luigi.parameter.Parameter method), 111
- set_global_from_args() (luigi.parameter.Parameter method), 112
- set_global_parameters() (in module luigi.interface), 106
- set_status() (luigi.scheduler.SimpleTaskState method), 121
- settings (luigi.contrib.esindex.CopyToIndex attribute), 64, 148
- setup() (luigi.contrib.spark.PySparkTask method), 85, 169
- setup_interface_logging() (in module luigi.interface), 105
- setup_remote() (luigi.contrib.spark.PySparkTask method), 85, 169
- SimpleTaskState (class in luigi.scheduler), 120
- SingleProcessPool (class in luigi.worker), 133
- snakebite_autoconfig (luigi.contrib.hdfs.config.hdfs attribute), 52, 136
- SnakebiteHdfsClient (class in luigi.contrib.hdfs.snakebite_client), 55, 139
- source() (luigi.contrib.scalding.ScaldingJobTask method), 82, 167
- source_format (luigi.contrib.bigquery.BigqueryLoadTask attribute), 60, 144
- source_uris (luigi.contrib.bigquery.BigqueryLoadTask attribute), 61, 145
- SourceFormat (class in luigi.contrib.bigquery), 58, 142
- Spark1xBackwardCompat (class in luigi.contrib.spark), 86, 170
- Spark1xJob (class in luigi.contrib.spark), 86, 170
- spark_command() (luigi.contrib.spark.Spark1xBackwardCompat method), 86, 170
- spark_command() (luigi.contrib.spark.SparkSubmitTask method), 84, 168
- spark_master_memory (luigi.contrib.spark.SparkJob attribute), 85, 170
- spark_options() (luigi.contrib.spark.Spark1xBackwardCompat method), 86, 170
- spark_submit (luigi.contrib.spark.SparkSubmitTask attribute), 83, 168
- spark_worker_memory (luigi.contrib.spark.SparkJob attribute), 85, 170
- spark_workers (luigi.contrib.spark.SparkJob attribute), 85, 170
- SparkeyExportTask (class in luigi.contrib.sparkey), 87, 171
- SparkJob (class in luigi.contrib.spark), 85, 169
- SparkJobError, 83, 167
- SparkRunContext (class in luigi.contrib.spark), 83, 167
- SparkSubmitTask (class in luigi.contrib.spark), 83, 167
- SQLAlchemyTarget (class in luigi.contrib.sqla), 89, 173
- SQLAlchemyTarget.Connection (class in luigi.contrib.sqla), 89, 174
- ssh() (luigi.contrib.hadoop_jar.HadoopJarJobTask method), 72, 156
- START (luigi.event.Event attribute), 102
- start (luigi.tools.range.RangeBase attribute), 94, 179
- start (luigi.tools.range.RangeDailyBase attribute), 95, 180
- start (luigi.tools.range.RangeHourlyBase attribute), 96, 181
- state_path (luigi.scheduler.scheduler attribute), 119
- StaticFileHandler (class in luigi.server), 123
- status_search() (luigi.tools.luigi_grep.LuigiGrep

method), 93, 178
 stop (luigi.tools.range.RangeBase attribute), 94, 179
 stop (luigi.tools.range.RangeDailyBase attribute), 95, 180
 stop (luigi.tools.range.RangeHourlyBase attribute), 96, 181
 stop() (in module luigi.server), 123
 stop() (luigi.worker.KeepAliveThread method), 134
 stop() (luigi.worker.Worker method), 134
 SUCCESS (luigi.event.Event attribute), 102
 supervise (luigi.contrib.spark.SparkSubmitTask attribute), 84, 168

T

table (luigi.contrib.hive.ExternalHiveTask attribute), 74, 158
 table (luigi.contrib.rdbms.CopyToTable attribute), 77, 161
 table (luigi.contrib.sqla.CopyToTable attribute), 90, 174
 table_attributes() (luigi.contrib.redshift.S3CopyToTable method), 79, 163
 table_exists() (luigi.contrib.bigquery.BigqueryClient method), 59, 143
 table_exists() (luigi.contrib.hive.HiveClient method), 72, 156
 table_exists() (luigi.contrib.hive.HiveCommandClient method), 72, 157
 table_exists() (luigi.contrib.hive.MetastoreClient method), 73, 157
 table_location() (luigi.contrib.hive.HiveClient method), 72, 156
 table_location() (luigi.contrib.hive.HiveCommandClient method), 72, 157
 table_location() (luigi.contrib.hive.MetastoreClient method), 73, 157
 table_schema() (luigi.contrib.hive.ApacheHiveCommandClient method), 73, 157
 table_schema() (luigi.contrib.hive.HiveClient method), 72, 156
 table_schema() (luigi.contrib.hive.HiveCommandClient method), 72, 157
 table_schema() (luigi.contrib.hive.MetastoreClient method), 73, 157
 Target (class in luigi.target), 123
 Task (class in luigi.scheduler), 120
 Task (class in luigi.task), 126
 Task (class in luigi.task_history), 130
 task_family (luigi.task.Task attribute), 127
 task_family (luigi.task_register.Register attribute), 131
 task_finished() (luigi.db_task_history.DbTaskHistory method), 100
 task_finished() (luigi.task_history.NopHistory method), 130
 task_finished() (luigi.task_history.TaskHistory method), 130
 task_history (luigi.scheduler.CentralPlannerScheduler attribute), 122
 task_id (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 68, 152
 task_id (luigi.db_task_history.TaskEvent attribute), 101
 task_id (luigi.db_task_history.TaskParameter attribute), 101
 task_limit (luigi.tools.range.RangeBase attribute), 95, 179
 task_limit (luigi.worker.worker attribute), 134
 task_list() (luigi.rpc.RemoteScheduler method), 116
 task_list() (luigi.scheduler.CentralPlannerScheduler method), 122
 task_module (luigi.task.Task attribute), 127
 task_names() (luigi.task_register.Register class method), 131
 task_namespace (luigi.contrib.bigquery.BigqueryLoadTask attribute), 61, 145
 task_namespace (luigi.contrib.bigquery.BigqueryRunQueryTask attribute), 61, 145
 task_namespace (luigi.contrib.esindex.CopyToIndex attribute), 65, 149
 task_namespace (luigi.contrib.hadoop.BaseHadoopJobTask attribute), 68, 152
 task_namespace (luigi.contrib.hadoop.hadoop attribute), 71, 155
 task_namespace (luigi.contrib.hadoop.JobTask attribute), 70, 154
 task_namespace (luigi.contrib.hadoop_jar.HadoopJarJobTask attribute), 72, 156
 task_namespace (luigi.contrib.hdfs.config.hadoopcli attribute), 53, 136
 task_namespace (luigi.contrib.hdfs.config.hdfs attribute), 52, 136
 task_namespace (luigi.contrib.hive.ExternalHiveTask attribute), 74, 159
 task_namespace (luigi.contrib.hive.HiveQueryTask attribute), 74, 158
 task_namespace (luigi.contrib.pig.PigJobTask attribute), 76, 160
 task_namespace (luigi.contrib.rdbms.CopyToTable attribute), 77, 162
 task_namespace (luigi.contrib.redshift.KillOpenRedshiftSessions attribute), 81, 166
 task_namespace (luigi.contrib.redshift.RedshiftManifestTask attribute), 81, 165
 task_namespace (luigi.contrib.redshift.S3CopyJSONToTable attribute), 80, 164
 task_namespace (luigi.contrib.redshift.S3CopyToTable attribute), 79, 164
 task_namespace (luigi.contrib.scalding.ScaldingJobTask attribute), 83, 167
 task_namespace (luigi.contrib.spark.PySparkIxJob attribute), 87, 171

`task_namespace` (`luigi.contrib.spark.PySparkTask` attribute), 85, 169

`task_namespace` (`luigi.contrib.spark.Spark1xBackwardCom` attribute), 86, 170

`task_namespace` (`luigi.contrib.spark.Spark1xJob` attribute), 86, 171

`task_namespace` (`luigi.contrib.spark.SparkJob` attribute), 86, 170

`task_namespace` (`luigi.contrib.spark.SparkSubmitTask` attribute), 84, 169

`task_namespace` (`luigi.contrib.sparkey.SparkeyExportTask` attribute), 87, 171

`task_namespace` (`luigi.contrib.sqla.CopyToTable` attribute), 91, 175

`task_namespace` (`luigi.interface.core` attribute), 106

`task_namespace` (`luigi.postgres.CopyToTable` attribute), 115

`task_namespace` (`luigi.s3.S3EmrTask` attribute), 118

`task_namespace` (`luigi.s3.S3FlagTask` attribute), 118

`task_namespace` (`luigi.s3.S3PathTask` attribute), 118

`task_namespace` (`luigi.scheduler.scheduler` attribute), 119

`task_namespace` (`luigi.task.Config` attribute), 130

`task_namespace` (`luigi.task.ExternalTask` attribute), 129

`task_namespace` (`luigi.task.Task` attribute), 129

`task_namespace` (`luigi.task WrapperTask` attribute), 129

`task_namespace` (`luigi.tools.deps.upstream` attribute), 93, 178

`task_namespace` (`luigi.tools.range.RangeBase` attribute), 95, 180

`task_namespace` (`luigi.tools.range.RangeDaily` attribute), 97, 181

`task_namespace` (`luigi.tools.range.RangeDailyBase` attribute), 96, 180

`task_namespace` (`luigi.tools.range.RangeHourly` attribute), 97, 182

`task_namespace` (`luigi.tools.range.RangeHourlyBase` attribute), 96, 181

`task_namespace` (`luigi.worker.worker` attribute), 134

`task_scheduled()` (`luigi.db_task_history.DbTaskHistory` method), 100

`task_scheduled()` (`luigi.task_history.NopHistory` method), 130

`task_scheduled()` (`luigi.task_history.TaskHistory` method), 130

`task_search()` (`luigi.rpc.RemoteScheduler` method), 116

`task_search()` (`luigi.scheduler.CentralPlannerScheduler` method), 122

`task_started()` (`luigi.db_task_history.DbTaskHistory` method), 100

`task_started()` (`luigi.task_history.NopHistory` method), 130

`task_started()` (`luigi.task_history.TaskHistory` method), 130

`task_value()` (`luigi.parameter.Parameter` method), 111

`task_wraps()` (in module `luigi.util`), 132

`TaskClassException`, 131

`TaskEvent` (class in `luigi.db_task_history`), 101

`TaskException`, 133

`TaskHistory` (class in `luigi.task_history`), 130

`TaskParameter` (class in `luigi.db_task_history`), 100

`TaskProcess` (class in `luigi.worker`), 133

`TaskRecord` (class in `luigi.db_task_history`), 101

`tasks_str()` (`luigi.task_register.Register` class method), 131

`temp_hadoop_output_file` (`luigi.contrib.spark.SparkJob` attribute), 85, 170

`text_target` (`luigi.contrib.redshift.RedshiftManifestTask` attribute), 80, 165

`TextFormat` (class in `luigi.format`), 104

`TextWrapper` (class in `luigi.format`), 104

`TimeDeltaParameter` (class in `luigi.parameter`), 113

`timeout` (`luigi.contrib.esindex.CopyToIndex` attribute), 64, 148

`timeout` (`luigi.worker.worker` attribute), 134

`tmp_dir` (`luigi.contrib.hdfs.config.hdfs` attribute), 52, 136

`tmp_path` (`luigi.contrib.ssh.AtomicRemoteFileWriter` attribute), 92, 176

`tmp_path` (`luigi.target.AtomicLocalFile` attribute), 125

`tmppath()` (in module `luigi.contrib.hdfs.config`), 53, 137

`to_str_params()` (`luigi.task.Task` method), 127

`to_string()` (`luigi.date_interval.Custom` method), 100

`to_string()` (`luigi.date_interval.Date` method), 99

`to_string()` (`luigi.date_interval.DateInterval` method), 99

`to_string()` (`luigi.date_interval.Month` method), 99

`to_string()` (`luigi.date_interval.Week` method), 99

`to_string()` (`luigi.date_interval.Year` method), 100

`total_executor_cores` (`luigi.contrib.spark.SparkSubmitTask` attribute), 84, 168

`touch()` (`luigi.contrib.esindex.ElasticsearchTarget` method), 62, 146

`touch()` (`luigi.contrib.mysql.MySQLTarget` method), 75, 159

`touch()` (`luigi.contrib.redis_store.RedisTarget` method), 78, 162

`touch()` (`luigi.contrib.sqla.SQLAlchemyTarget` method), 89, 174

`touch()` (`luigi.postgres.PostgresTarget` method), 114

`touchz()` (`luigi.contrib.hdfs.abstract_client.HdfsFileSystem` method), 52, 136

`touchz()` (`luigi.contrib.hdfs.hadoopcli_clients.HdfsClient` method), 55, 138

`touchz()` (`luigi.contrib.hdfs.snakebite_client.SnakebiteHdfsClient` method), 57, 141

`TracebackWrapper` (class in `luigi.worker`), 133

`track_and_progress()` (`luigi.contrib.pig.PigJobTask` method), 76, 160

`track_progress()` (`luigi.contrib.spark.SparkJob` method), 85, 170

trigger_event() (luigi.task.Task method), 127
 truncate_table() (luigi.contrib.redshift.S3CopyToTable method), 79, 163
 ts (luigi.db_task_history.TaskEvent attribute), 101
 tunnel() (luigi.contrib.ssh.RemoteContext method), 91, 176

U

UnknownConfigException, 110
 UnknownParameterException, 110
 update() (luigi.scheduler.CentralPlannerScheduler method), 121
 update() (luigi.scheduler.Worker method), 120
 update_id() (luigi.contrib.esindex.CopyToIndex method), 64, 148
 update_id() (luigi.contrib.rdbms.CopyToTable method), 77, 161
 update_id() (luigi.contrib.redshift.KillOpenRedshiftSessions method), 81, 165
 update_id() (luigi.contrib.sqla.CopyToTable method), 90, 175
 update_resources() (luigi.rpc.RemoteScheduler method), 116
 update_resources() (luigi.scheduler.CentralPlannerScheduler method), 121
 upload() (luigi.contrib.webhdfs.WebHdfsClient method), 92, 177
 upstream (class in luigi.tools.deps), 93, 177
 UPSTREAM_SEVERITY_KEY() (in module luigi.scheduler), 119
 use_cmdline_section (luigi.interface.core attribute), 106
 use_cmdline_section (luigi.task.Task attribute), 127
 use_db_timestamps (luigi.contrib.redshift.RedshiftTarget attribute), 78, 162
 use_db_timestamps (luigi.postgres.PostgresTarget attribute), 114
 user (luigi.contrib.rdbms.CopyToTable attribute), 77, 161
 user (luigi.contrib.redshift.KillOpenRedshiftSessions attribute), 81, 165

V

value (luigi.db_task_history.TaskParameter attribute), 101
 value (luigi.parameter.Parameter attribute), 111
 version (luigi.contrib.hdfs.config.hadoopcli attribute), 53, 136
 visualization_graph (luigi.scheduler.scheduler attribute), 119

W

wait_interval (luigi.worker.worker attribute), 134
 walk() (luigi.contrib.webhdfs.WebHdfsClient method), 92, 177
 WebHdfsClient (class in luigi.contrib.webhdfs), 92, 177

WebHdfsTarget (class in luigi.contrib.webhdfs), 92, 177
 Week (class in luigi.date_interval), 99
 Worker (class in luigi.scheduler), 120
 Worker (class in luigi.worker), 134
 worker (class in luigi.worker), 134
 worker_disconnect_delay (luigi.scheduler.scheduler attribute), 119
 worker_list() (luigi.rpc.RemoteScheduler method), 116
 worker_list() (luigi.scheduler.CentralPlannerScheduler method), 122
 worker_timeout (luigi.task.Task attribute), 127
 workers (luigi.interface.core attribute), 106
 WorkerSchedulerFactory (class in luigi.interface), 106
 wrap_traceback() (in module luigi.notifications), 109
 WrappedFormat (class in luigi.format), 104
 wrapper_cls (luigi.format.MixedUnicodeBytesFormat attribute), 104
 wrapper_cls (luigi.format.NewlineFormat attribute), 104
 wrapper_cls (luigi.format.TextFormat attribute), 104
 WrapperTask (class in luigi.task), 129
 writable() (luigi.format.InputPipeProcessWrapper method), 103
 writable() (luigi.format.OutputPipeProcessWrapper method), 103
 writable() (luigi.s3.ReadableS3File method), 117
 write() (luigi.format.MixedUnicodeBytesWrapper method), 104
 write() (luigi.format.NewlineWrapper method), 103
 write() (luigi.format.OutputPipeProcessWrapper method), 103
 WRITE_APPEND (luigi.contrib.bigquery.WriteDisposition attribute), 58, 142
 write_disposition (luigi.contrib.bigquery.BigqueryLoadTask attribute), 60, 144
 write_disposition (luigi.contrib.bigquery.BigqueryRunQueryTask attribute), 61, 145
 WRITE_EMPTY (luigi.contrib.bigquery.WriteDisposition attribute), 58, 142
 write_pid() (in module luigi.process), 115
 WRITE_TRUNCATE (luigi.contrib.bigquery.WriteDisposition attribute), 58, 142
 WriteDisposition (class in luigi.contrib.bigquery), 58, 142
 writeLine() (luigi.format.OutputPipeProcessWrapper method), 103
 writelines() (luigi.format.MixedUnicodeBytesWrapper method), 104
 writelines() (luigi.format.NewlineWrapper method), 103
 writer() (luigi.contrib.hadoop.JobTask method), 70, 154
 WRITES_BEFORE_FLUSH (luigi.format.OutputPipeProcessWrapper attribute), 103

Y

Year (class in luigi.date_interval), 99